A Review of Large Language Models and Autonomous Agents in Chemistry

Mayk Caldas Ramos

FutureHouse Inc., San Francisco, CA

Department of Chemical Engineering

University of Rochester, Rochester, NY

mcaldasr@ur.rochester.edu

Christopher J. Collison

School of Chemistry and Materials Science

Rochester Institute of Technology, Rochester, NY

cjcscha@rit.edu

Andrew D. White[*]

FutureHouse Inc., San Francisco, CA

Department of Chemical Engineering

University of Rochester, Rochester, NY

andrew@futurehouse.org

November 18, 2024

ABSTRACT

Large language models (LLMs) have emerged as powerful tools in chemistry, significantly impacting molecule design, property prediction, and synthesis optimization. This review highlights LLM capabilities in these domains and their potential to accelerate scientific discovery through automation. We also review LLM-based autonomous agents: LLMs with a broader set of tools to interact with their surrounding environment. These agents perform diverse tasks such as paper scraping, interfacing with automated laboratories, and synthesis planning. As agents are an emerging topic, we extend the scope of our review of agents beyond chemistry and discuss across any scientific domains. This review covers the recent history, current capabilities, and design of LLMs and autonomous agents, addressing specific challenges, opportunities, and future directions in chemistry. Key challenges include data quality and integration, model interpretability, and the need for standard benchmarks, while future directions point towards more sophisticated multi-modal agents and enhanced collaboration between agents and experimental methods. Due to the quick pace of this field, a repository has been built to keep track of the latest studies: `https://github.com/ur-whitelab/LLMs-in-science`.

Keywords Large Language Model, LLM, LLM agent, agent, science, chemistry

Contents

---

[*]Corresponding author

1 Introduction

The integration of Machine Learning (ML) and Artificial Intelligence (AI) into chemistry has spanned several decades.[1–10] Although applications of computational methods in quantum chemistry and molecular modeling from the 1950s-1970s were not considered AI, they laid the groundwork. Subsequently in the 1980s expert systems like DENDRAL[11,12] were expanded to infer molecular structures from mass spectrometry data.[13] At the same time, Quantitative Structure-Activity Relationship (QSAR) Models were developed[5] that would use statistical methods to predict the effects of chemical structure on activity.[14–17] In the 1990s, neural networks, and associated Kohonen Self-Organizing Maps were introduced to domains such as drug design,[18,19] as summarized well by Yang et al.[5] and Goldman and Walters[20], although they were limited by the computational resources of the time. With an explosion of data from High-Throughput Screening (HTS),[21,22] models then started to benefit from vast datasets of molecular structures and their biological activities. Furthermore, ML algorithms such as Support Vector Machines and Random Forests became popular for classification and regression tasks in cheminformatics,[1] offering improved performance over traditional statistical methods.[23]

Deep learning transformed the landscape of ML in chemistry and materials science in the 2010s.[24] Recurrent Neural Networks (RNNs),[25–29] Convolutional Neural Networks (CNNs)[30–32] and later, Graph Neural Networks (GNNs),[33–38] made great gains in their application to molecular property prediction, drug discovery,[39] and synthesis prediction.[40] Such methods were able to capture complex patterns in data, and therefore enabled the identification of novel materials for high-impact needs such as energy storage and conversion.[41,42]

In this review, we explore the next phase of AI in chemistry, namely the use of Large Language Models (LLMs) and autonomous agents. Inspired by successes in natural language processing (NLP), LLMs were adapted for chemical language (e.g., Simplified Molecular Input Line Entry System (SMILES)[43]) to tackle tasks from synthesis prediction to molecule generation.[44–46] We will then explore the integration of LLMs into autonomous agents as illustrated by M. Bran et al.[47] and Boiko et al.[48], which may be used for data interpretation or, for example, to experiment with robotic systems. We are at a crossroads where AI enables chemists to solve major global problems faster and streamline routine lab tasks. This enables, for instance, the development of larger, consistent experimental datasets and shorter lead times for drug and material commercialization. As such, language has been the preferred mechanism for describing and disseminating research results and protocols in chemistry for hundreds of years.[49]

1.1 Challenges in Chemistry

We categorize some key challenges that can be addressed by AI in chemistry as: Property Prediction, Property-Directed Molecule Generation, and Synthesis Prediction. These categories, as illustrated in Figure 1 can be connected to a fourth challenge in automation. The first task is to predict a property for a given compound to decide if it should be synthesized for a specific application, such as an indicator,[50] light harvester,[51] or catalyst.[52] To achieve better models for property prediction, high-quality data is crucial. We discuss the caveats and issues with the current datasets in Section 3.1 and illustrate state-of-the-art findings in Section 3.2.
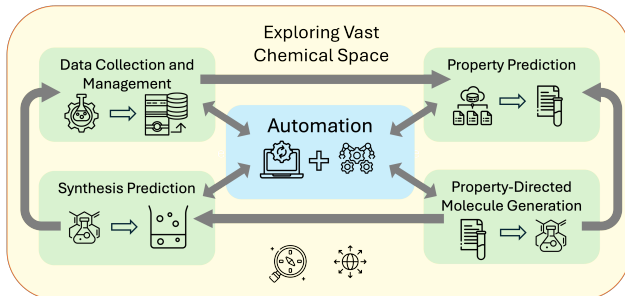
Figure 1: AI-powered LLMs accelerate chemical discovery with models that address key challenges in Property Prediction, Property Directed Molecule Generation, and Synthesis Prediction. Autonomous agents connect these models and additional tools thereby enabling rapid exploration of vast chemical spaces.

The second task is to generate novel chemical structures that meet desired chemical profiles or exhibit properties.[53] Success in this area would accelerate progress in various chemical applications, but reliable reverse engineering (inverse

design)[54] is not yet feasible over the vast chemical space.[55] For instance, inverse design, when coupled with automatic selection of novel structures (*de novo* molecular design) could lead to the development of drugs targeting specific proteins while retaining properties like solubility, toxicity, and blood-brain barrier permeability.[56] The complexity of connecting *de novo* design with property prediction is high and we show how state-of-the-art models currently perform in Section 3.3.

Once a potential target molecule has been identified, the next challenge is predicting its optimal synthesis using inexpensive, readily available, and non-toxic starting materials. In a vast chemical space, there will always be an alternative molecule "B" that has similar properties to molecule "A" but is easier to synthesize. Exploring this space to find a new molecule with the right properties and a high-yield synthesis route brings together these challenges. The number of possible stable chemicals is estimated to be up to 10^{180}.[57–60] Exploring this vast space requires significant acceleration beyond current methods.[61] As Restrepo[57] emphasizes, cataloguing failed syntheses is essential to building a comprehensive dataset of chemical features. Autonomous chemical resources can accelerate database growth and tackle this challenge. Thus, automation is considered a fourth major task in chemistry.[62–65] The following discussion explores how LLMs and autonomous agents can provide the most value. Relevant papers are discussed in Section 3.4

This review is organized within the context of these categories. The structure of the review is as follows. Section 2 provides an introduction to transformers, including a brief description of encoder-only, decoder-only and encoder-decoder architectures. Section 3 provides a detailed survey of work with LLMs, where we connect each transformer architecture to the areas of chemistry that it is best suited to support. We then progress into a description of autonomous agents in section 4, and a survey of how such LLM-based agents are finding application in chemistry-centered scientific research, section 5. After providing some perspective on future challenges and opportunities in section 6, and we conclude in section 7. We distinguish between "text-based" and "mol-based" inputs and outputs, with "text" referring to natural language and "mol" referring to the chemical syntax for material structures, as introduced by Zhang et al.[66].

2   Large Language Models

The prior state-of-the-art for sequence-to-sequence (seq2seq) tasks had been the Recurrent Neural Network (RNN),[67] typically as implemented by Hochreiter and Schmidhuber[68]. In a seq2seq task, an input sequence, such as a paragraph in English, is processed to generate a corresponding output sequence, such as a translation into French. The RNN retains "memory" of previous steps in a sequence to predict later parts. However, as sequence length increases, gradients can become vanishingly small or explosively large,[69,70] preventing effective use of earlier information in long sequences. Due to these limitations, RNNs have thus fallen behind Large Language Models (LLMs), which primarily implement transformer architectures, introduced by Vaswani et al.[71]. LLMs are deep neural networks (NN) characterized by their vast number of parameters and, though transformers dominate, other architectures for handling longer input sequences are being actively explored.[72–75] A detailed discussion of more generally applied LLMs can be found elsewhere.[76] Since transformers are well-developed in chemistry and are the dominant paradigm behind nearly all state-of-the-art sequence modeling results, they are a focus in this review.

2.1   The Transformer

The transformer was introduced in, "Attention is all you need" by Vaswani et al.[71] in 2017. A careful line-by-line review of the model can be found in "The Annotated Transformer".[77] The transformer was the first seq2seq model based entirely on attention mechanisms, although attention had been a feature for RNNs some years prior.[78] The concept of "attention" is a focus applied to certain words of the input, which would convey the most importance, or the context of the passage, and thereby would allow for better decision-making and greater accuracy. However, in a practical sense, "attention" is implemented simply as the dot-product between token embeddings and a learned non-linear function, which will be described further below.

Context Window   Large language models are limited by the size of their context window, which represents the maximum number of input tokens they can process at once. This constraint arises from the quadratic computational cost of the transformer's attention mechanism, which restricts effective input to a few thousand tokens.[79] Hence, LLM-based agents struggle to maintain coherence and capture long-range dependencies in extensive texts or complex dialogues, impacting their performance in applications requiring deep contextual understanding.[80] These limitations and strategies to overcome them are better discussed in Section 4.

Tokenization   In NLP tasks, the natural language text sequence, provided in the context window, is first converted to a list of tokens, which are integers that each represent a fragment of the sequence. Hence the input is numericized according to the model's vocabulary following a specific tokenization scheme.[81–85]
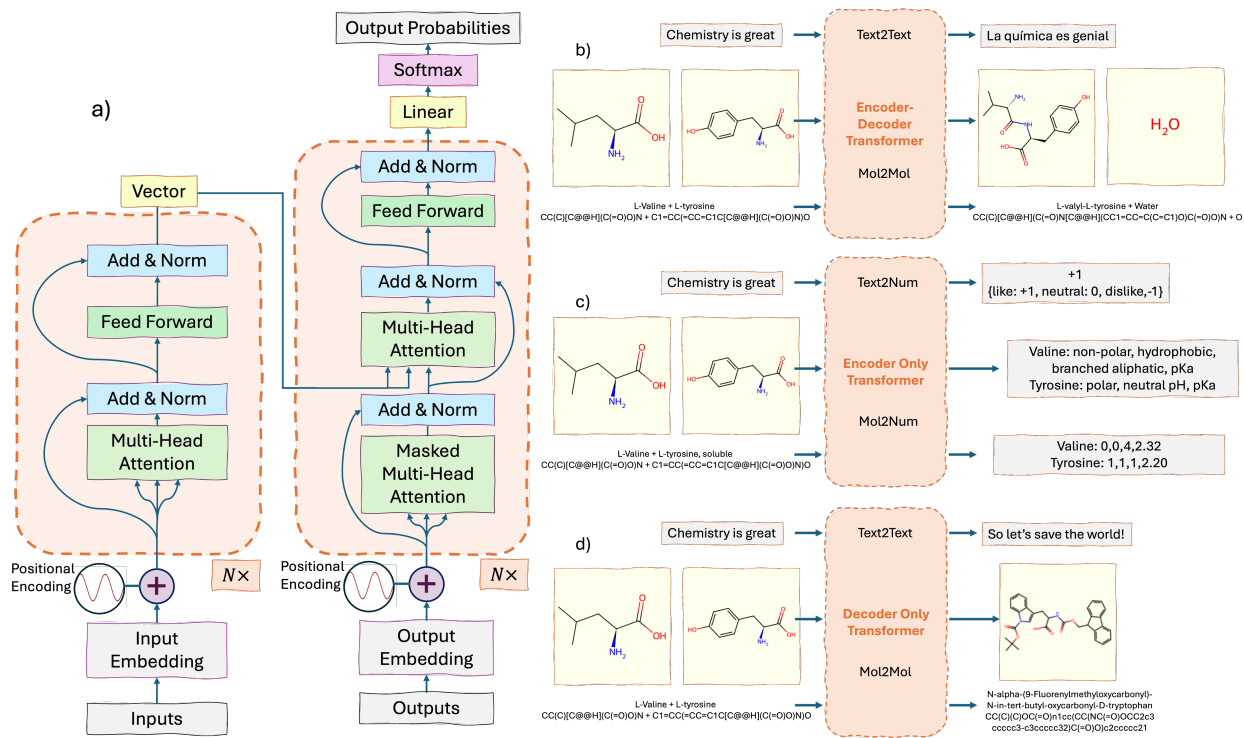
Figure 2: a) The generalized encoder-decoder transformer: The encoder on the left converts an input into a vector, while the decoder on the right predicts the next token in a sequence. b) Encoder-decoder transformers are traditionally used for translation tasks and, in chemistry, for reaction prediction, translating reactants into products. c) Encoder-only transformers provide a vector output and are typically used for sentiment analysis. In chemistry, they are used for property prediction or classification tasks. d) Decoder-only transformers generate likely next tokens in a sequence. In chemistry, they are used to generate new molecules given an instruction and description of molecules.

Input Embeddings   Each token is then converted into a vector in a process called input embedding. This vector is a learned representation that positions tokens in a continuous space based on their semantic relationships. This process allows the model to capture similarities between tokens, which is further refined through mechanisms like attention (discussed below) that weigh and enhance these semantic connections.

Positional Encoding   A positional encoding is then added, which plays a major role in transformer success. It is added to the input embeddings to provide information about the order of elements in a sequence, as transformers lack a built-in notion of sequence position. Vaswani et al.[71] reported similar performance with both fixed positional encoding based on sine and cosine functions, and learned encodings. However, many options for positional embeddings exist.[86] In fixed positional encoding, the position of each element in a sequence is encoded using sine and cosine functions with different frequencies, depending on the element's position. This encoding is then added to the word's vector representation (generated during the tokenization and embedding process). The result is a modified vector that encodes both the meaning of the word and its position within the sequence. These sine and cosine functions generate values within a manageable range of -1 to 1, ensuring that each positional encoding is unique and that the encoding is unaffected by sequence length.

Attention   The concept of "attention" is central to the transformer's success, especially during training. Attention enables the model to focus on the most relevant parts of the input data. It operates by comparing each element in a sequence, such as a word, to every other element. Each element serves as a *query*, compared against other elements called *keys*, each associated with a corresponding value. The alignment between a *query* and a *keys*, determines the strength of their connection, represented by an *attention weight*.[87] These weights highlight the importance of certain elements by scaling their associated values accordingly. During training, the model learns to adjust these weights, capturing relationships and contextual information within the sequence. Once trained, the model uses these learned weights to integrate information from different parts of the sequence, ensuring that its output remains coherent and contextually aligned with the input.

5

The transformer architecture is built around two key modules: the encoder and the decoder. Figure 2a provides a simplified diagram of the general encoder-decoder transformer architecture. The input is The input is tokenized, from the model's vocabulary,[81–85] embedded and positionally encoded, as described above. The encoder consists of multiple stacked layers (six layers in the original model),[71] with each layer building on the outputs of the previous one. Each token is represented as a vector, that gets passed through these layers. At each encoder layer, a self-attention mechanism is applied, which calculates the attention between tokens, as discussed earlier. Afterward, the model uses normalization and adds the output back to the input through what's called a residual connection. Residual connection is represented in Figure 2a by the "by-passing" arrow. This bypass helps prevent issues with vanishing gradients,[69,70] ensuring that information flows smoothly through the model. The final step in each encoder layer is a feed-forward neural network with an activation function (such as ReLU,[88] SwiGLU,[89] GELU,[90] etc) that further refines the representation of the input.

The decoder works similarly to the encoder but with key differences. It starts with an initial input token – usually a special start token—embedded into a numerical vector. This token initiates the output sequence generation. Positional encodings are applied to preserve the token order. The decoder is composed of stacked layers, each containing a masked self-attention mechanism that ensures the model only attends to the current and previous tokens, preventing access to future tokens. Additionally, an encoder-decoder attention mechanism aligns the decoder's output with relevant encoder inputs, as depicted by the connecting arrows in Figure 2a. This alignment helps the model focus on the most critical information from the input sequence. Each layer also employs normalization, residual connections, and a feed-forward network. The final layer applies a softmax function, converting the scores into a probability density over the vocabulary of tokens. The decoder generates the sequence autoregressively, predicting each token based on prior outputs until an end token signals termination.

2.2 Model training

The common lifetime of an LLM consists of being first pretrained using self-supervised techniques, generating what is called a base model. Effective prompt engineering may lead to successful task completion but this base model is often fine-tuned for specific applications using supervised techniques and this creates the "instruct model." It is called the "instruct model" because the fine-tuning is usually done for it to follow arbitrary instructions, removing the need to specialize fine-tuning for each downstream task.[91] Finally, the instruct model can be further tuned with reward models to improve human preference or some other non-differentiable and sparse desired character.[92] These concepts are expanded on below.

Self-supervised Pretraining   A significant benefit implied in all the transformer models described in this review is that self-supervised learning takes place with a vast corpus of text. Thus, the algorithm learns patterns from unlabeled data, which opens up the model to larger datasets that may not have been explicitly annotated by humans. The advantage is to discover underlying structures or distributions without being provided with explicit instructions on what to predict, nor with labels that might indicate the correct answer.

Prompt Engineering   The model's behavior can be guided by carefully crafting input prompts that leverage the pretrained capabilities of LLMs. Since the original LLM remains unchanged, it retains its generality and can be applied across various tasks.[93] However, this approach relies heavily on the assumption that the model has adequately learned the necessary domain knowledge during pretraining to achieve an appropriate level of accuracy in a specific domain. Prompt engineering can be sensitive to subtle choices of language; small changes in wording can lead to significantly different outputs, making it challenging to achieve consistent results and to quantify the accuracy of the outputs.[94]

Supervised Fine-tuning   After this pretraining, many models described herein are fine-tuned on specific downstream tasks (e.g., text classification, question answering) using supervised learning. In supervised learning, models learn from labeled data, and map inputs to known outputs. Such fine-tuning allows the model to be adjusted with a smaller, task-specific dataset to perform well on that downstream task.

LLM Alignment   A key step after model training is aligning the output with human preferences. This process is critical to ensure that the large language model (LLM) produces outputs that are not only accurate but also reflect appropriate style, tone, and ethical considerations. Pretraining and fine-tuning often do not incorporate human values, so alignment methods are essential to adjust the model's behavior, including reducing harmful outputs.[95]

One important technique for LLM alignment is instruction tuning. This method refines the model by training it on datasets that contain specific instructions and examples of preferred responses. By doing so, the model learns to generalize from these examples and follow user instructions more effectively, leading to outputs that are more relevant

and safer for real-world applications.[96,97] Instruction tuning establishes a baseline alignment, which can then be further improved in the next phase using reinforcement learning (RL).[98]

In RL-based alignment, the model generates tokens as actions and receives rewards based on the quality of the output, guiding the model to optimize its behavior over time. Unlike post-hoc human evaluations, RL actively integrates preference feedback during training, refining the model to maximize cumulative rewards. This approach eliminates the need for token-by-token supervised fine-tuning by focusing on complete outputs, which better capture human preferences.[99–101]

The text generation process in RL is typically modeled as a Markov Decision Process (MDP), where actions are tokens, and rewards reflect how well the final output aligns with human intent.[102] A popular method, Reinforcement Learning with Human Feedback (RLHF),[103] leverages human input to shape the reward system, ensuring alignment with user preferences. Variants such as reinforcement learning with synthetic feedback (RLSF),[104] Proximal Policy Optimization (PPO),[105] and REINFORCE[106] offer alternative strategies for assigning rewards and refining model policies.[99,102,107,108] A broader exploration of RL's potential in fine-tuning LLMs is available in works by Cao et al.[109] and Shen et al.[95]

There are ways to reformulate the RLHF process into a direct optimization problem with a different loss. This is known as reward-free metods. Among the main examples of reward-free methods, we have the direct preference optimization (DPO),[110] Rank Responses to align Human Feedback (RRHF),[111] and Preference Ranking Optimization (PRO).[112] These models are popular competitors to PPO and other reward-based methods due to its simplicity. It overcomes the lack of token-by-token loss signal by comparing two completions at a time. The discussions about which technique is superior remain very active in the literature.[113]

Finally, the alignment may not be to human preferences but to downstream tasks that do not provide token-by-token rewards. For example, Bou et al.[114] and Hayes et al.[115] both use RL on a language model for improving its outputs on a downstream scientific task.

2.3 Model types

While the Vaswani Transformer[71] employed an encoder-decoder structure for sequence-to-sequence tasks, the encoder and decoder were ultimately seen as independent models, leading to "encoder-only", and "decoder-only" models described below.

Examples of how such models can be used are provided in Figures 2b, c, and d. Figure 2b illustrates the encoder-decoder model's capability to transform sequences, such as translating from English to Spanish or predicting reaction products by mapping atoms from reactants (amino acids) to product positions (a dipeptide and water). This architecture has large potential on sequence-to-sequence transformations.[116,117] Figure 2c highlights the strengths of an encoder-only model in extracting properties or insights directly from input sequences. For example, in text analysis, it can assign sentiment scores or labels, such as tagging the phrase "Chemistry is great" with a positive sentiment. In chemistry, it can predict molecular properties, like hydrophobicity or pKa, from amino acid representations, demonstrating its applications in material science and cheminformatics.[118–120] Finally, Figure 2d depicts a decoder-only architecture, ideal for tasks requiring sequence generation or completion. This model excels at inferring new outputs from input prompts. For instance, given that "chemistry is great," it can propose broader implications or solutions. It can also generate new peptide sequences from smaller amino acid fragments, showcasing its ability to create novel compounds. This generative capacity is particularly valuable in drug design, where the goal is to discover new molecules or expand chemical libraries.[44,121–123]

2.3.1 Encoder-only Models

Beyond Vaswani's transformer,[71] used for sequence-to-sequence tasks, another significant evolutionary step forward came in the guise of the Bidirectional Encoder Representations from Transformers, or "BERT", described in October 2018 by Devlin et al.[87] BERT utilized only the encoder component, achieving state-of-the-art performance on sentence-level and token-level tasks, outperforming prior task-specific architectures.[87] The key difference was BERT's bidirectional transformer pretraining on unlabeled text, meaning the model processes the context both to the left and right of the word in question, facilitated by a Masked Language Model (MLM). This encoder-only design allowed BERT to develop more comprehensive representations of input sequences, rather than mapping input sequences to output sequences. In pretraining, BERT also uses Next Sentence Prediction (NSP). "Sentence" here means an arbitrary span of contiguous text. The MLM task randomly masks tokens and predicts them by considering both preceding and following contexts simultaneously, inspired by Taylor.[124] NSP predicts whether one sentence logically follows another, training the model to understand sentence relationships. This bidirectional approach allows BERT to recognize greater nuance and richness in the input data.

Subsequent evolutions of BERT include, for example, RoBERTa, (Robustly optimized BERT approach), described in 2019 by Liu et al.[125]. RoBERTa was trained on a larger corpus, for more iterations, with larger mini-batches, and longer sequences, improving model understanding and generalization. By removing the NSP task and focusing on the MLM task, performance improved. RoBERTa dynamically changed masked positions during training and used different hyperparameters. Evolutions of BERT also include domain-specific pretraining and creating specialist LLMs for fields like chemistry, as described below (see Section 3).

2.3.2 Decoder-only Models

In June 2018, Radford et al.[126] proposed the Generative Pretrained Transformer (GPT) in their paper, "Improving Language Understanding by Generative Pretraining". GPT used a decoder-only, left-to-right unidirectional language model to predict the next word in a sequence based on previous words, without an encoder. Unlike earlier models, GPT could predict the next sequence, applying a general language understanding to specific tasks with smaller annotated datasets.

GPT employed positional encodings to maintain word order in its predictions. Its self-attention mechanism prevented tokens from attending to future tokens, ensuring each word prediction depended only on preceding words. Hence a decoder-only architecture represents a so-called causal language model, one that generates each item in a sequence based on the previous items. This approach is also referred to as "autoregressive", meaning that each new word is predicted based on the previously generated words, with no influence from future words. The generation of each subsequent output is causally linked to the history of generated outputs and nothing ahead of the current word affects its generation.

2.3.3 Encoder-decoder Models

Evolving further, BART (Bidirectional and Auto-Regressive Transformers) was introduced by Lewis et al. in 2019.[127] BART combined the context learning strengths of the bidirectional BERT, and the autoregressive capabilities of models like GPT, which excel at generating coherent text. BART was thus a hybrid seq2seq model, consisting of a BERT-like bidirectional encoder and a GPT-like autoregressive decoder. This is nearly the same architecture as Vaswani et al.[71]; the differences are in the pretraining. BART was pretrained using a task that corrupted text by, for example, deleting tokens, and shuffling sentences. It then learned to reconstruct the original text with left-to-right autoregressive decoding as in GPT models.

2.3.4 Multi-task and Multi-modal Models

In previous sections, we discussed LLMs that take natural language text as input and then output either a learned representation or another text sequence. These models traditionally perform tasks like translation, summarization, and classification. However, multi-task models are capable of performing several different tasks using the same model, even if those tasks are unrelated. This allows a single model to be trained on multiple objectives, enhancing its versatility and efficiency, as it can generalize across various tasks during inference.

Multi-task models, such as the Text-to-Text Transfer Transformer (T5) developed by Raffel et al.[128] demonstrate that various tasks can be reframed into a text-to-text format, allowing the same model architecture and training procedure to be applied universally. By doing so, the model can be used for diverse tasks, but all with the same set of weights. This reduces the need for task-specific models and increases the model's adaptability to new problems. The relevance of this approach is particularly significant as it enables researchers to tackle multiple tasks without needing to retrain separate models, saving both computational resources and time. For instance, Flan-T5[129] used instruction fine-tuning with chain-of-thought prompts, enabling it to generalize to unseen tasks, such as generating rationales before answering. This fine-tuning expands the model's ability to tackle more complex problems. More advanced approaches have since been proposed to build robust multi-task models that can flexibly switch between tasks at inference time.[130–133]

Additionally, LLMs have been extended to process different input modalities, such as image and sound, even though they initially only processed text. For example, Fuyu[134] uses linear projection to adapt image representations into the token space of an LLM, allowing a decoder-only model to generate captions for figures. Expanding on this, next-GPT[135] was developed as an "any-to-any" model, capable of processing multiple modalities, such as text, audio, image, and video, through modality-specific encoders. The encoded representation is projected into a decoder-only token space, and the LLM's output is processed by a domain-specific diffusion model to generate each modality's output. Multitask or multimodel methods are further described below as these methods start to connect LLMs with autonomous agents.
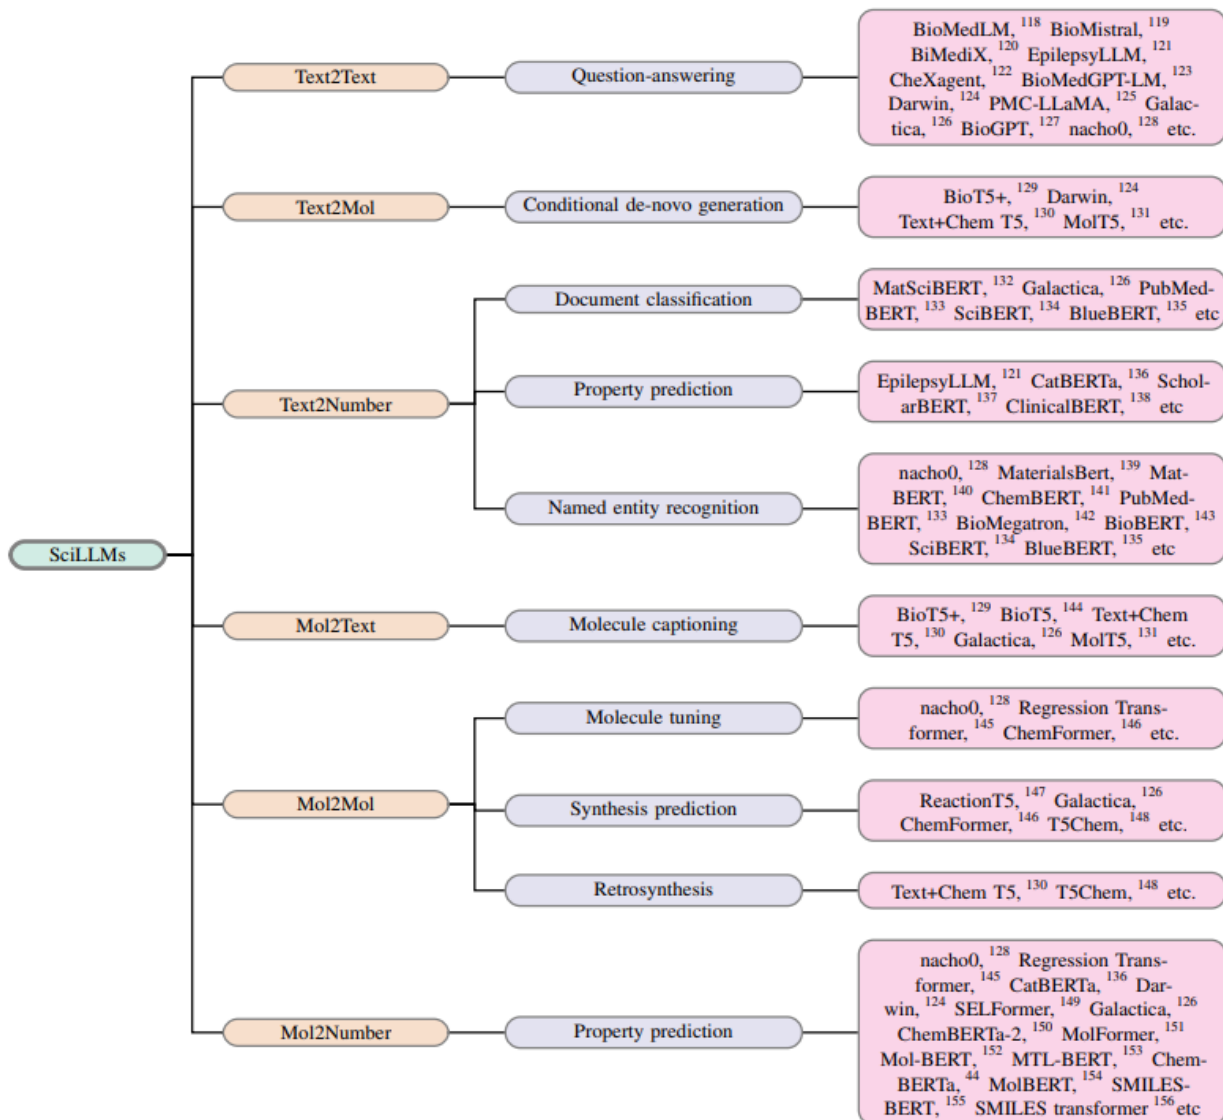
Figure 3: Classification of LLMs in chemistry and biochemistry according to their application.

3 LLMs for Chemistry and Biochemistry

The integration of large language models (LLMs) into chemistry and biochemistry is opening new frontiers in molecular design, property prediction, and synthesis. As these models evolve, they increasingly align with specific chemical tasks, capitalizing on the strengths of their architectures. Specifically, encoder-only models excel at property prediction,[118] decoder-only models are suited for inverse design,[136] and encoder-decoder models are applied to synthesis prediction.[137] However, with the development improvement of decoder-only models[138] and the suggestion that regression tasks can be reformulated as a text completion task,[139] decoder-only models started being also applied for property prediction.[140–143] This section surveys key LLMs that interpret chemical languages like SMILES and InChI, as well as those that process natural language descriptions relevant to chemistry.

We provide a chronological perspective on the evolution of LLMs in this field (Figure 4), presenting broadly on the design, functionality, and value of each model. Our approach primarily centers on models that use chemical representations like SMILES strings as inputs, but we also examine how natural language models extract valuable data from scientific literature to enhance chemical research.
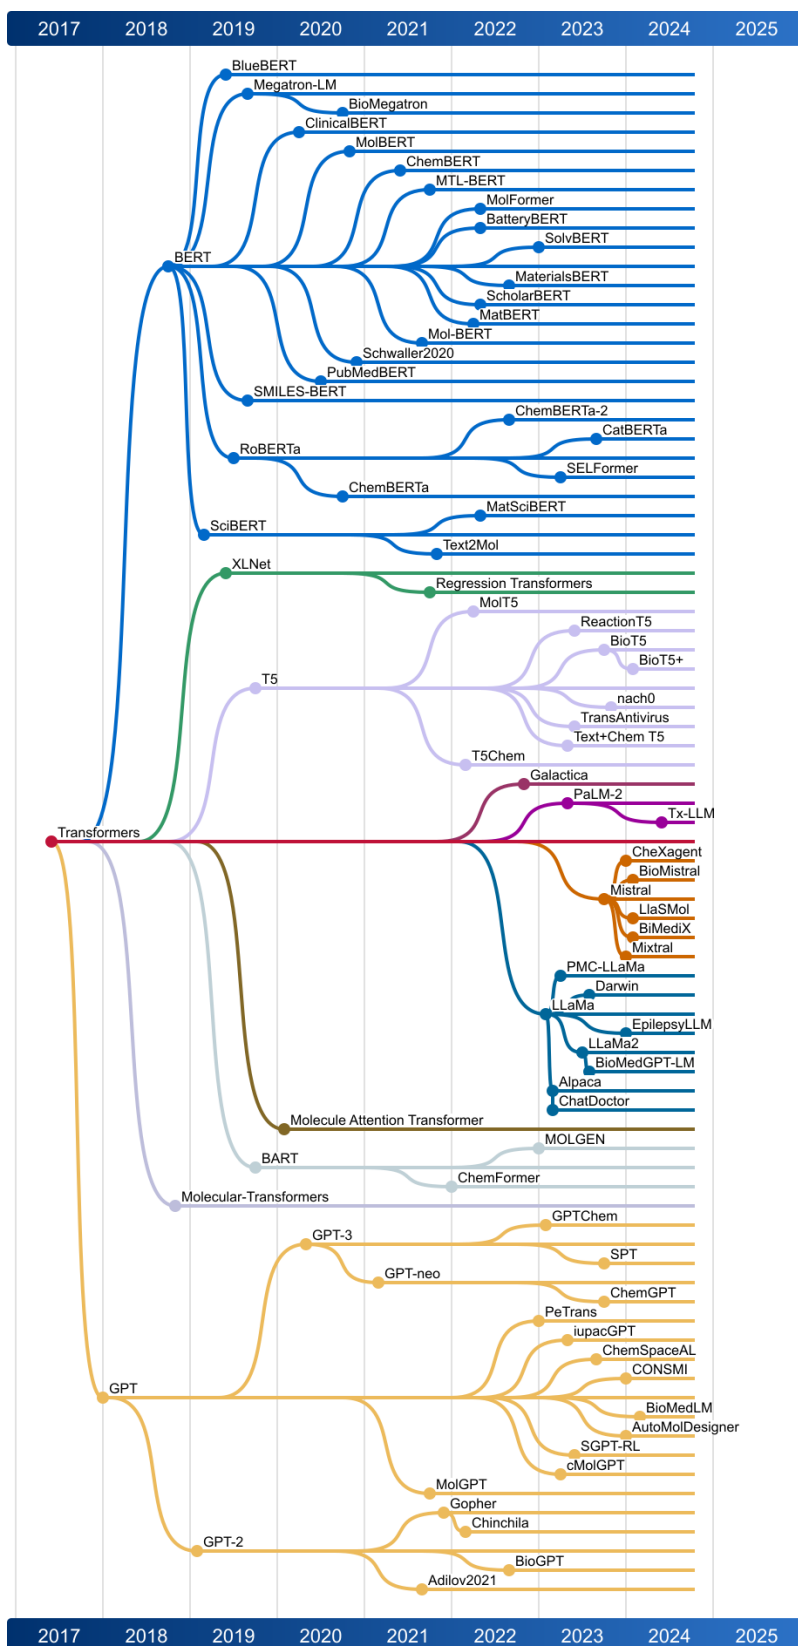
Figure 4: Illustration of how Large Language Models (LLMs) evolved chronologically. The dates display the first publication of each model.

Ultimately, this discussion underscores the potential for mol-based and text-based LLMs to work together, addressing the growing opportunity for automation in chemistry. This sets the stage for a broader application of autonomous agents in scientific discovery. Figure 3 illustrates the capabilities of different LLMs available currently, while Figure 4 presents a chronological map of LLM development in chemistry and biology.

Of critical importance, this section starts by emphasizing the role of trustworthy datasets and robust benchmarks. Without well-curated, diverse datasets, models may fail to generalize across real-world applications. Benchmarks that are too narrowly focused can limit the model's applicability, preventing a true measure of its potential. While natural language models take up a smaller fraction of this section, these models will be increasingly used to curate these datasets, ensuring data quality becomes a key part of advancing LLM capabilities in chemistry.

3.1   Molecular Representations, Datasets, and Benchmarks

Molecules can be described in a variety of ways, ranging from two-dimensional structural formulas to more complex three-dimensional models that capture electrostatic potentials. Additionally, molecules can be characterized through properties such as solubility, reactivity, or spectral data from techniques like NMR or mass spectrometry. However, to leverage these descriptions in machine learning, they must be converted into a numerical form that a computer can process. Given the diversity of data in chemistry-based machine learning, multiple methods exist for representing molecules,[144–149] highlighting this heterogeneity. Common representations include molecular graphs,[150–152] 3D point clouds,[153–156] and quantitative feature descriptors.[145,157–160] In this review, we focus specifically on string-based representations of molecules, given the interest in language models. Among the known string representations, we can cite IUPAC names, SMILES,[43] DeepSMILES,[161] SELFIES,[162] and InChI,[163] as recently reviewed by Das et al.[164]
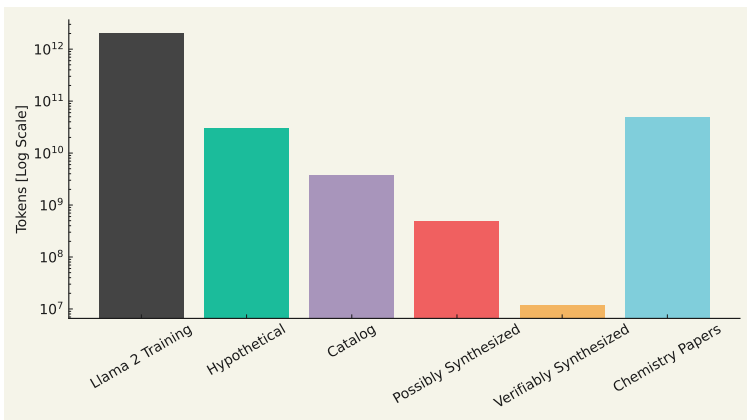

Figure 5: Number of training tokens (on log scale) available from various chemical sources compared with typical LLM training runs. The numbers are drawn from ZINC,[165] PubChem,[166] Touvron et al.[167], ChEMBL,[168] and Kinney et al.[169]

.

Regarding datasets, there are two types of data used for training LLMs, namely training data and evaluation data. Training data should be grounded in real molecular structures to ensure the model develops an accurate representation of what constitutes a valid molecule. This is similar to how natural language training data, such as that used in models like GPT-4, must be based on real sentences or code to avoid generating nonsensical outputs. Figure 5 shows a comparison of the number of tokens in common chemistry datasets with those used to train LLaMA2, based on literature data.[165–169] With this in mind, we note the largest chemical training corpus, which largely comprises hypothetical chemical structures, amounts to billions of tokens, almost two orders of magnitude fewer than the trillions of tokens used to train LLaMA2. When excluding hypothetical structures from datasets like ZINC,[165] (Figure 5), the number of tokens associated with verifiably synthesized compounds is over five orders of magnitude lower than that of LLaMA2's training data. To address this gap, efforts such as the Mol-instructions dataset, curated by Fang et al.[170], prioritize quality over quantity, providing ∼2M biomolecular and protein-related instructions. Mol-instructions[170] was selectively built from multiple data sources,[56,171–180] with rigorous quality control. Given the success of literature-based LLMs, one may naturally assume that large datasets are of paramount importance for chemistry. However, it is crucial not to overlook the importance of data quality. Segler et al.[181] demonstrated that even using the Reaxys dataset, a very small, human-curated collection of chemical reactions, was sufficient to achieve state-of-the-art results in retrosynthesis. Therefore, the issue is not merely a lack of data, but rather a lack of high-quality data that may be the pivotal factor

11

holding back the development of better scientific LLMs. Ultimately, the focus must shift from sheer quantity to the curation of higher-quality datasets to advance these models.

To evaluate the accuracy of these models, we compare their performance against well-established benchmarks. However, if the benchmarks are not truly representative of the broader chemistry field, it becomes difficult to gauge the expected impact of these models. Numerous datasets, curated by the scientific community, are available for this benchmarking.[182,183] Among them, MoleculeNet,[56] first published in 2017, is the most commonly used labeled dataset for chemistry. However, MoleculeNet has several limitations: it is small, contains errors and inconsistencies, and lacks relevance to a larger number of real-world chemistry problems.[184–187] Pat Walters, a leader in ML for drug discovery, has emphasized, "I think the best way to make progress on applications of machine learning to drug discovery is to fund a large public effort that will generate high-quality data and make this data available to the community".[188]

Walters provides several constructive critiques noting, for example, that the QM7, QM8, and QM9 datasets, intended for predicting quantum properties from 3D structures, are often misused with predictions based incorrectly on their 1D SMILES strings, which inadequately represent 3D molecular conformations. He also suggests more relevant benchmarks and also datasets with more valid entries. For example, he points to the Absorption, Distribution, Metabolism, and Excretion (ADME) data curated by Fang et al.[189], as well as the Therapeutic Data Commons (TDC)[190,191] and TDC-2.[192] These datasets contain measurements of real compounds, making them grounded in reality. Moreover, ADME is crucial for determining a drug candidate's success, while therapeutic results in diverse modalities align with metrics used in drug development.

Here, we hypothesize that the lack of easily accessible, high-quality data in the correct format for training foundational chemical language models is a major bottleneck to the development of the highly desired "super-human" AI-powered digital chemist. A more optimistic view is presented by Rich and Birnbaum[193] They argue that we do not need to wait for the creation of new benchmarks. Instead, they suggest that even the currently available, messy public data can be carefully curated to create benchmarks that approximate real-world applications. In addition, we argue that extracting data from scientific chemistry papers might be an interesting commitment to generating data of high quality, grounded to the truth, and on a large scale.[194] Some work has been done in using LLMs for data extraction.[195,196] Recently, a few benchmarks following these ideas were created for evaluating LLMs' performance in biology (LAB-Bench[197]) and material science (MatText,[198] MatSci-NLP[199] and MaScQA[200]).

3.2 Property Prediction and Encoder-only Mol-LLMs

Encoder-only transformer architectures are primarily composed of an encoder, making them well-suited for chemistry tasks that require extracting meaningful information from input sequences, such as classification and property prediction. Since encoder-only architectures are mostly applied to capturing the underlying structure-property relationships, we describe here the relative importance of the property prediction task. Sultan et al.[201] also discussed the high importance of this task, the knowledge obtained in the last years, and the remaining challenges regarding molecular property prediction using LLMs.

Table 1: Encoder-only scientific LLMs. The release date column displays the date of the first publication for each paper. When available, the publication date of the last updated version is displayed between parentheses. a:"Model Size" is reported as the number of parameters. b: The authors report they not used as many encoder layers as it was used in the original BERT paper. But the total number of parameters was not reported.

LLM	Model Size^a	Training Data	Architecture	Application	Release date
CatBERTa[202]	355M	OpenCatalyst2020 (OC20)	RoBERTa	Property prediction	2023.09 (2023.11)
SELFormer[203]	~86M	~2M compounds from ChEMBL	RoBERTa	Property prediction	2023.04 (2023.06)
ChemBERTa-2[122]	5M - 46M	77M SMILES from Pub-Chem	RoBERTa	Property prediction	2022.09
MaterialsBERT[204]	110M	2.4M material science abstracts + 750 annotated abstract for NER	BERT	NER and property extraction	2022.09 (2023.04)

<div align="right">Continued on next page</div>

Table 1 – continued from previous page

LLM	Model Size[a]	Training Data	Architecture	Application	Release date
SolvBERT[205]	b	1M SMILES of solute-solvent pairs from CombiSolv-QM and LogS from Boobier et al.[206]	BERT	Property prediction	2022.07 (2023.01)
ScholarBERT[207]	340M, 770M	Public.Resource.Org, Inc	BERT	Property prediction	2022.05 (2023.05)
BatteryBERT[208]	∼ 110M	∼ 400k papers from RSC, Elsevier and Springer	BERT	Document classification	2022.05
MatBERT[209]	110M	Abstracts from solid state articles and abstracts and methods from gold nanoparticle articles	BERT	NER	2022.04
MatSciBERT[210]	110M	∼150K material science paper downloaded from Elsevier	BERT	NER and text classification	2021.09 (2022.05)
Mol-BERT[118]	110M	∼4B SMILES from ZINC15 and ChEMBL27	BERT	Property prediction	2021.09
MolFormer[211]	b	PubChem and ZINC	BERT	Property prediction	2021.06 (2022.12)
ChemBERT[212]	110M	∼200k extracted using ChemDataExtractor	BERT	NER	2021.06
MolBERT[213]	∼85M	ChemBench	BERT	Property prediction	2020.11
ChemBERTa[44]		10M SMILES from PubChem	RoBERTa	Property prediction	2020.10
BioMegatron[214]	345M, 800M, 1.2B	Wikipedia, CC-Stories, Real-News, and OpenWebtext	Megatron-LM	NER and QA	2020-10
PubMedBERT[215]	110M	14M abstracts from PubMed	BERT	NER, QA, and document classification	2020.07 (2021.10)
Molecule Attention Transformer[216]	b	ZINC15	Encoder with GCN features	Property prediction	2020.02
SMILES-BERT[217]	b	∼18M SMILES from ZINC	BERT	Property prediction	2019.09
BlueBERT[218]	110M	PubMed and MIMIC-III	BERT	NER, and document classification	2019.06
ClinicalBERT[219]	110M	MIMIC-III	BERT	Patient readmission probability	2019.04
SciBERT[220]	110M	1.14M papers from Semantic Scholar	BERT	NER and sentence classification	2019.03 (2019.11)
BioBERT[221]	110M	PubMed and PMC	BERT	NER and QA	2019.01 (2019.09)

3.2.1 Property Prediction

The universal value of chemistry lies in identifying and understanding the properties of compounds to optimize their practical applications. In the pharmaceutical industry, therapeutic molecules interact with the body in profound

ways.[222–224] Understanding these interactions and modifying molecular structures to enhance those therapeutic benefits can lead to significant medical advancements.[225] Similarly, in polymer science, material properties depend on chemical structure, polymer chain length, and packing,[226] and a protein's function similarly depends on its structure and folding. Historically, chemists have identified new molecules from natural products[227] and screened them against potential targets[228] to test their properties for diseases. Once a natural product shows potential, chemists synthesize scaled-up quantities for further testing or derivatization,[229–231] a costly and labor-intensive process.[232,233] Traditionally, chemists have used their expertise to hypothesize the properties of new molecules derived from those natural products, hence aiming for the best investment of synthesis time and labor. Computational chemistry has evolved to support the chemical industry in more accurate property prediction.[234] Techniques such as quantum theoretical calculations and force-field-based molecular dynamics offer great support for property prediction and the investigation of molecular systems, though both require substantial computational resources.[235–239] Property prediction can now be enhanced through machine learning tools,[159,240–242] and more recent advancements in LLMs lead to effective property prediction without the extensive computational demands of quantum mechanics and MD calculations. Combined with human insight, AI can revolutionize material development, enabling the synthesis of new materials with a high likelihood of possessing desired properties for specific applications.

3.2.2 Encoder-only Mol-LLMs

Encoder-only models are exemplified by the BERT architecture, which is commonly applied in natural language sentiment analysis to extract deeper patterns from prose.[243] The human chemist has been taught to look at a 2D image of a molecular structure and to recognize its chemical properties or classify the compound. Therefore, encoder-only models would ideally convert SMILES strings, empty of inherent chemical essence, into a vector representation, or latent space, which would reflect those chemical properties. This vector representation can then be used directly for various downstream tasks.

While encoder-only LLMs are predominantly used for property prediction, they are also applicable for synthesis classification. Schwaller et al.[244] used a BERT model to more accurately classify complex synthesis reactions by generating reaction fingerprints from raw SMILES strings, without the need to separate reactants from reagents in the input data, thereby simplifying data preparation. The BERT model achieved higher accuracy (98.2%) compared to the encoder-decoder model (95.2%) for classifying reactions. Accurate classification aids in understanding reaction mechanisms, vital for reaction design, optimization, and retrosynthesis. Toniato et al.[245] also used a BERT architecture to classify reaction types for downstream retrosynthesis tasks that would enable the manufacture of any molecular target. Further examples of BERT use include self-supervised reaction atom-to-atom mapping.[246,247] These chemical classifications would accelerate research and development in organic synthesis, described further below.

Beyond synthesis classification, encoder-only models like BERT have shown great promise for molecular property prediction, especially when labeled data is limited. Recognizing this, Wang et al. introduced a semi-supervised SMILES-BERT model, which was pretrained on a large unlabeled dataset with a Masked SMILES Recovery task.[248] The model was then fine-tuned for various molecular property prediction tasks, outperforming state-of-the-art methods in 2019 on three chosen datasets varying in size and property. This marked a shift from using BERT for reaction classification towards property prediction and drug discovery. Maziarka et al.[216] also claimed state-of-the-art performance in property prediction after self-supervised pretraining in their Molecule Attention Transformer (MAT), which adapted BERT to chemical molecules by augmenting the self-attention with inter-atomic distances and molecular graph structure.

Zhang et al.[249] also tackled the issue of limited property-labeled data and the lack of correlation between any two datasets labeled for different properties, hindering generalizability. They introduced multitask learning BERT (MTL-BERT), which used large-scale pretraining and multitask learning with unlabeled SMILES strings from ChEMBL,[168] which is a widely-used database containing bioactive molecules with drug-like properties, designed to aid drug discovery. The MTL-BERT approach mined contextual information and extracted key patterns from complex SMILES strings, improving model interpretability. The model was fine-tuned for relevant downstream tasks, achieving better performance than state-of-the-art methods in 2022 on 60 molecular datasets from ADMETlab[250] and MoleculeNet.[56]

In 2021, Li and Jiang[118] introduced Mol-BERT, pretrained on four million unlabeled drug SMILES from the ZINC15[251] and ChEMBL27[168] databases to capture molecular substructure information for property prediction. Their work leveraged the underutilized potential of large unlabeled datasets like ZINC, which contains over 230 million commercially available compounds, and is designed for virtual screening and drug discovery. Mol-BERT consisted of three components: a PretrainingExtractor, Pretraining Mol-BERT, and Fine-Tuning Mol-BERT. It treated Morgan fingerprint fragments as "words" and complete molecular compounds as "sentences," using RDKit and the Morgan algorithm for canonicalization and substructure identification. This approach generated comprehensive molecular fingerprints from SMILES strings, used in a Masked Language Model (MLM) task for pretraining. Mol-BERT was fine-tuned on labeled samples, providing outputs as binary values or continuous scores for classification or regression, and it

outperformed existing sequence and graph-based methods by at least 2% in ROC-AUC scores on Tox21, SIDER, and ClinTox benchmark datasets.[56]

Ross et al.[252] introduced MoLFormer, a large-scale self-supervised BERT model, with the intention to provide molecular property predictions with competitive accuracy and speed when compared to Density Functional Theory calculations or wet-lab experiments. They trained MoLFormer with rotary positional embeddings on SMILES sequences of 1.1 billion unlabeled molecules from ZINC,[251] and PubChem,[166] another database of chemical properties and activities of millions of small molecules, widely used in drug discovery and chemical research. The rotary positional encoding captures token positions more effectively than traditional methods,[71] improving modeling of sequence relationships. MoLFormer outperformed state-of-the-art GNNs on several classification and regression tasks from ten MoleculeNet[56] datasets, while performing competitively on two others. It effectively learned spatial relationships between atoms, predicting various molecular properties, including quantum-chemical properties. Additionally, the authors stated how MoLFormer represents an efficient and environment-friendly use of computational resources, claiming a reduced GPU usage in training by a factor of 60 (16 GPUs instead of 1000).

With ChemBERTa, Chithrananda et al.[44] explored the impact of pretraining dataset size, tokenization strategy, and the use of SMILES or SELFIES, distinguishing their work from other BERT studies. They used HuggingFace's RoBERTa transformer,[253] and referenced a DeepChem[56] tutorial for accessibility. Their results showed improved performance on downstream tasks (BBBP, ClinTox, HIV, Tox21 from MoleculeNet[56]) as the pretraining dataset size increased from 100K to 10M. Although ChemBERTa did not surpass state-of-the-art GNN-based baselines like Chemprop (which used 2048-bit Morgan Fingerprints from RDKit),[254] the authors suggested that with expansion to larger datasets they would eventually beat those baselines. The authors compared Byte-Pair Encoder (BPE) with a custom SmilesTokenizer and its regular expression developed by[255] while exploring tokenization strategies. They found the SmilesTokenizer slightly outperformed BPE, suggesting more relevant sub-word tokenization is beneficial. No difference was found between SMILES and SELFIES, but the paper highlighted how attention heads in transformers could be visualized with BertViz,[256] showing certain neurons selective for functional groups. This study underscored the importance of appropriate benchmarking and addresses the carbon footprint of AI in molecular property prediction.

In ChemBERTa-2, Ahmad et al.[122] aimed to create a foundational model applicable across various tasks. They addressed a criticism that LLMs were not so generalizable because the training data was biased or non-representative. They addressed this criticism by training on 77M samples and adding a Multi-Task Regression component to the pretraining. ChemBERTa-2 matched state-of-the-art architectures on MoleculeNet.[56] As with ChemBERTa, the work was valuable because of additional exploration, in this case into how pretraining improvements affected certain downstream tasks more than others, depending on the type of fine-tuning task, the structural features of the molecules in the fine-tuning task data set, or the size of that fine-tuning dataset. The result was that pretraining the encoder-only model is important, but gains could be made by considering the chemical application itself, and the associated fine-tuning dataset.

In June 2023, Yuksel et al.[203] introduced SELFormer, building on ideas from ChemBERTa2[122] and using SELFIES for large data input. Yuksel et al.[203] argue that SMILES strings have validity and robustness issues, hindering effective chemical interpretation of the data, although this perspective is not universally held.[257] SELFormer uses SELFIES and is pretrained on two million drug-like compounds, fine-tuned for diverse molecular property prediction tasks (BBBP, SIDER, Tox21, HIV, BACE, FreeSolv, ESOL, PDBbind from MoleculeNet).[56] SELFormer outperformed all competing methods for some tasks and produced comparable results for the rest. It could also discriminate molecules with different structural properties. The paper suggests future directions in multimodal models combining structural data with other types of molecular information, including text-based annotations. We will discuss such multimodal models below.

In 2022, Yu et al.[205] published SolvBERT, a multi-task BERT-based regression model that could predict both solvation free energy and solubility from the SMILES notations of solute-solvent complexes. It was trained on the CombiSolv-QM dataset,[258] a curation of experimental solvent free energy data called CombiSolv-Exp-8780,[259–262] and the solubility dataset from Boobier et al.[206]. SolvBERT's performance was benchmarked against advanced graph-based models[263,264] This work is powerful because there is an expectation that solvation free energy depends on 3-dimensional conformational properties of the molecules, or at least 2D properties that would be well characterized by graph-based molecular representations. It shows an overachieving utility of using SMILES strings in property prediction, and aligns with other work by Winter et al.[265], regarding activity coefficients. SolvBERT showed comparable performance to a Directed Message Passing Neural Network (DMPNN) in predicting solvation free energy, largely due to its effective clustering feature in the pretraining phase as shown by TMAP (Tree Map of All Possible) visualizations. Furthermore, SolvBERT outperformed Graph Representation Of Molecular Data with Self-supervision (GROVER)[264] in predicting experimentally evaluated solubility data for new solute-solvent combinations. This underscores the significance of SolvBERT's ability to capture the dynamic and spatial complexities of solvation interactions in a text-based model.

While models like SolvBERT have achieved impressive results in solvation free energy prediction, challenges such as limited labeled data continue to restrict the broader application of transformer models in chemistry. Recognizing

this issue, Jiang et al. introduced INTransformer in 2024,[266] a method designed to enhance property prediction by capturing global molecular information more effectively, even when data is scarce. By incorporating perturbing noise and using contrastive learning to artificially augment smaller datasets, INTransformer delivered improved performance on several tasks. Ongoing work continues to explore various transformer strategies for smaller datasets. Again using contrastive learning, which maximizes the difference between representations of similar and dissimilar data points, but in a different context, MoleculeSTM[267] uses LLM encoders to create representations for SMILES and for descriptions of molecules extracted from PubChem.[268] Similar work was performed by Xu et al.[269] The authors curated a dataset with descriptions of proteins. Subsequently, to train ProtST, a protein language model (PLM) was used to encode amino acid sequences and LLMs to encode the descriptions.

In this section, we outlined the advancements of encoder-only models like BERT and their evolution for property prediction and synthesis classification. Chemists traditionally hypothesize molecular properties, but these models, ranging from Mol-BERT to SolvBERT, showcase the growing efficiency of machine learning in property prediction. Approaches such as multitask learning and contrastive learning, as seen in INTransformer, offer solutions to challenges posed by limited labeled data.

3.3   Property Directed Inverse Design and Decoder-only mol-LLMs

Decoder-only GPT-like architectures offer significant value for property-directed molecule generation and *de novo* chemistry applications because they excel at generating novel molecular structures by learning from vast datasets of chemical compounds. These models can capture intricate patterns and relationships within molecular sequences, proposing viable new compounds that adhere to desired chemical properties and constraints. This enables rapid exploration and innovation within an almost infinite chemical space. Moreover, such large general-purpose models can be fine-tuned with small amounts of domain-specific scientific data,[142,270] allowing them to support specific applications efficiently. In this section, we first describe property-directed inverse design from a chemistry perspective and then examine how decoder-only LLMs have propelled inverse design forward.

Table 2: Decoder-only scientific LLMs. The release date column displays the date of the first publication for each paper. When available, the publication date of the last updated version is displayed between parentheses. a:"Model Size" is reported as the number of parameters. "PubMed" refer to the PubMed abstracts dataset, while PMC (PubMed Corpus) refers to the full-text corpus dataset. b: The total number of parameters was not reported.

LLM	Model Size^a	Training Data	Architecture	Application	Release date
Tx-LLM[271]	b	TDC datasets	PaLM-2	Property prediction and retrosynthesis	2024.06
BioMedLM[272]	2.7B	PubMed abstracts and full articles	GPT	QA	2024.03
LlasMol[273]	$\sim$ 7B	SMolInstruct	Galactica, LLaMa, Mistral	Property prediction, molecule captioning, molecule generation, retrosynthesis, name conversion	2024.02 (2024.08)
BioMistral[274]	7B	PubMed Central (PMC)	Mistral	QA	2024.02 (2024.08)
BiMediX[275]	8x7B	1.3M Arabic-English instructions (BiMed)	Mixtral	QA	2024.02
EpilepsyLLM[276]	7B	Data from the Japan Epilepsy Association, Epilepsy Information Center, and Tenkan Net	LLaMa	QA	2024.01

Table 2 – continued from previous page

LLM	Model Size[a]	Training Data	Architecture	Application	Release date
CheXagent[277]	7B	28 publicly available datasets, including PMC, MIMIC, wikipedia, PadChest, and BIMCV-COVID-19	Mistral	QA, Image understanding	2024.01
ChemSpaceAL[278]	b	ChEMBL 33, GuacaMol v1, MOSES, and BindingDB 08-2023	GPT	Molecule Generation	2023.09 (2024.02)
BioMedGPT-LM[279]	7B and 10B	5.5M biomedical papers from S2ORC	LLaMA2	QA	2023.08
Darwin[280]	7B	SciQ and Web of Science	LLaMA	QA, Property prediction, NER, and Molecule Generation	2023.08
cMolGPT[46]	b	MOSES	GPT	Molecule Generation	2023.05
PMC-LLaMA[281]	7B and 13B	MedC-k and MedC-I	LLaMA	QA	2023.04 (2024.04)
GPTChem[142]	175B	Curation of multiple classification and regression benchmarks	GPT-3	Property prediction and inverse design	2023.02 (2024.02)
Galactica[123]	125M, 1.3B, 6.7B, 30B, 120B	The galactica corpus, a curation with 62B scientific documents	Decoder-only	QA, NER, Document Summarization, Property Prediction	2022.11
BioGPT[282]	355M	15M of Title and abstract from PubMed	GPT-2	QA, NER, and Document Classification	2022-09 (2023.04)
SMILES-to-properties-transformer[265]	6.5M	Synthetic data generated with the thermodynamic model COSMO-RS	GPT-3	Property prediction	2022.06 (2022.09)
ChemGPT[283]	∼ 1B	10M molecules from PubChem	GPT-neo	Molecule generation	2022.05 (2023.11)
Regression Transformer[139]	∼27M	ChEMBL, MoleculeNet, USPTO, etc	XLNet	Property prediction, Molecule tuning, Molecule generation	2022.02 (2023.04)
MolGPT[284]	6M	MOSES and GuacaMol	GPT	Molecule Generation	2021.10
Adilov2021[285]	13.4M	5M SMILES from ChemBERTa's PubChem-10M.	GPT-2	Property prediction and molecule generation	2021.09

3.3.1 Property Directed Inverse Design

Nature has long been a rich source of molecules that inhibit disease proliferation, because organisms have evolved chemicals for self-defense. Historically, most pharmaceuticals are derived from these natural products,[286,287] which offer benefits such as cell permeability, target specificity, and a vast chemical diversity.[288] However, the high costs and complexities associated with high-throughput screening and synthesizing natural products limit the exploration of this space.[286,288]

While natural products have been a valuable starting point, we are not confined to their derivatives. AI, particularly generative LLMs, allows us to go beyond nature and explore a much larger chemical space. *In-silico* molecular design enables rapid modification, akin to random mutation,[289] where only valid, synthesizable molecules that meet predefined property criteria remain in the generated set.[242,290] This approach allows us to test modifications *in-silico*, expanding exploration beyond the boundaries of natural products.

The true innovation of AI-driven molecular design, however, lies in its ability to directly generate candidate molecules based on desired properties, without the need for iterative stepwise modifications.[291] This "inverse design" capability allows us to start with a target property and directly generate candidate molecules that meet the predefined property requirements. Generative LLMs applied to sequences of atoms and functional groups offer a powerful opportunity for out-of-the-box exploration, tapping into the vast chemical space that extends far beyond the confines of nature. This accelerates the path from concept to viable therapeutic agents, aligning seamlessly with decoder-only LLM architectures.

3.3.2 Decoder-only Mol-LLMs

One of the first applications of decoder-only models in chemistry was Adilov's (2021) "Generative pretraining from Molecules".[285] This work pretrained a GPT-2-like causal transformer for self-supervised learning using SMILES strings. By introducing "adapters" between attention blocks for task-specific fine-tuning,[292] this method provided a versatile approach for both molecule generation and property prediction, requiring minimal architectural changes. It aimed to surpass encoder-only models, such as ChemBERTa,[44] with a more scalable and resource-efficient approach, demonstrating the power of decoder-only models in chemical generation.

A key advancement then came with MolGPT,[284] a 6-million-parameter decoder-only model designed for molecular generation. MolGPT introduced masked self-attention, enabling the learning of long-range dependencies in SMILES strings. The model ensured chemically valid SMILES representations, respecting structural rules like valency and ring closures. It also utilized salience measures for interpretability, aiding in predicting SMILES tokens and understanding which parts of the molecule were most influential in the model's predictions. MolGPT outperformed many existing Variational Auto-Encoder (VAE)-based approaches,[293–300] in predicting novel molecules with specified properties, being trained on datasets like MOSES[301] and GuacaMol.[302]

While MolGPT's computational demands may be higher than traditional VAEs, its ability to generate high-quality, novel molecules justifies this trade-off. MolGPT demonstrated strong performance on key metrics such as validity, which measures the percentage of generated molecules that are chemically valid according to bonding rules; uniqueness, the proportion of generated molecules that are distinct from one another; Frechet ChemNet Distance (FCD),[303] which compares the distribution of generated molecules to that of real molecules in the training set, indicating how closely the generated molecules resemble real-world compounds; and KL divergence,[302] a measure of how the probability distribution of generated molecules deviates from the true distribution of the training data. These metrics illustrate MolGPT's ability to generate high-quality, novel molecules while maintaining a balance between diversity and similarity to known chemical spaces. A brief summary of advancements in transformer-based models for *de-novo* molecule generation from 2023 and 2024 follows, which continue to refine and expand upon the foundational work laid by models like MolGPT.

Haroon et al.[304] further developed a GPT-based model with relative attention for *de novo* drug design, showing improved validity, uniqueness, and novelty. This work was followed by Frey et al.[283], who introduced ChemGPT to explore hyperparameter tuning and dataset scaling in new domains. ChemGPT's contribution lies in refining generative models to better fit specific chemical domains, advancing the understanding of how data scale impacts generative performance. Both Wang et al.[305] and Mao et al.[306] presented work that surpassed MolGPT. Furthermore, Mao et al.[140] showed that decoder-only models could generate novel compounds using IUPAC names directly.

This marked a departure from typical SMILES-based molecular representations, as IUPAC names offer a standardized, human-readable format that aligns with how chemists conceptualize molecular structures. By integrating these chemical semantics into the model, iupacGPT[140] bridges the gap between computational predictions and real-world chemical applications. The IUPAC name outputs are easier to understand, validate, and apply, facilitating smoother integration into workflows like regulatory filings, chemical databases, and drug design. Focusing on pretraining with a vast dataset of IUPAC names and fine-tuning with lightweight networks, iupacGPT excels in molecule generation, classification, and regression tasks, providing an intuitive interface for chemists in both drug discovery and material science.

In a similar vein, Zhang et al.[307] proposed including target 3D structural information in molecular generative models, even though their approach is not LLM-based. However, it serves as a noteworthy contribution to the field of structure-based drug design. Integrating biological data, such as 3D protein structures, can significantly improve the relevance and specificity of generated molecules, making this method valuable for future LLM-based drug design. Similarly, Wang et al.[308] discussed PETrans, a deep learning method that generates target-specific ligands using protein-specific encoding and transfer learning. This study further emphasizes the importance of using transformer models for generating molecules with high binding affinity to specific protein targets. The significance of these works lies in their demonstration that integrating both human-readable formats (like IUPAC names) and biological context (such as protein structures) into generative models can lead to more relevant, interpretable, and target-specific drug

candidates. This reflects a broader trend in AI-driven chemistry to combine multiple data sources for more precise molecular generation, accelerating the drug discovery process.

In 2024, Yoshikai et al.[309] discussed the limitations of transformer architectures in recognizing chirality from SMILES representations, which impacts the prediction accuracy of molecular properties. To address this, they coupled a transformer with a VAE. Using contrastive learning from NLP to generate new molecules with multiple SMILES representations, enhancing molecular novelty and validity.Kyro et al.[278] presented ChemSpaceAL, an active learning method for protein-specific molecular generation, efficiently identifying molecules with desired characteristics without prior knowledge of inhibitors. Yan et al.[310] proposed the GMIA framework, which improves prediction accuracy and interpretability in drug-drug interactions through a graph mutual interaction attention decoder. These innovations represent significant strides in addressing key challenges in molecular generation, such as chirality recognition, molecular novelty, and drug-drug interaction prediction. By integrating new techniques like VAEs, contrastive learning, and active learning into transformer-based models, they have improved both the accuracy and interpretability of molecular design.

Building on these developments, Shen et al.[311] reported on AutoMolDesigner, an open-source tool for small-molecule antibiotic design, further emphasizing the role of automation in molecular generation. This work serves as a precursor to more complex models, such as Taiga[101] and cMolGPT,[46] which employ advanced methods like autoregressive mechanisms and reinforcement learning for molecular generation and property optimization.

For a deeper dive into decoder-only transformer architecture in chemistry, we highlight the May 2023 "Taiga" model by Mazuz et al.[101], and cMolGPT by Wang et al.[46]. Taiga first learns to map SMILES strings to a vector space, and then refines that space using a smaller, labeled dataset to generate molecules with targeted attributes. It uses an autoregressive mechanism, predicting each SMILES character in sequence based on the preceding ones. For property optimization, Taiga employs the REINFORCE algorithm,[106] which helps refine molecules to enhance specific features. While this reinforcement learning (RL) approach may slightly reduce molecular validity, it significantly improves the practical applicability of the generated compounds. Initially evaluated using the Quantitative Estimate of Drug-likeness (QED) metric,[312] Taiga has also demonstrated promising results in targeting IC50 values,[168] the BACE protein,[313] and anti-cancer activities they collected from a variety of sources. This work underscores the importance of using new models to address applications that require a higher level of chemical sophistication, to illustrate how such models could ultimately be applied outside of the available benchmark datasets. It also builds on the necessary use of standardized datasets and train-validation-test splitting, to demonstrate progress, as explained by Wu et al.[56]. Yet, even the MoleculeNet benchmarks[56] are flawed, and we point the reader here to a more detailed discussion on benchmarking,[188] given that a significant portion of molecules in the BACE dataset have undefined stereo centers, which, at a deeper level, complicates the modeling and prediction accuracy.

While models like Taiga demonstrate the power of autoregressive learning and reinforcement strategies to generate molecules with optimized properties, the next step in molecular design incorporates deeper chemical domain knowledge. This approach is exemplified by Wang et al.[46]. They introduced cMolGPT, a conditional generative model that brings a more targeted focus to drug discovery by integrating specific protein-ligand interactions, which underscores the importance of incorporating chemical domain knowledge to effectively navigate the vast landscape of drug-like molecules. Using self-supervised learning and an auto-regressive approach, cMolGPT generates SMILES guided by predefined conditions based on target proteins and binding molecules. Initially trained on the MOSES dataset[301] without target information, the model is fine-tuned with embeddings of protein-binder pairs, focusing on generating compound libraries and target-specific molecules for the EGFR, HTR1A, and S1PR1 protein datasets.[314–317]

Their approach employs a QSAR model[5] to predict the activity of generated compounds, achieving a Pearson correlation coefficient over 0.75. However, despite the strong predictive capabilities, this reliance on a QSAR model, with its own inherent limitations, highlights the need for more extensive experimental datasets. cMolGPT[46] tends to generate molecules within the sub-chemical space represented in the original dataset, successfully identifying potential binders but struggling to broadly explore the chemical space for novel solutions. This underscores the challenge of generating diverse molecules with varying structural characteristics while maintaining high binding affinity to specific targets. While cMolGPT advances the integration of biological data and fine-tuned embeddings for more precise molecular generation, models like Taiga and cMolGPT differ in their approach. Taiga[101] employs reinforcement learning to optimize generative models for molecule generation, while cMolGPT uses target-specific embeddings to guide the design process. Both highlight the strengths of decoder-only models but emphasize distinct strategies; Taiga optimizes molecular properties through autoregressive learning, and cMolGPT focuses on conditional generation based on protein-ligand interactions.

In contrast, Yu et al.[273] follow a different approach with LlaSMol,[273] which utilizes pretrained models (for instance Galactica, LlaMa2, and Mistral) and performs parameter efficient fine-tuning (PEFT) techniques[318,319] such as LoRa.[320] PEFT enables fine-tuning large language models with fewer parameters, making the process more resource-efficient

while maintaining high performance. LlaSMol demonstrated its potential by achieving state-of-the-art performance in property prediction tasks, particularly when fine-tuned on benchmark datasets like MoleculeNet.[56]

There continue to be significant advancements being made in using transformer-based models to tackle chemical prediction tasks with optimized computational resources, including more generalist models, such as Tx-LLM,[271] designed to streamline the complex process of drug discovery. For additional insights on how these models are shaping the field, we refer the reader to several excellent reviews,[164,321–323] with Goel et al.[324] highlighting the efficiency of modern machine learning methods in sampling drug-like chemical space for virtual screening and molecular design. Goel et al.[324] discussed the effectiveness of generative models, including large language models (LLMs), in approximating the vast chemical space, particularly when conditioned on specific properties or receptor structures.

We provide a segue from this section by introducing the work by Jablonka et al.[142], which showcases a decoder-only GPT model that, despite its training on natural language rather than specialized chemical languages, competes effectively with decoder-only LLMs tailored to chemical languages. The authors finetuned GPT-3 to predict properties and conditionally generate molecules and, therefore, highlight its potential as a foundational tool in the field. This work sets the stage for integrating natural language decoder-only LLMs, like GPT, into chemical research, where they could serve as central hubs for knowledge discovery.

Looking ahead, this integration foreshadows future developments that pair LLMs with specialized tools to enhance their capabilities, paving the way for the creation of autonomous agents that leverage deep language understanding in scientific domains. Decoder-only models have already significantly advanced inverse molecular design, from improving property prediction to enabling target-specific molecular generation. Their adaptability to various chemical tasks demonstrates their value in optimizing drug discovery processes and beyond. As models like LlaSMol and cMolGPT continue to evolve, integrating chemical domain knowledge and biological data, they offer exciting opportunities for more precise molecular generation. The growing potential for combining large language models like GPT-4 with specialized chemical tools signals a future where AI-driven autonomous agents could revolutionize chemical research, making these models indispensable to scientific discovery.

3.4 Synthesis Prediction and Encoder-decoder Mol-LLMs

The encoder-decoder architecture is designed for tasks involving the translation of one sequence into another, making it ideal for predicting chemical reaction outcomes or generating synthesis pathways from given reactants. We begin with a background on optimal synthesis prediction and describe how earlier machine learning has approached this challenge. Following that, we explain how LLMs have enhanced chemical synthesis prediction and optimization. Although, our context below is aptly chosen to be synthesis prediction, other applications exist. For example, SMILES Transformer (ST)[325] is worth a mention, historically, because it explored the benefits of self-supervised pretraining to produce continuous, data-driven molecular fingerprints from large SMILES-based datasets.

3.4.1 Synthesis Prediction

Once a molecule has been identified through property-directed inverse design, the next challenge is to predict its optimal synthesis, including yield. Shenvi[332] describe how the demanding and elegant syntheses of natural products has contributed greatly to organic chemistry. However, in the past 20 years, the focus has shifted away from complex natural product synthesis towards developing new reactions applicable for a broader range of compounds, especially in reaction catalysis.[332] Yet, complex synthesis is becoming relevant again as it can be digitally encoded, mined by LLMs,[333] and applied to new challenges. Unlike property prediction, reaction prediction is particularly challenging due to the involvement of multiple molecules. Modifying one reactant requires adjusting all others, with different synthesis mechanisms or conditions likely involved. Higher-level challenges exist for catalytic reactions and complex natural product synthesis. Synthesis can be approached in two ways. Forward synthesis involves building complex target molecules from simple, readily available substances, planning the steps progressively. Retrosynthesis, introduced by E.J. Corey in 1988,[334] is more common. It involves working backward from the target molecule, breaking it into smaller fragments whose re-connection is most effective. Chemists choose small, inexpensive, and readily available starting materials to achieve the greatest yield and cost-effectiveness. As a broad illustration, the first total synthesis of discodermolide[335] involved 36 such steps, a 24-step longest linear sequence, and a 3.2% yield. There are many possible combinations for the total synthesis of the target molecule, and the synthetic chemist must choose the most sensible approach based on their expertise and knowledge. However, this approach to total synthesis takes many years. LLMs can now transform synthesis such that structure-activity relationship predictions can be coupled in lock-step with molecule selection based on easier synthetic routes. This third challenge of predicting the optimal synthesis can also lead to the creation of innovative, non-natural compounds, chosen because of such an easier predicted synthesis but for which the properties are still predicted to meet the needs of the application. Thus, these three challenges introduced above are interconnected.

Table 3: Encoder-decoder scientific LLMs. The release date column displays the date of the first publication for each paper. When available, the publication date of the last updated version is displayed between parentheses. a:"Model Size" is reported as the number of parameters. b: The total number of parameters was not reported.

LLM	Model Size[a]	Training Data	Architecture	Application	Release date
BioT5+[117]	252M	ZINC20, UniRef50, 33M PubMed articles, 339K mol-text pairs from PubChem, 569K FASTA-text pairs from Swiss-prot	T5	Molecule Captioning, Molecule Generation, Property Prediction,	2024.02 (2024.08)
nach0[187]	250M	MoleculeNet, USPTO, ZINC	T5	Property prediction, Molecule generation, Question answering, NER	2023.11 (2024.05)
ReactionT5[326]	220M	ZINC and ORD	T5	Property prediction and Reaction prediction	2023.11
BioT5[116]	252M	ZINC20, UniRef50, full-articles from BioRxiv and PubMed, mol-text-IUPAC information from PubChem	T5	Molecule Captioning, Property Prediction	2023-10 (2024.12)
MOLGEN[327]	b	ZINC15	BART	Molecule Generation	2023.01 (2024.03)
Text+Chem T5[328]	60M, 220M	11.5M or 33.5M samples curated from Vaucher et al.[329], Toniato et al.[245], and CheBI-20	T5	Molecule Captioning, Product Prediction, Retrosynthesis, Molecule Generation	2023.01 (2023.06)
MolT5[330]	60M, 770M	C4 dataset	T5	Molecule Captioning and Molecule Generation	2022.04 (2022.12)
T5Chem[179]	220M	USPTO	T5	Product Prediction, Retrosynthesis, Property Prediction	2022.03
Text2Mol[331]	b	CheBI-20	SciBERT w/ decoder	Molecule captioning and conditional molecule generation	2021.11
ChemFormer[185]	45M, 230M	100M SMILES from ZINC-15	BART	Product Prediction, Property Prediction, Molecular Generation	2021.07 (2022.01)
SMILES transformer[325]	b	ChEMBL24	Transformer	Property prediction	2019.11
Molecular Transformer[255]	12M	USPTO	Transformer	Product prediction	2018.11 (2019.08)

3.4.2 Encoder-decoder mol-LLMs

Before we focus on transformer use, some description is provided on the evolution from RNN and Gated Recurrent Unit (GRU) approaches in concert with the move from template-based to semi-template-based to template-free models. Nam and Kim[336] pioneered forward synthesis prediction using a GRU-based translation model. In contrast, Liu et al.[337] reported retro-synthesis prediction with a Long Short-Term Memory (LSTM) based seq2seq model incorporating an attention mechanism, achieving 37.4% accuracy on the USPTO-50K dataset. The reported accuracies of these early models highlighted the challenges of synthesis prediction, particularly retrosynthesis. Schneider et al.[338] further advanced retrosynthesis by assigning reaction roles to reagents and reactants based on the product.

Building on RNNs and GRUs, the field advanced with the introduction of template-based models. In parallel with the development of the Chematica tool[339,340] for synthesis mapping, Segler and Waller[341] highlighted that traditional rule-based systems often failed by neglecting molecular context, leading to "reactivity conflicts." Their approach emphasized transformation rules that capture atomic and bond changes, applied in reverse for retrosynthesis. Trained on 3.5 million reactions, their model achieved 95% top-10 accuracy in retrosynthesis and 97% for reaction prediction on a validation set of nearly 1 million reactions from the Reaxys database (1771–2015). Although not transformer-based, this work laid the foundation for large language models (LLMs) in synthesis. However, template-based models depend on

explicit reaction templates from known reactions, limiting their ability to predict novel reactions and requiring manual updates to incorporate new data.

Semi-template-based models offered a balance between rigid template-based methods and flexible template-free approaches. They used interpolation or extrapolation within template-defined spaces to predict a wider range of reactions and to adjust based on new data. In 2021, Somnath et al.[342] introduced a graph-based approach recognizing that precursor molecule topology is largely unchanged during reactions. Their model broke the product molecule into "synthons" and added relevant leaving groups, making results more interpretable.[343] Training on the USPTO-50k dataset,[338] they achieved a top-1 accuracy of 53.7%, outperforming previous methods.

However, the template-free approaches align well with transformer-based learning approaches because they learn retrosynthetic rules from raw training data. This provides significant flexibility and generalizability across various types of chemistry. Template-free models are not constrained by template libraries and so can uncover novel synthetic routes that are undocumented or not obvious from existing reaction templates. To pave the way for transformer use in synthesis, Cadeddu et al.[344] drew an analogy between fragments in a compound and words in a sentence due to their similar rank distributions. Schwaller et al.[345] further advanced this with an LSTM network augmented by an attention-mechanism-based encoder-decoder architecture, using the USPTO dataset.[338] They introduced a new "regular expression" (or regex) for tokenizing molecules, framing synthesis (or retrosynthesis) predictions as translation problems with a data-driven, template-free sequence-to-sequence model. They tracked which starting materials were actual reactants, distinguishing them from other reagents like solvents or catalysts, and used the regular expression to uniquely tokenize recurring reagents, as their atoms were not mapped to products in the core reaction. This regex for tokenizing molecules is commonly used today in all mol-based LLMs.

In 2019, going beyond the "neural machine" work of Nam and Kim[336], Schwaller et al.[255] first applied a transformer for synthesis prediction, framing the task as translating reactants and reagents into the final product. Their model inferred correlations between chemical motifs in reactants, reagents, and products in the dataset (USPTO-MIT,[346] USPTO-LEF,[347] USPTO-STEREO[345]). It required no handcrafted rules and accurately predicted subtle chemical transformations, outperforming all prior algorithms on a common benchmark dataset. The model handled inputs without a reactant-reagent split, following their previous work,[345] and accounted for stereochemistry, making it valuable for universal application. Then, in 2020, for automated retrosynthesis, Schwaller et al.[348] developed an advanced Molecular Transformer model with a hyper-graph exploration strategy. The model set a standard for predicting reactants and other entities, evaluated using four new metrics. "Coverage" measured how comprehensively the model could predict across the chemical space, while "class diversity" assessed the variety of chemical types the model could generate, ensuring it was not limited to narrow subsets of reactions. "Round-trip accuracy" checked whether the retrosynthetically predicted reactants could regenerate the original products, ensuring consistency in both directions. "Jensen–Shannon divergence" compared the predicted outcomes to actual real-world distributions, indicating how closely the model's predictions matched reality. Constructed dynamically, the hypergraph allowed for efficient expansion based on Bayesian-like probability scores, showing high performance despite training data limitations. Notably, accuracy improved when the re-synthesis of the target product from the generated precursors was factored in, a concept also employed by Chen and Jung[349] and Westerlund et al.[350]. Also in 2020, Zheng et al.[351] developed a "template-free self-corrected retrosynthesis predictor" (SCROP) using transformer networks and a neural network-based syntax corrector, achieving 59.0% accuracy on a benchmark dataset.[338,352] This approach outperformed other deep learning methods by over 2% and template-based methods by over 6%.

We now highlight advancements in synthesis prediction using the BART Encoder-Decoder architecture, starting with Chemformer by Irwin et al.[185]. This paper emphasized the computational expense of training transformers on SMILES and the importance of pretraining for efficiency. It showed that models pretrained on task-specific datasets or using only the encoder stack were limited for sequence-to-sequence tasks. After transfer learning, Chemformer achieved state-of-the-art results in both sequence-to-sequence synthesis tasks and discriminative tasks, such as optimizing molecular structures for specific properties. They studied the effects of small changes on molecular properties using pairs of molecules from the ChEMBL database[168] with a single structural modification. Chemformer's performance was tested on the ESOL, Lipophilicity, and Free Solvation datasets.[56] Irwin et al.[185] also described their use of an in-house property prediction model, but when models train on calculated data for ease of access and uniformity, they abstract away from real-world chemical properties. We again emphasize the importance of incorporating experimentally derived data into Chemistry LLM research to create more robust and relevant models. Continuously curating new, relevant datasets that better represent real-world chemical complexities will enhance the applicability and transferability of these models.

In 2023, Toniato et al.[245] also applied LLMs to single-step retrosynthesis as a translation problem, but increased retrosynthesis prediction diversity by adding classification tokens, or "prompt tokens," to the target molecule's language representation, guiding the model towards different disconnection strategies. Increased prediction diversity has

high value by providing out-of-the-box synthetic strategies to complement the human chemist's work. To measure retrosynthesis accuracy, Li et al.[353] introduced Retro-BLEU, a metric adapted from the BLEU (Bilingual Evaluation Understudy) score used in machine translation.[354] Despite progress in computer-assisted synthesis planning (CASP), not all generated routes are chemically feasible due to steps like protection and deprotection needed for product formation. Widely accepted NLP metrics like BLEU[354] and ROUGE[355] focus on precision and recall by computing n-gram overlaps between generated and reference texts. Similarly, in retrosynthesis, reactant-product pairs can be treated as overlapping bigrams. Retro-BLEU uses a modified BLEU score, emphasizing precision over recall, as there is no absolute best route for retrosynthesis. Although not yet applied to LLM-based predictions, this approach has value by allowing future performance comparison with a single standard.

Finally, by expanding the use of encoder-decoder architectures outside synthesis prediction into molecular generation, Fang et al.[327] introduced MOLGEN, a BART-based pretrained molecular language model, in a 2023 preprint updated in 2024. MOLGEN addressed three key challenges: generating valid SMILES strings, avoiding an observed bias that existed against natural product-like molecules, and preventing hallucinations of molecules that didn't retain the intended properties. Pretrained on 100 million molecules using SELFIES[162] and a masked language model approach, MOLGEN predicts missing tokens to internalize chemical grammar. An additional highlight of this work is how MOLGEN uses "domain-agnostic molecular prefix tuning." This technique integrates domain knowledge directly into the model's attention mechanisms by adding molecule-specific prefixes, trained simultaneously with the main model across various molecular domains. The model's parameters would thus be adjusted to better capture the complexities and diversities of molecular structures, and domain-specific insights would be seamlessly integrated. To prevent molecular hallucinations, MOLGEN employs a chemical feedback mechanism, to autonomously evaluate generated molecules for appropriate properties, to guide learning and optimization. Such feedback foreshadows a core aspect of autonomous agents, which is their capacity for reflection. We will explore this further below.

The advancements in synthesis prediction and molecular generation using encoder-decoder architectures have revolutionized the field, moving from rigid, template-based models to more flexible, template-free approaches. Early work with LSTMs and GRUs laid the foundation, while transformer-based models like Molecular Transformer and Chemformer set new benchmarks in accuracy and versatility. New metrics, such as Retro-BLEU, and domain-aware techniques, like MOLGEN's prefix tuning, have further refined predictions and molecular design. These innovations, coupled with self-correcting mechanisms, point to a future of autonomous molecular design, where AI agents can predict, evaluate, and optimize synthetic pathways and molecular properties, accelerating chemical discovery.

3.5   Multi-Modal LLMs

We have demonstrated the impact of LLMs on chemistry through their ability to process textual representations of molecules and reactions. However, LLMs can also handle diverse input modalities, representing molecular and chemical data in various formats.[356–358] In chemistry, data can be represented in various forms, each providing unique insights and information (see Section 3.1). Chemical representations can be broadly classified into 1D, 2D, and 3D categories, depending on how much structural detail they convey.[148,149] 1D representations include basic numerical descriptors, such as molecular features and fingerprints, as well as textual representations like SMILES,[43] SELFIES,[162] and IUPAC names. These descriptors vary in the amount of chemical information they carry.[359] 2D representations involve graph-based structures and visual formats, which can be extended with geometric information to produce 3D representations. Examples of 3D representations include molecular graphs enriched with spatial data, molecular point clouds, molecular grids, and 3D geometry files.[360]

Some of these representations can be input into models in different ways. For instance, a point cloud can be expressed either as a vector of coordinates (numerical input) or as a text-based PDB file. However, due to the distinct nature of the information conveyed, we treat textual descriptions of different molecular representations as separate modalities, even though both are technically strings. Additionally, molecule images have been utilized to train transformer-based models.[361] However, spectral data—such as Nuclear Magnetic Resonance (NMR), Infrared (IR) spectroscopy, and mass spectrometry, remain underexplored as inputs for LLM-based applications.

Multi-modal LLMs leverage and integrate these diverse data types to enhance their predictive and analytical capabilities. This integration improves the accuracy of molecular property predictions and facilitates the generation of novel compounds with desired properties. A key example is Text2Mol proposed by Edwards et al.[331] in 2021, which integrates natural language descriptions with molecular representations, addressing the cross-lingual challenges of retrieving molecules using text queries. The researchers created a paired dataset linking molecules with corresponding text descriptions and developed a unified semantic embedding space to facilitate efficient retrieval across both modalities. This was further enhanced with a cross-modal attention-based model for explainability and reranking. One stated aim was to improve retrieval metrics, which would further advance the ability for machines to learn from chemical literature.

In their 2022 follow-up, MolT5, Edwards et al.[330] expanded on earlier work by utilizing both SMILES string representations and textual descriptions to address two tasks: generating molecular captions from SMILES and predicting molecular structures from textual descriptions of desired properties. However, several key challenges remain. Molecules can be described from various perspectives, such as their therapeutic effects, applications (e.g., aspirin for pain relief or heart attack prevention), chemical structure (an ester and a carboxylic acid connected to a benzene ring in ortho geometry), or degradation pathways (e.g., breaking down into salicylic acid and ethanoic acid in moisture).[362] This complexity demands expertise across different chemistry domains, unlike typical image captioning tasks involving everyday objects (e.g., cats and dogs), which require minimal specialized knowledge. Consequently, building large, high-quality datasets pairing chemical representations with textual descriptions is a challenging task.

Moreover, standard metrics like BLEU, effective in traditional NLP, are insufficient for evaluating molecule-text tasks. To address these challenges, Edwards et al.[330] employed a denoising objective, training the model to reconstruct corrupted input data, thereby learning the structure of both text and molecules. Fine-tuning on gold-standard annotations further improved the model's performance, enhancing previous Text2Mol metrics[331] and enabling MolT5 to generate accurate molecular structures and their corresponding captions.

Other multimodal approaches similarly target the fusion of chemical and linguistic data to advance applications in molecular design. Seidl et al.[363] developed CLAMP, which combines separate chemical and language modules to predict biochemical activity, while Xu et al.[364] presented BioTranslator, a tool that translates text descriptions into non-text biological data to explore novel cell types, protein function, and drug targets. These examples highlight the growing trend of using language-based interfaces to enhance molecular exploration. The potential of multimodal LLMs extends beyond chemistry into more interactive and accessible tools. ChatDrug, by Liu et al.[365], integrates multimodal capabilities through a prompt module, a retrieval and domain feedback module, and a conversation module for systematic drug editing. It identifies and manipulates molecular structures for better interpretability in pharmaceutical research. Similarly, Christofidellis et al.[328] introduced a multi-domain, multi-task language model capable of handling tasks across both chemical and natural language domains without requiring task-specific pretraining.describe Joint Multi-domain Pre-training (JMP), which operates on the hypothesis that pre-training across diverse chemical domains enhances generalization toward a more robust foundational model. In this context, Liu et al.[366] developed MolXPT, introduced MolXPT, which further demonstrated the power of multimodal learning by achieving robust zero-shot molecular generation.

Finally, models that integrate even more diverse data types, such as GIT-Mol,[367] which combines graphs, images, and text, and MolTC,[368] which integrates graphical information for molecular interaction predictions illustrate how multimodal data improves accuracy and generalizability. Moreover, multimodal fusion models like PremuNet[369] and 3M-Diffusion, Zhu et al.[370] which use molecular graphs and natural language for molecule generation, represent a significant leap forward in the creation of novel compounds. Gao et al.[371] advanced targeted molecule generation with DockingGA, combining transformer neural networks with genetic algorithms and docking simulations for optimal molecule generation, utilizing Self-referencing Chemical Structure Strings to represent and optimize molecules. Zhou et al.[372] developed TSMMG, a teacher-student LLM designed for multi-constraint molecular generation, leveraging a large set of text-molecule pairs to generate molecules that satisfy complex property requirements. Gong et al.[373] introduced TGM-DLM, a diffusion model for text-guided molecule generation that overcomes limitations of autoregressive models in generating precise molecules from textual descriptions. These advances culminate in works like MULTIMODAL-MOLFORMER by Soares et al.[374], which integrates chemical language and physicochemical features with molecular embeddings from MOLFORMER,,[211] significantly enhancing prediction accuracy for complex tasks like biodegradability and PFAS toxicity.

Overall, the shift to multimodal LLMs represents a robust approach to molecular design. By integrating diverse data sources, these models significantly enhance accuracy, interpretability, and scalability, opening new avenues for drug discovery, material design, and molecular property prediction. Combining linguistic, chemical, and graphical data into unified frameworks enables AI-driven models to make more informed predictions and generate innovative molecular structures.

3.6   Textual Scientific LLMs

LLMs are large neural networks known for their performance across various machine learning tasks, with the main advantage of not requiring well-structured data like molecular descriptors.[375] Their true power lies in their ability to handle more challenging tasks, such as extracting insights from less structured data sources like scientific texts or natural language descriptions. In chemistry, this opens doors to new methods of data extraction, classification, and generation, although it depends heavily on the availability of high-quality and diverse datasets (as discussed in Section 3.1. Unfortunately, many datasets are locked behind paywalls or are not machine-readable, limiting the full

potential of LLMs in scientific applications. Encouraging open data initiatives and standardization of formats will play a vital role in expanding LLM applications in chemistry and related fields.

3.6.1 Text Classification

One of the key uses of LLMs in science is text classification, where models sift through vast amounts of scientific literature to extract structured data. For example, Huang et al.[219] applied LLMs to predict patient readmission using clinical data from MIMIC-III.[376] ClinicalBERT[219] used a combination of masked language modeling and next-sentence prediction, followed by fine-tuning on the readmission prediction task. Similarly, Zhao et al.[276] developed EpilepsyLLM by fine-tuning LLaMA using epilepsy data, demonstrating how instruction-based fine-tuning enables models to specialize in highly specific fields. In another application, SciBERT[220] and ScholarBERT[207] adapted BERT to handle scientific literature. SciBERT, developed by Beltagy et al.[220] utilized a specialized tokenizer built for scientific texts from Semantic Scholar,[169] and demonstrated superior performance over fine-tuned BERT models[87] on scientific tasks. This improvement highlighted the importance of tailored vocabularies in model performance.Hong et al.[207] later developed ScholarBERT by pretraining on scientific articles from `Public.Resource.Org` and using RoBERTa optimizations[377] to improve pretraining performance. ScholarBERT was further fine-tuned on the tasks used for evaluation. Despite using a larger dataset, ScholarBERT did not outperform LLMs trained on narrower domain datasets. However, ScholarBERT performed well on specific tasks, such as named entity recognition (NER) within the ScienceExamCER dataset,[378] which involved 3rd to 9th grade science exam questions.

Guo et al.[212] argue that manually curating structured datasets is a sub-optimal, time-consuming, and labor-intensive task. Therefore, they automated data extraction and annotation from scientific papers using `ChemDataExtractor`[379] and their in-house annotation tool.[380] Text extraction tasks, like NER, can be formulated as multi-label classification tasks, which motivates using NER-like approaches and LLMs to extract structured data directly from unstructured text. LLMs developed for data mining include the work of Zhang et al.[381] and Chen et al.[382]

Text extraction tasks, like NER, can be formulated as multi-label classification tasks, which motivates using NER-like approaches and LLMs to extract structured data directly from unstructured text. LLMs developed for data mining include the work of Zhang et al.[381] and Chen et al.[382] Building upon this, Wang et al.[383] conducted a study comparing GPT-4 and ChemDataExtractor[379] for extracting band gap information from materials science literature. They found that GPT-4 achieved a higher level of accuracy (Correctness 87.95% vs. 51.08%) without the need for training data, demonstrating the potential of generative LLMs in domain-specific information extraction tasks. Additionally, LLMs with support for image inputs have been shown to enable accurate data extraction directly from images of tables.[196] A detailed discussion can be found in the study by Schilling-Wilhelmi et al.[384].

In contrast to broad domain models, some LLMs focus on narrow, specialized fields to improve performance. Chem-BERT[212] was pretrained using a BERT model to encode chemical reaction information, followed by fine-tuning a NER head. ChemBERT outperformed other models such as BERT[87] and BioBERT[221] in the product extraction task, presenting an improvement of $\sim 6\%$ in precision. For product role labeling, that is by identifying the role an extracted compound plays in a reaction, ChemBERT showed a $\sim 5\%$ improvement in precision. This suggests that training on narrower datasets enables models to learn specific patterns in the data more effectively.

This trend continued with MatSciBERT,[210] and MaterialsBERT.[385] With MatSciBERT, Gupta et al.[210] fine-tuned SciBERT[220] on the Material Science Corpus (MSC), a curated dataset of materials extracted from Elsevier's scientific papers and improved article subject classification accuracy by 3% compared to SciBERT. In a similar vein, with MaterialsBERT, Shetty et al.[385] fine-tuned PubMedBERT[215] on 2.4 million abstracts, showing incremental precision improvements in NER tasks. BatteryBERT[208] also followed this strategy, outperforming baseline BERT models in battery-related tasks.

Considerable effort has also been devoted to developing LLMs for biology tasks, following a similar trend of training models on large corpora such as Wikipedia, scientific databases, and textbooks, and then fine-tuning them for specific downstream tasks. Shin et al.[214] pretrained various sizes of Megatron-LM,[386] another BERT-like LLM, to create the BioMegatron family of models. These models, which had 345M, 800M, and 1.2B parameters and vocabularies of either 30k or 50k tokens, were pretrained using abstracts from the PubMed dataset and full-text scientific articles from PubMed Central (PMC), similar to BioBERT.[221]

Surprisingly, the largest 1.2B model did not perform better than the smaller ones, with the 345M parameter model using the 50k tokens vocabulary consistently outperforming others in tasks like Named Entity Recognition (NER) and Relation Extraction (RE). NER identifies specific entities, such as chemicals or diseases, while RE determines the relationships between them—both crucial for structuring knowledge from unstructured data. These processes streamline research by converting raw textual information into structured, usable formats for further analysis. This suggests that, for certain tasks, increasing model size does not necessarily lead to better performance. The relevance of model size

was more apparent in the SQuAD[387] dataset, suggesting that LLMs trained on smaller, domain-specific datasets may face limitations in broader generalization.

BioBERT[221] pretrained using data from Wikipedia, textbooks, PubMed abstracts, and the PMC full-text corpus, outperformed the original BERT in all tested benchmarks, and in some cases even achieved state-of-the-art (SOTA) performance in benchmarks such as NCBI disease, 2010 i2b2/VA, BC5CDR, BC4CHEMD, BC2GM, JNLPBA, LINNAEUS, and Species-800. Peng et al.[388] developed BlueBERT, a multi-task BERT model, which was evaluated on the Biomedical Language Understanding Evaluation (BLUE) benchmark.[218] BlueBERT was pretrained on PubMed abstracts and MIMIC-III,[376] and fine-tuned on various BLUE tasks, showing performance similar to BioBERT across multiple benchmarks.

PubMedBERT,[215] following the approach adopted in SciBERT, created a domain-specific vocabulary using 14M abstracts from PubChem for pretraining. In addition to pretraining, the team curated and grouped biomedical datasets to develop BLURB, a comprehensive benchmark for biomedical natural language processing (NLP) tasks, including NER, sentence similarity, document classification, and question-answering. Gu et al.[215] demonstrated that PubMedBERT significantly outperformed other LLMs in the BLURB benchmark, particularly in the PubMedQA and BioQSA datasets. The second-best model in these datasets was BioBERT, emphasizing the importance of domain-specific training for high-performance LLMs in biomedical applications.

Text classification using LLMs, particularly in biomedicine and materials science, has demonstrated that domain-specific pretraining is most effective for enhancing model performance. Models like BioBERT, BlueBERT, and PubMedBERT highlight how focusing on specialized datasets, such as PubMed and MIMIC-III, improves accuracy in tasks like NER, RE, and document classification. These advances illustrate how narrowing the training scope to relevant data enables more effective extraction of structured information from unstructured scientific texts.

In the broader context of this work, text classification serves as a key element that allows AI models to interface with chemical, biological, and medical literature, thereby accelerating progress in drug design, materials discovery, and other research fields. This ability to classify and extract relevant information from scientific texts directly impacts the efficiency and precision of data interpretation, facilitating real-world applications across multiple domains.

3.6.2 Text Generation

Text generation in scientific LLMs offers unique capabilities beyond simply encoding and retrieving information. Unlike encoder-only models, which focus primarily on extracting insights from structured data, decoder models introduce generative abilities that allow them to create new text, answer questions, and classify documents with generated labels. This capability is particularly valuable in scientific fields, where LLMs must not only interpret data but also generate coherent and contextually accurate outputs based on domain-specific instructions. The following models demonstrate how decoder-based architectures enhance generative tasks in natural science, biology, and medical applications.

The Darwin model, as outlined by Xie et al.[280], is one such example. It fine-tunes LLaMA-7B on FAIR, a general QA dataset, followed by specific scientific QA datasets. Instructions for scientific QA were sourced from SciQ[389] and generated using the Scientific Instruction Generation (SIG) model, a tool fine-tuned from Vicuna-7B that converts full-text scientific papers into question-answer pairs. This multi-step training process significantly improved Darwin's performance on regression and classification benchmarks. Notably, LLaMA-7B fine-tuned only on FAIR achieved nearly the same results as the fully fine-tuned model on six out of nine benchmarks, indicating that the integration of domain-specific datasets may not always require extensive fine-tuning for performance gains.

Similarly, Song et al.[390] created HoneyBee by fine-tuning LLaMA-7B and LLaMa-13B on MatSci-Instruct, a dataset with ~52k instructions curated by the authors. HoneyBee outperformed other models, including MatBERT, MatSciB-ERT, GPT, LLaMa, and Claude, within its specialized dataset. However, Zhang et al.[391] showed that HoneyBee did not generalize well to other benchmarks, such as MaScQA[200] and ScQA,[392] highlighting the limitations of models trained on narrow domains in terms of broader applicability.

In biology, BioGPT[282] pretrained a GPT-2 model architecture using 15M abstracts from PubChem corpus. BioGPT was evaluated across four tasks and five benchmarks, including end-to-end relation extraction on BC5CDR, KD-TI, and DDI, question-answering on PubMedQA, document classification on HoC, and text generation on all these benchmarks. After fine-tuning on these tasks (excluding text generation), BioGPT consistently outperformed encoder-only models like BioBERT and PubMedBERT, particularly in relation extraction and document classification. Focusing specifically on text generation, the authors compared BioGPT's outputs to those of GPT-2, concluding that BioGPT was superior, although no quantitative metric was provided for this comparison.

Building on these ideas, Wu et al.[281] pretrained LLaMA2 with the MedC-k dataset, which included 4.8M academic papers and 30K textbooks. This model was further refined through instruction tuning using the MedC-I dataset,

a collection of medical QA problems. PMC-LLaMA [281] outperformed both LLaMa-2 and ChatGPT on multiple biomedical QA benchmarks, even though it was $\sim$ 10 times smaller in size. Notably, the model's performance on MedQA,[393] MedMCQA,[394] and PubMedQA [121] benchmarks improved progressively as additional knowledge was incorporated, the model size increased, and more specific instructions were introduced during tuning.

Text generation through decoder models has significantly expanded the applications of LLMs in scientific fields by enabling the generation of contextual answers and labels from scientific data. Unlike encoder-only models that rely on predefined classifications, decoder models such as Darwin, HoneyBee, and BioGPT can produce outputs tailored to domain-specific needs. This capability is important in fields like biomedicine, where accurate question-answering and document generation are highly valued. By leveraging multi-step pretraining and fine-tuning on specialized datasets, decoder models offer greater flexibility in handling both general and domain-specific tasks.

In the broader context of this work, text generation marks a key methodological advance that complements other LLM tasks, such as classification and extraction. The ability to generate structured responses and create new text from scientific literature accelerates research and discovery across chemistry, biology, and medicine. This generative capacity bridges the gap between raw data and meaningful scientific insights, equipping AI-driven models with a more comprehensive toolkit for addressing complex research challenges.

3.7 The use of ChatGPT in Chemistry

With the rise of ChatGPT, we review here how many researchers have wanted to test the capability of such an accessible decoder-only LLM. Castro Nascimento and Pimentel [395] wrote the first notable paper on ChatGPT's impact on Chemistry. The authors emphasize that LLMs, trained on extensive, uncurated datasets potentially containing errors or secondary sources, may include inaccuracies limiting their ability to predict chemical properties or trends. The paper highlighted that while LLMs could generate seemingly valid responses, they lacked true reasoning or comprehension abilities and would perpetuate existing errors from their training data. However, the authors suggested that these limitations could be addressed in the future. The work serves as a benchmark to qualitatively assess improvements in generative pretrained transformers. For example, five tasks were given to ChatGPT (GPT-3). The accuracy for converting compound names to SMILES representations and vice versa was about 27%, with issues in differentiating alkanes and alkenes, benzene and cyclohexene, or *cis* and *trans* isomers. ChatGPT found reasonable octanol-water partition coefficients with a 31% mean relative error, and a 58% hit rate for coordination compounds' structural information. It had a 100% hit rate for polymer water solubility and a 60% hit rate for molecular point groups. Understandably, the best accuracies were achieved with widely recognized topics. The authors concluded that neither experimental nor computational chemists should fear the development of LLMs or task automation; instead, they advocated for enhancing AI tools tailored to specific problems and integrating them into research as valuable facilitators.

The use of ChatGPT in chemistry remains somewhat limited. Studies by Humphry and Fuller [396], Emenike and Emenike [397], and Fergus et al. [398] focus on its role in chemical education. Some research also explores ChatGPT's application in specific areas, such as the synthesis and functional optimization of Metal-Organic Frameworks (MOFs), where computational modeling is integrated with empirical chemistry research. [399–402] Deb et al. [403] offer a detailed yet subjective evaluation of ChatGPT's capabilities in computational materials science. They demonstrate how ChatGPT assisted with tasks like identifying crystal space groups, generating simulation inputs, refining analyses, and finding relevant resources. Notably, the authors emphasize ChatGPT's potential to write code that optimizes processes and its usefulness for non-experts, particularly in catalyst development for CO_2 capture.

Three key points emerge regarding the use of ChatGPT alone. First, reliable outputs depend on precise and detailed input, as Deb et al. [403] found when ChatGPT struggled to predict or mine crystal structures. Second, standardized methods for reproducing and evaluating GPT-based work remain underdeveloped. Third, achieving complex reasoning likely requires additional chemical tools or agents, aligning with Bloom's Taxonomy. [404,405] Bloom's Taxonomy organizes educational objectives into hierarchical levels: Remembering, Understanding, Applying, Analyzing, Evaluating, and Creating. These range from recalling facts to constructing new concepts from diverse elements. While LLMs and autonomous agents can support lower-level tasks, they currently fall short of replicating higher-order cognitive skills comparable to human expertise.

Currently, LLMs and autonomous agents are limited in replicating higher-level thinking compared to human understanding. To better assess LLMs' capabilities in this domain, we propose using Bloom's Taxonomy as a quality metric. [404,405] This framework offers a structured approach for evaluating the sophistication of LLMs and autonomous agents, especially when addressing complex chemical challenges. It can help quantify their ability to engage in higher-level reasoning and problem-solving.

3.7.1 Automation

The evolution of artificial intelligence in chemistry has fueled the potential for automating scientific processes. For example, in 2019, Coley et al.[406] developed a flow-based synthesis robot proposing synthetic routes and assembling flow reaction systems, tested on medically relevant molecules, and in 2020, Gromski et al.[407] provided a useful exploration of how chemical robots could outperform humans when executing chemical reactions and analyses. They developed the Chemputer, a programmable batch synthesis robot handling reactions like peptide synthesis and Suzuki coupling. In 2021, Grisoni et al.[408] combined deep learning-based molecular generation with on-chip synthesis and testing. The Automated Chemical Design (ACD) framework by Goldman et al.[409] provides a useful taxonomy for automation and experimental integration levels. Thus, automation promises to enhance productivity through increased efficiency, error reduction, and the ability to handle complex problems, as described in several excellent reviews regarding automation in chemistry,[410–416]

This increased productivity may be the only possible approach to exploring the vastness of all chemical space. To fully leverage AI in property prediction, inverse design, and synthesis prediction, it must be integrated with automated synthesis, purification, and testing. This automation should be high-throughput and driven by AI-based autonomous decision-making (sometimes called "lights-out" automation). Janet et al.[411] highlighted challenges in multi-step reactions with intermediate purifications, quantifying uncertainty, and the need for standardized recipe formats. They also stated the limitations of automated decision-making. Organa[417] addresses some of these challenges. It can significantly reduce physical workload and improve users' lab experience by automating diverse common lab routine tasks such as solubility assessment, pH measurement, and recrystallization. Organa interacts with the user through text and audio. The commands are converted into a detailed LLM prompt and used to map the goal to the robot's instructions. Interestingly, Organa is also capable of reasoning over the instructions, giving feedback about the experiments, and producing a written report with the results.

Other limitations exist, like a machine being restricted to pre-defined instructions, its inability to originate new materials, and the lower likelihood of lucky discoveries. Yet, when dedicated tools can be connected to address each step of an automated chemical design, these limitations can be systematically addressed through advancements in LLMs and autonomous agents, discussed in the next section.

4 LLM-based Autonomous Agents

The term "agent" originates in philosophy, referring to entities capable of making decisions.[418] Hence, in artificial intelligence, an "agent" is a system that can perceive its environment, make decisions, and act upon them in response to external stimuli.[419] Language has enabled humans to decide and act to make progress in response to the environment and its stimuli, and so LLMs are naturally ideal for serving as the core of autonomous agents. Thus, in agreement with Gao et al.[420], we define a "language agent" as a model or program (typically based on LLMs) that receives an observation from its environment and executes an action in this environment. Here, environment means a set of tools and a task. Hence, "LLM-based autonomous agents" refer to language agents whose core is based on an LLM model. Comprehensive analyses of these agents are available in the literature,[419–421] but this section highlights key aspects to prepare the reader for future discussions.

There is no agreed definition of the nomenclature to be used to discuss agents. For instance, Gao et al.[420] created a classification scheme that aims to group agents by their autonomy in biological research. This means a level 0 agent has no autonomy and can only be used as a tool, while a level 3 agent can independently create hypotheses, design experiments, and reason.

Following this perspective, Wang et al.[421] categorizes agent components into four modules: profiling, memory, planning, and action. In contrast, Weng[422] also identifies four elements — memory, planning, action, and tools — but with a different emphasis. Meanwhile, Xi et al.[419] proposes a division into three components: brain, perception, and action, integrating profiling, memory, and planning within the brain component, where the brain is typically an LLM. Recently, Sumers et al.[423] proposed Cognitive Architectures for Language Agents (CoALA), a conceptual framework to generalize and ease the design of general-purpose cognitive language agents. In their framework, a larger cognitive architecture composed of modules and processes is defined. CoALA defines a memory, decision-making, and core processing module, in addition to an action space composed of both internal and external tools. While internal tools mainly interact with the memory to support decision-making, external tools make up the environment, as illustrated in Figure 6. Given a task that initiates the environment, the "decision process" runs continuously in a loop, receiving observations and executing actions until the task is completed. For more details, read Sumers et al.[423].

In this review, we define an autonomous agent system as a model (typically an LLM) that continuously receives observations from the environment and executes actions to complete a provided task, as described by Gao et al.[420].
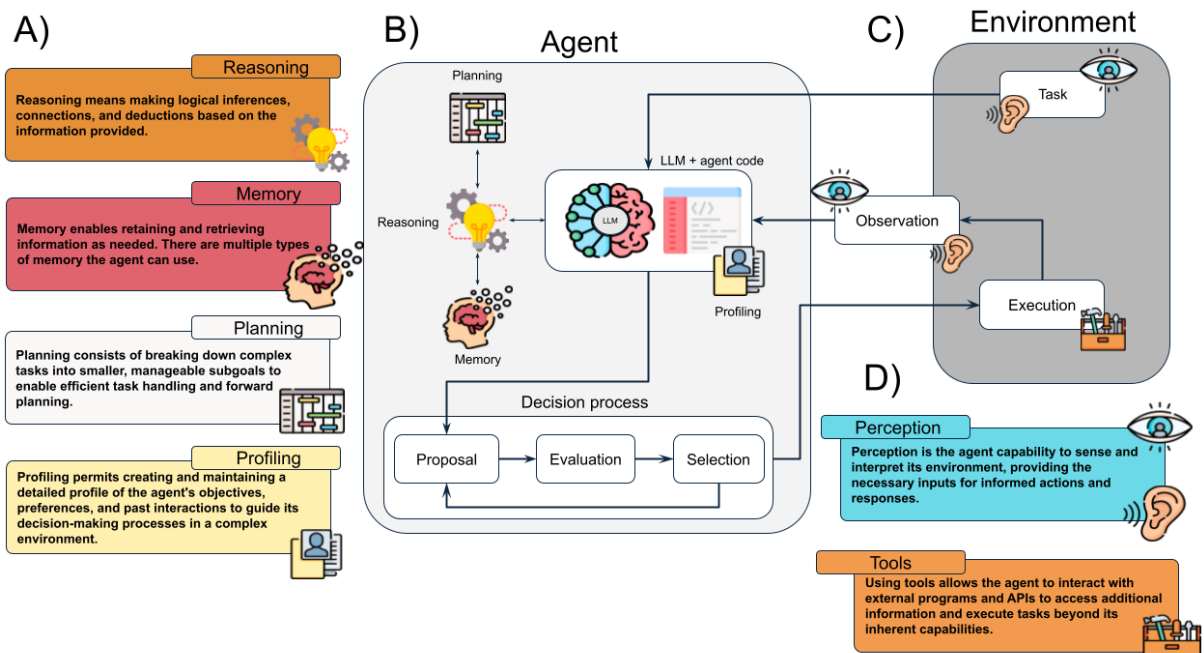
Figure 6: Agent's architecture as defined in this review. According to our definition, an agent is composed of a central program (typically an LLM and the code to implement the agent's dynamic behavior) and the agent modules. The agent continuously receives observations from the environment and decides which action should be executed to complete the task given to it. Here, we define the agent as the set of elements whose decision is trainable, that is, the LLM, the agent code, the decision process, and the agent modules. Given a task, the agent uses the agent modules (memory, reasoning, planning, profiling) and the LLM to decide which action should be executed. This action is executed by calling a tool from the environment. After the action is executed, an observation is produced and fed back to the agent. The agent can use perception to receive inputs in different modalities from the environment. A) Description of agent modules, B) illustration of the agent architecture, C) illustration of the environment components, D) description of tools elements present in the environment.

Nevertheless, in contrast to CoALA,[420] we will rename "internal tools" as "agent modules" and "external tools" simply as "tools", for clarity. The agent consists of trainable decision-making components such as the LLM itself, policy, memory, and reasoning scheme. In contrast, the environment comprises non-trainable elements like the task to be completed, Application Programming Interface (API) access, interfaces with self-driving labs, dataset access, and execution of external code. By referring to decision-making components as agent modules, we emphasize their inclusion as parts of the agent. By referring to non-trainable elements as tools, we highlight their role as part of the environment. We discuss six main types of actions. As shown in Figure 6, four of the six, memory, planning, reasoning, and profiling are agent modules. The remaining two actions (or tools) and perception are part of the environment. Since the perception is how the agent interacts with the environment and is not a trainable decision, we therefore included it as part of the environment.

4.1 Memory Module

The role of the memory module is to store and recall information from past interactions and experiences to inform future decisions and actions. There are multiple types of memory in agents, namely sensory memory, short-term memory, and long-term memory. A major challenge in using agents is the limited context window, which restricts the amount of in-context information and can lead to information loss, thereby impacting the effectiveness of short-term and long-term memory. Solutions involve summarizing memory content,[424] compressing memories into vectors,[425–427] and utilizing vector databases[428] or combinations thereof,[429] with various databases available such as ChromaDB, FAISS, Pinecone, Weaviate, Annoy, and ScaNN.[430] Addressing these challenges to enhance agent memory continues to be a significant area of research.[431] Sensory, or procedural memory is knowledge embedded into the model's parameters during pretraining and/or in heuristics implemented into the agent's code. Short-term, or working, memory includes

the agent's finite knowledge during a task, incorporating interaction history and techniques like in-context learning[93] (ICL), which leverages the limited input's context length for information retention. Long-term memory involves storing information externally, typically through an embedded vector representation in an external database. In the original CoALA[420] paper, long-term memory is further categorized as episodic, which registers previous experiences, and semantic, which stores general information about the world.

4.2 Planning and Reasoning Modules

The planning and reasoning module is made of two components. Planning involves identifying a sequence of actions required to achieve a specified goal. In the context of language agents, this means generating steps or strategies that the model can follow to solve a problem or answer a question, which can be enhanced with retrieval from previous experiences,[432] and from feedback from post-execution reasoning.[433,434] We note that Retrieval-Augmented Generation (RAG) enhances the planning phase by enabling models to access external knowledge bases, integrating retrieved information into the generation process. This approach improves accuracy and relevance, especially when handling complex or knowledge-intensive tasks. Reasoning refers to the process of drawing conclusions or making decisions based on available information and logical steps. For example, there are studies that demonstrate the benefits of LLM reasoning for question answering, where new context tokens can be integrated in a step-by-step way to guide the model towards more accurate answers.[435–440] One popular reasoning strategy is Chain-of-Thought (CoT),[107,437,441–444] a reasoning strategy which substantially boosts QA performance by generating intermediate reasoning steps in a sequential manner. CoT involves breaking down complex problems into smaller, manageable steps, allowing the model to work through reasoning one step at a time rather than attempting to solve the entire problem at once. CoT thereby reduces hallucinations and enhances interpretability, as demonstrated by improved results in models like PaLM[445] and GPT-3 with benchmarks like GSM8K,[446] SVAMPs,[447] and MAWPS.[448]

In advanced reasoning, final tasks are often decomposed into intermediary ones using a cascading approach, similar to Zero-shot-CoT[436] and RePrompt.[433] However, while CoT is considered as single-path reasoning, CoT extensions like Tree-of-Thoughts,[439] Graph-of-Thoughts,[449] Self-consistent CoT,[438] and Algorithm-of-Thoughts[450] offer multi-path reasoning. Furthermore, other models have pitted multiple agents against each other to debate or discuss various reasoning paths,[451–453] while others use external planners to create plans.[454,455] A feedback step during the execution of the plan was a further extension of the CoT ideas; this enables agents to refine their actions based on environmental responses adaptively, which is crucial for complex tasks.[456,457]

Another interesting reasoning scheme is the Chain-of-Verification(CoVe),[434] where once an answer is generated, another LLM is prompted to generate a set of verification questions to check for agreement between the original answer and the answers to the verification questions such that the final answer can be refined. The ReAct[439] – Reason+Act – model proposes adding an observation step after acting. This means the LLM first reasons about the task and determines the necessary step for its execution, it performs the action and then observes the action's result. Reasoning on that result, it can subsequently perform the following step. Similarly, Reflexion[107] also implements a reasoning step after executing an action. However, Reflexion implements an evaluator and self-reflection LLMs to not only reason about each step but also to evaluate the current trajectory the agent is following using a long-term memory module. As the context increases, it may become challenging for agents to deal with the long prompt. Aiming to solve this issue, the Chain-of-Agents (CoA)[80] extends reasoning schemes that leverage multi-agent collaboration to reason over long contexts. This framework employs workers and manager agents to process and synthesize information to generate the final response. CoA demonstrated improvements of up to 10% when compared against an RAG baseline.

ReAct and Reflexion are closed-ended approaches where the agent starts with all the tools and must determine which to use. To address more open-world challenges, Wang et al.[458] introduced the Describe, Explain, Plan, and Select (DEPS) method, which extends this approach. Lastly, human inputs can also be used to provide feedback to the agent. Providing feedback using a human-in-the-loop approach is particularly interesting in fields where safety is a main concern.

4.3 Profiling Module

LLMs can be configured to perform in specific roles, such as coders, professors, students, and domain experts, through a process known as profiling. Language agents can thus incorporate the profile through the LLM or through the agent code. The profiling approach involves inputting psychological characteristics to the agent, significantly impacting its decision-making process[459–462]. Profiling enables the creation of multi-agent systems that simulate societal interactions, with each agent embodying a unique persona within the group[432,463]. The most prevalent technique for profiling, called "handcrafting", requires manually defining the agent's profile, often through prompts or system messages[464,465]. While profiling can also be automated with LLMs[466], that automation method may only be suited for generating large

numbers of agents since it offers less control over their overall behavior. An interesting application of profiling is the development of agent sets that reflect demographic distributions[467].

4.4 Perception

Perception is an analog to the human sensory system, which interprets multimodal information such as text, images, or auditory data, transforming it into a format comprehensible by LLMs, as demonstrated by SAM,[468] GPT4-V,[469] LLaVa,[470] Fuyu8B,[134] and BuboGPT.[471] In our proposed architecture, the perception is responsible for converting the task and the observations to a data representation that can be understood by the agent. Moreover, advancements in LLMs have led to the development of even more versatile models, such as the any-to-any Next-GPT[135] and the any-to-text Macaw-LLM.[472] Employing such multimodal LLMs in decision-making processes can simplify perception tasks for agents, with several studies exploring their use in autonomous systems.[473,474]

4.5 Tools

In our proposed definition (see Figure 6b), tools or actions are part of the environment. The agent can interact with this environment by deciding which action to execute through the decision-making process. The set of all possible actions that can be selected is also known as the "action space".

The decision process is composed of three main steps: proposal, evaluation, and selection. During the proposal, one or more action candidates are selected using reasoning,[439] code structures,[432,475] or simply by selecting every tool available.[476][435,438,477] The evaluation process consists of evaluating each selected action according to some metric to predict which action would bring more value to the agent. Lastly, the action is selected and executed.

Given that pretrained parameters (sensory memory) are limited, the model must use tools for complex tasks in order to provide reliable answers. However, LLMs need to learn how to interact with the action space and how and when to use those tools most accurately.[478] LLMs can be pretrained or fine-tuned with examples of tool use, enabling them to operate tools and directly retrieve tool calls from sensory memory during a zero-shot generation.[479] Recent studies investigate this approach, particularly focusing on open-source LLMs.[480–482]

As foundational AI models become more advanced, their abilities can be expanded. It was shown that general-purpose foundation models can reason and select tools even with no fine-tuning. For example, MRKL[483] implements an extendable set of specialized tools known as neuro-symbolic modules and a smart "router" system to retrieve the best module based on the textual input. Specifically, this router smartly parses the agent's output and selects which neuro-symbolic module is more suitable to perform the task following some heuristic. These neuro-symbolic modules are designed to handle specific tasks or types of information and are equipped with built-in capabilities and task-relevant knowledge. This pre-specialization allows the model to perform domain-specific tasks without needing a separate, domain-specific dataset. This design addresses the problem of LLMs lacking domain-specific knowledge and eliminates the need for the costly and time-consuming LLM fine-tuning step, using specialized data annotation.[484] The router can receive support from a reasoning strategy to help select the tools[484] or follow a previously created plan.[458] Recent advances have shown that LLMs can develop new tools of their own,[485–487] enabling agents to operate, as needed, in dynamic and unpredictable "open-worlds", on unseen problems as illustrated by Voyager.[475] This capability allows agents to evolve and improve continually.

5 LLM-Based Autonomous Agents in Scientific Research

The previous section introduced key concepts relevant to any description of the development of autonomous agents. Here, we now focus on which agents were developed for scientific purposes, and ultimately for chemistry. Previous sections of this review have discussed how LLMs could be powerful in addressing challenges in molecular property prediction, inverse design, and synthesis prediction. When we consider the value of agents in chemistry and the ability to combine tools that, for example, search the internet for established synthetic procedures, look up experimental properties, and control robotic synthesis and characterization systems, we can see how autonomous agents powerfully align with the broader theme of automation, which will lead to an acceleration of chemical research and application.

Table 4: Scientific LLM systems and agents. We identify the studies we classified as an agent with the icon 🤖[488] and multi-agent systems with the icon 🤖🤖. Ⓦ, Ⓢ, and Ⓛ[489] mean the agent bases his behavior on sensory, short, and long memory components, respectively. Besides the textual capabilities of LLM-based agents, 🔊[490] and 👁[491] mean the agent has additional audio and visual perception, respectively. The release date column displays the date of the first publication for each paper. When available, the publication date of the last updated version is displayed between parentheses.

Agent	Memory	Planning	Reasoning	Action	Release date
PaperQA2[492]🤖	ⓈⓁ	✓	✓	Tools to search the scientific literature, gather evidence, and answer questions	2024.09
LLaMP[493]🤖	ⓈⓁ		✓	Tools for database access, literature search, and atomistic simulations	2024.01 (2024.06)
SGA[494]🤖	Ⓢ			Employ the LLMs in a optimization loop	2024.05
CRISPR-GPT[495]🤖🤖	Ⓢ	✓	✓	Tool for gene editing experiments design	2024.04
TAIS[496]🤖🤖	Ⓢ	✓	✓	Tools for gene expression data analysis	2024.02
ChemReasoner[497]🤖	ⓌⓈ	✓	✓	Tools for heuristic search, 3D structure generation, and prediction using GNNs	2024.02 (2024.06)
SciAgent[498]🤖	ⓌⓁ	✓	✓	Trained Mistral for tool usage. Evaluated it using MathToolBench's tools	2024.02
STORM[499]	ⓈⓁ	✓	✓	Article writing using retrieval from multi-LLM conversations and pre-generated outline	2024.02 (2024.04)
Völker et al.[500]	Ⓢ			Regression with ICL and text retrieval	2024.02
ProtAgent[501]🤖🤖	ⓈⓁ	✓	✓	Tools for proteins information retrieval, analyzing, *de novo* design, and 3D folded structure generation	2024.01 (2024.05)
Organa[417]🤖🔊👁	ⓈⓁ	✓	✓	Tools for common lab procedures, reasoning about experimental results, and report writing	2024.01
PaperQA[502]🤖	ⓈⓁ	✓	✓	Tools to search the scientific literature, gather evidence, and answer questions	2023.12
WikiCrow[503]🤖	ⓈⓁ	✓	✓	Uses PaperQA as a tool	2023.12
Coscientist[48]🤖	Ⓢ	✓	✓	Tools for running Python code, web-searching, and interacting with lab equipment	2023.04 (2023.12)
Eunomia[504]🤖	Ⓢ		✓	Tools for literature and dataset searching and a chain–of–verification loop	2023.12
CALMS[505]🤖	Ⓛ			Tools for using the Materials Project API, designing experiments, and using a hardware API to perform the experiment	2023.12
CoQuest[506]🤖	Ⓢ		✓	Research question generations and tools for literature visualization using a graph organization	2023.10 (2024.03)
eXpertAI[343]🤖	Ⓢ		✓	Tools for applying XAI methods	2023.11
BioPlanner[507]🤖	Ⓢ	✓	✓	Tool for protocol searching in the BioProt dataset	2023.10
IBM ChatChem[508]🤖	Ⓢ			Tools for cheminformatics and accessing GT4SD and HuggingFace models	2023.09

Continued on next page

Table 4 – continued from previous page

Agent	Memory	Planning	Reasoning	Action	Release date
ChatMOF [509] 🤖	Ⓦ Ⓢ		✓	Tools for database search, property prediction, and MOF's structure generation	2023.08 (2024.06)
AmadeusGPT [47] 🤖👁	Ⓢ Ⓛ		✓	Tools for writing and executing code for computer vision, machine learning, and spatial-temporal reasoning	2023.07
i-Digest [510] 🎧	Ⓢ			Uses the whisper model to process audio transcription from classes and write summaries and following up questions	2023.06
BOLLaMa [511]	Ⓦ			Implements an LLM interface to ease the usage of their BO code	2023.06
text2concrete [512]	Ⓢ			Uses ICL to predict compressive strength from concrete formulation	2023.06
MAPI_LLM [513] 🤖	Ⓢ Ⓛ		✓	Database access and LLM prediction using ICL	2023.06
BO-LIFT [514]	Ⓢ			Regression using ICL and text retrieval	2023.04
ChemCrow [47] 🤖	Ⓢ		✓	Molecular, cheminformatics, search and critique tools	2023.04 (2024.05)

It was Hocky and White [515] who discussed the early stages of models that could automate programming and, hence, the expected impacts in chemistry. Then, early work by White et al. [516] applied LLMs that could generate code to a benchmark set of chemical problems. In that case, not only were LLMs demonstrated to possess a notable understanding of chemistry, based on accurate question answering, but White et al. [516] imagined a potential to use them as base models to control knowledge augmentation and a variety of other tools. Thus, these LLMs could be used to execute routine tasks, optimize procedures, and enhance the retrieval of information from scientific literature across a range of scientific domains. To the best of our knowledge, this is the first review of autonomous agents in chemistry that have evolved since these two visionary conceptual perspectives. A deeper exploration follows below. One driving motivation for the need to augment LLMs with a more pertinent and dedicated knowledge base is the need to circumvent problems of a limited context prompt window, and the restriction that once an LLM is trained, any new information is beyond it's reach since it necessarily has fallen outside its corpus of training data. Furthermore, LLMs are also known to hallucinate. Their predictions are probabilistic and, in science, if experimental evidence is available, then there is great value in building from known domain-specific information. Some improved prompt engineering can aid in the generation of results that are more likely to be accurate, but the use of autonomous agents may solve such problems completely in this next phase of AI in chemistry. In fact, even adding one or two components when building an agent, as opposed to a whole suite, has shown some significant gains.

Building on this foundation, Ramos et al. [514] illustrated that LLMs could directly predict experimental outcomes from natural language descriptions, incorporating this ability into a Bayesian optimization (BO) algorithm to streamline chemical processes. Using in-context learning (ICL), where a model learns from examples provided during inference without requiring retraining, their approach avoided additional model training or fine-tuning, simplifying the optimization process. In a similar vein, Kristiadi et al. [517] demonstrated similar results with a smaller, domain-specific model, using parameter-efficient fine-tuning (PEFT) rather than ICL. Ranković and Schwaller [518] also explored BO using natural language. They used an LLM to encode chemical reaction procedures, described using natural language, and then trained a Gaussian process (GP) head to predict the reaction yield from the latent encoded representation of the procedure. By keeping the LLM frozen and only updating the multilayer perceptron (MLP) head, this approach minimized training time. Völker et al. [500] extended these ideas by sampling multiple model completions and adding a verifier model to select the next best step in the BO algorithm. They also used ICL and a short-term memory component to optimize alkali-activated concrete mix design. These examples demonstrate how agent-based systems can execute complex optimization algorithms step by step, directly contributing to automation and more efficient experimental design.

To better promote new ideas regarding AI in scientific research, Jablonka et al. [519] organized a one-day hackathon in March 2023 where participants developed 14 innovative projects addressing chemical problems centered on predictive modeling, automation, knowledge extraction, and education. Several agent-based approaches emerged from this hackathon. First, MAPI_LLM [513] is an agent with access to the Materials Project API (MAPI) database that receives a

query asking for a property of a material and then retrieves the relevant information from the dataset. If the material is not available on MAPI, the agent can search for similar materials and use in-context learning (ICL) to provide a prediction of the requested property. Additionally, MAPI_LLM also has a reaction module for synthesis proposal. Second, Rankovic et al.[511] used LLMs to make BO algorithms more accessible to a broader group of scientists; BOLLaMa implements a natural language interface to easily interact with BO software developed by their group.[520] Third, and similar to Ramos et al.[514] and Ranković and Schwaller[518] who employed LLMs in BO, Weiser et al.[521] focused on genetic algorithms (GA), a different optimization algorithm. In GA, pieces of information are stochastically combined and evaluated to guide the algorithm during the optimization. For chemistry, these pieces are often molecular fragments that are combined to compose a final whole molecular structure. Thus, Weiser et al.[521] used LLMs to implement common GA operators under the hypothesis that LLMs can generate new combined molecules better than random cross-over due to their sensory memory. Fourth, InsightGraph[522] can draw general relationships between materials and their properties from JSON files. Circi and Badhwar[522] showed that LLMs can understand the structured data from a JSON format and reorganize the information in a knowledge graph. Further refinement of this tool could automate the process of describing relationships between materials across various scientific reports, a task that remains labor-intensive today. Fifth, Kruschwitz et al.[512] used ICL and LLMs to accurately predict the compressive strength of concrete formulations; Text2Concrete achieved predictive accuracy comparable with a Gaussian process regression (GPR) model, with the advantage that design principles can be easily added as context. This model was successfully applied in a BO algorithm following the Ramos et al.[514] approach.[500] For education purposes, multiple authors have raised the discussion about how LLMs can be used to support educators' and instructors' daily work.[523–527] Finally, in this direction, Mouriño et al.[510] developed i-Digest, an agent whose perception module can understand audio tracks and video recordings. These audio recordings are transcribed to text using the Whisper[138] model, and therefore, i-Digest is a digital tutor that generates questions to help students test their knowledge about the course material. These are just a few examples to showcase the capabilities of AI systems to innovate and generate solutions rapidly.

More recently, Ma et al.[498] showed that agents can be trained to use tools. SciAgent[498] was developed under the premise that finetuning LLMs for domain-specific applications is often impractical. Nevertheless, the agent can be fine-tuned with a set of tools that will enable them to perform well in a domain-specific task. These tools, typically Python functions, enable SciAgent to plan, retrieve, and use these tools to facilitate reasoning and answer domain-related questions effectively. The benchmark developed for SciAgent, known as `SciToolBench`, includes five distinct domains, each equipped with a set of questions and corresponding tools. The development of its retrieval and planning modules involved finetuning different LLMs on the MathFunc benchmark, resulting in a notable performance improvement of approximately $\sim 20\%$ across all domains in `SciToolBench` compared to other LLMs.

These examples demonstrate the rapidly growing potential of autonomous agents to drive innovation and automation across scientific tasks, from optimizing experiments and materials discovery to enhancing education. As these tools advance, they streamline processes, generate new insights, and empower researchers to tackle complex challenges. By combining reasoning, optimization, and tool usage in real time, agents mark a significant leap in AI-driven research. In the next section, we focus on how agents are transforming literature review processes, a critical aspect of scientific discovery.

5.1 Agents for Literature Review

Another fantastic opportunity for automation in the sciences is associated with high-quality literature review, a pivotal aspect of scientific research that requires reading and selecting relevant information from large numbers of papers, and thereby distilling the current state of knowledge relevant to a particular research direction. This extremely time-consuming task is being revolutionized by advanced AI tools designed to automate and enhance such analysis and summarization.

PaperQA introduces a robust model that significantly reduces misinformation while improving the efficiency of information retrieval. This agent retrieves papers from online scientific databases, reasons about their content, and performs question-answering (QA) tasks. Its mechanism involves three primary components—"`search`", "`gather_evidence`", and "`answer_question`" and the authors adapted the Retrieval-Augmented Generation (RAG)[528] algorithm to include inner loops on each step. For instance, PaperQA can perform multiple rounds of `search` and `gather_evidence` if, upon reflection, not have enough evidence has been acquired to successfully `answer_question`.

To further validate its capabilities, the authors developed a new benchmark called LitQA, specifically designed to evaluate the performance of models like PaperQA in solving complex, real-world scientific questions. LitQA focuses on tasks that mimic the intricacy of scientific inquiry, comprising 50 multiple-choice questions derived from biomedical papers published post-September 2021, ensuring that these papers were not included in the training data of LLMs. In this challenging setting, PaperQA not only meets but exceeds human performance, achieving a precision rate of 87.9% and an accuracy score of 69.5%, compared to the human baseline of 66.8%.[502] By applying the RAG technique

to full-text scientific papers, PaperQA sets a new standard in QA capabilities, achieving human-like performance in curated datasets without hallucination or selecting irrelevant citations.[502]

Building on top of PaperQA, WikiCrow exemplifies the practical application of AI in generating concise and relevant Wikipedia-style summaries. The authors show that while 16% of a human-created Wikipedia article comprises irrelevant statements, WikiCrow displays irrelevant information only 3% of the time. Their system also added 5% more correct citations when compared with original articles. Moreover, thanks to its foundation in the PaperQA framework,[502] WikiCrow achieves remarkable cost-efficiency. The authors estimate that WikiCrow can accomplish in a few days what would take humans approximately 60,000 hours, or about 6.8 years, thereby underscoring its ability to rapidly produce extensive scientific content. This efficiency exemplifies the reliability and transformative potential of AI in content creation.[503]

Following a different approach, the STORM model also addressed the problem of writing Wikipedia-like summaries, where the STORM acronym represents the Synthesis of Topic Outlines through Retrieval and Multi-perspective questions.[499] This approach implements a two-step procedure. First, STORM retrieves multiple articles on a topic and uses an LLM to integrate various perspectives into a cohesive outline. Second, this outline is used to write each section of the Wikipedia-like summary individually. To create the outline, multiple articles discussing the topic of interest are retrieved by an "expert" LLM, which processes each one to create N perspectives. Each perspective is then fed to a "writer" LLM, and a conversation is initiated between writer and expert. Finally, the N conversations are used to design the final outline. The outline and the set of references, accessed by RAG, are given to the writer LLM. The writer LLM is prompted to use these inputs to generate each section of the article sequentially. Following this, all sections are merged and refined to eliminate redundancies and enhance coherence. Upon human evaluation, STORM is reported to be $\sim 25\%$ more organized and present $\sim 10\%$ better coverage when compared to a pure RAG approach. However, it was also less informative than human-written Wikipedia pages, and STORM presented a transfer of internet-borne biases, producing emotional articles, which is a major concern.

5.2 Agents for Chemical Innovation

Transitioning from literature synthesis to practical chemistry applications, we next explore how LLM-based agents have proven their capabilities to revolutionize routine chemical tasks toward an acceleration of molecular discovery and scientific research. Agents are flexible entities capable of developing prompt-specific workflows and executing a plan toward accomplishing a specific task. ChemCrow[47] introduced a significant shift in how LLMs would be applied in chemistry, given that LLMs alone do not access information outside of their training data nor can they directly perform chemistry-related tasks.

By augmenting LLMs with common chemical tools, computational or robotic, ChemCrow automates a broad spectrum of routine chemical tasks, demonstrating a significant leap in LLM applicability. Under human evaluation, ChemCrow consistently outperformed GPT-4, achieving an accuracy score of 9.24/10 compared to 4.79/10.[47] The developers of ChemCrow have also considered the ethical implications and potential risks associated with its capabilities. ChemCrow's high potential could be misused and exploited for malicious objectives, and therefore the authors have implemented safety checks and guidelines to prevent such misuse, or "dual usage". Additionally, they acknowledge that ChemCrow, relying on an LLM, may not always provide completely accurate answers due to gaps in its chemical knowledge. As such, they recommend careful and responsible use of the tool, along with thorough scrutiny of its outputs. In summary, while ChemCrow presents a powerful new chemical assistant,[47] oversight of its use is required, and this agent's access to tools has been deliberately limited to enhance security and avoid misuse.

Similarly to ChemCrow,[47] Chemist-X[529] uses RAG to get up-to-date literature information and use it to reliably solve a user's questions. Nevertheless, Chemist-X focuses on designing chemical reactions to achieve a given molecule. It works in three phases: (1) First, the agent searches molecule databases for similar molecules to the given molecule, then (2) it searches online literature searching for chemical reactions capable of converting the list of similar molecules in the target. Lastly, (3) machine learning models are used to propose the reaction conditions. To validate their agent, the authors used Chemist-X to design a High-Throughput Screening (HTS) experiment aiming to produce 6-(1-methyl-1H-indazol-4-yl), resulting in a maximum yield of 98.6%.

Another system called Coscientist[48] system exemplifies the integration of semi-autonomous robots in planning, conceiving, and performing chemical reactions with minimal human intervention. At its core, the system features a main module named 'PLANNER', which is supported by four submodules. These submodules, or tools, are responsible for performing actions such as searching the web for organic synthesis, executing Python code, searching the hardware documentation, and performing a reaction in an automated lab.[48] Utilizing this framework, the Coscientist successfully conducted two types of chemical coupling reactions, Suzuki-Miyaura and Sonogashira, in a semi-automated fashion, with manual handling of initial reagents and solvents. Additionally, Coscientist was also used to optimize reaction

conditions. In contrast to Ramos et al.[514], who used LLMs within a Bayesian Optimization (BO) algorithm as a surrogate model, Boiko et al.[48] approached the optimization task as a strategic "game" aimed at maximizing reaction yield by selecting optimal reaction conditions. This demonstrates the ability of GPT-4 to effectively reason about popular chemical reactions – possibly via comprehensive coverage in pretraining. The authors have indicated that the code for their agent will be released following changes in U.S. regulations on AI and its scientific applications. At the time of writing, the code remains unreleased, but a simple example that calculates the square roots of random numbers has been provided to illustrate their approach.[48] These examples underscore the transformative role of LLMs in enhancing and automating chemical processes, which will likely accelerate chemical discovery.

Automated workflows in protein research have also been explored. ProtAgent[501] is a multi-agent system designed to automate and optimize protein design with minimal human intervention. This system comprises three primary agents: `Planner`, `Assistant`, and `Critic`. The `Planner` is tasked with devising a strategy to address the given problem, the `Assistant` executes the plan using specialized tools and API calls, and the `Critic` supervises the entire process, providing feedback and analyzing outcomes. These agents collaborate through a dynamic group chat managed by a fourth agent, the `Chat Manager`. Tasks executed by this team include protein retrieval and analysis, *de novo* protein design, and conditioned protein design using Chroma[530] and OmegaFold.[531]

Similarly to ProtAgent, Liu et al.[496] created a team of AI-made scientists (TAIS) to conduct scientific discovery without human intervention. However, their agents have roles analogous to human roles, such as project manager, data engineer, code reviewer, statistician, and domain expert. While in ProtAgent[501] agents interact through the `Chat Manager` only, TAIS[496] enables AI scientists to interact between themselves directly using pre-defined collaboration pipelines. To evaluate TAIS, the authors curated the Genetic Question Exploration (GenQEX) benchmark, which consists of 457 selected genetic data questions. As a case study, the authors show TAIS's answer to the prompt ''What genes are associated with Pancreatic Cancer when considering conditions related to Vitamin D Levels?''. The system identified 20+ genes with a prediction accuracy of 80%.

Innovation can also be achieved by looking into data from a different point-of-view to get new insights. Automating querying databases was investigated by Ramos et al.[513] with a ReAct agent with access to the MAPI dataset. This concept was extended by Chiang et al.[493] using LLaMP,[493] which is a RAG-based ReAct agent that can interact with MAPI, arXiv, Wikipedia, and has access to atomistic simulation tools. The authors showed that grounding the responses on high-fidelity information (a well-known dataset) enabled the agent to perform inferences without fine-tuning.

The agents in chemistry, as exemplified by ChemCrow[47] and Coscientist,[48] highlight a significant shift towards automation and enhanced efficiency in molecular discovery and scientific research. These systems demonstrate the potential of integrating LLMs with chemical tools and automation frameworks, achieving impressive accuracy and effectiveness in tasks ranging from routine chemical operations to complex reaction optimizations. Similarly, ProtAgent[501] and TAIS[496] systems showcase the versatility of multi-agent frameworks in automating protein design and genetic research, pushing the boundaries of what AI-driven scientific discovery can achieve. These studies collectively showcase the incredible potential of agents in chemical and biological research, promising automation of routine tasks, easing the application of advanced techniques and analyses, and accelerating discoveries. However, they also underscore the necessity for meticulous oversight and responsible development to harness their full potential while mitigating risks.[532]

5.3 Agents for Experiments Planning

Building on the capabilities of ChemCrow and Coscientist in automating chemistry-related tasks, recent advances have focused on bridging the gap between virtual agents and physical laboratory environments. For example, Context-Aware Language Models for Science (CALMS),[505] BioPlanner,[507] and CRISPR-GPT[495] focus on giving support to researchers with wet-lab experimental design and data analysis.

CALMS[505] focuses on improving laboratory efficiency through the operation of instruments and management of complex experiments, employing conversational LLMs to interact with scientists during experiments. In addition, this agent can perform actions using lab equipment after lab equipment APIs have been provided to the agent as tools. CALMS was designed to enhance instrument usability and speed up scientific discovery, providing on-the-spot assistance for complex experimental setups, such as tomography scans, and enabling fully automated experiments. For instance, its capability was showcased through the operation of a real-world diffractometer. Although CALMS excelled in several tasks, a comparison between GPT-3.5 and Vicuna 1.5 revealed Vicuna's limitations in handling tools.

In contrast, BioPlanner[507] significantly improves the efficiency of scientific experimentation by creating pseudocode representations of experimental procedures, showcasing AI's capacity to streamline scientific workflows. Therefore, Rather than interacting directly with lab equipment through APIs, BioPlanner creates innovative experimental protocols that can be expanded upon within a laboratory setting. The initial step in BioPlanner's process involves assessing the

capability of LLMs to produce structured pseudocode based on detailed natural language descriptions of experimental procedures.

In testing, BioPlanner successfully generated correct pseudocode for 59 out of 100 procedures using GPT-4, although the most common errors involved omitted units. Afterward, the authors used BioPlanned to generate a procedure for culturing an E.coli bacteria colony and storing it with cryopreservation, which ran successfully.

Focusing on gene editing experiments, CRISPR-GPT[495] is an agent developed to design experiments iteratively with constant human feedback. CRISPR-GPT[495] aims to bridge the gap for non-experts by simplifying this process into manageable steps solvable by an LLM with access to useful tools. This agent operates in three modes based on user prompts: "Meta mode" provides predefined pipelines for common gene-editing scenarios; "Auto mode" uses the LLM to plan a sequence of tasks; and "Q&A mode" answers general questions about the experimental design. The authors demonstrate that based on human evaluations, CRISPR-GPT outperforms GPT-3.5 and GPT-4 in accuracy, reasoning, completeness, and conciseness. Additionally, they applied CRISPR-GPT to design real-world experiments for knocking out TGFBR1, SNAI1, BAX, and BCL2L1 in the human A375 cell line, achieving an editing efficiency of approximately 70% for each gene.

Following the ideas of developing agents for automating experimental protocol generation, Ruan et al.[533] created a multi-agents system composed of 6 agents: `Literature Scouter`, `Experiment Designer`, `Hardware Executor`, `Spectrum Analyzer`, `Separation Instructor`, and `Result Interpreter`. The Large Language Models-based Reaction Development Framework (LLM-RDF)[533] automates every step of the synthesis workflow. While other studies focus on the literature review,[381,534,535] HTS,[529] and reaction optimization,[48,536] LLM-RDF can support researchers from literature search until the product purification. Using this system, the authors showed they could design a copper/TEMPO catalyzed alcohol oxidation reaction, optimize reaction conditions, engineer a scale-up, and purify the products, obtaining a yield of 86% and a purity >98% while producing 1 gram of product.

Interestingly, despite covering different fields and having diverse goals, all of these studies, from the fully automated systems like CALMS,[505] and LLM-RDF,[533] to human-driven protocols in BioPlanner[507] and CRISPR-GPT,[495] share a "human-in-the-loop" approach. This ensures the researcher remains integral to the development process, enhancing reliability and mitigating potential agent limitations, such as errors or hallucinations. Moreover, this approach addresses risks and dual-use concerns, as humans can assess whether the agents' suggestions are safe.[421,537] On a slightly different track, Organa[417] fully automates the laboratory workload while providing feedback to the researcher and producing reports with the results, as discussed on Section 3.7.1.

Autonomous agents significantly enhance productivity and efficiency in scientific research, but human creativity and decision-making remain vital to ensure quality and safety. In the next section, we explore agents designed to automate cheminformatics tasks, continuing our focus on how AI systems are reshaping the chemical sciences.

5.4 Agents for Automating Cheminformatics Tasks

Cheminformatics consists of applying information technology techniques to convert physicochemical information into knowledge. The process of solving cheminformatics problems commonly involves retrieving, processing, and analyzing chemical data.[538] Getting inspiration from ChemCrow[47] ideas, Chemistry Agent Connecting Tool Usage to Science (CACTUS)[539] focused on assisting scientists by automating cheminformatics tasks. CACTUS automates the applications of multiple cheminformatics tools, such as property prediction and calculation, while maintaining the human-in-the-loop for molecular discovery. The authors investigated the performance of a diverse set of open-source LLMs, where Gemma-7B and Mistral-7B demonstrated superior performance against LLaMA-7B and Falcon-7B. In addition, the authors reported that adding domain-specific information in the prompt to align the agent to chemistry problems considerably increases a model's performance. For instance, predicting drug-likeness with a Gemma-7B agent improves the accuracy of $\sim 60\%$ when aligning the agent in this way, and prompt alignment improved the prediction of all properties they studied.

Further illustrating the versatility of AI in scientific research and domain-specific tools usage is ChatMOF,[509] which focuses on the prediction and generation of Metal-Organic Frameworks (MOFs). ChatMOF integrates MOF databases with its MOFTransformer[540] predictor module, thereby showcasing the innovative use of genetic algorithms in guiding generative tasks from associated predictions. The authors showed that ChatMOF achieved an accuracy of $\sim 90\%$ in search and prediction tasks while generative tasks have an accuracy of $\sim 70\%$. The genetic algorithm used by ChatMOF allows for the generation of a diverse array of MOF structures, which can be further refined based on specific properties requested by users. For instance, when prompted to, "generate structures with the largest surface area", the system initially generated a broad distribution of structures with surface area centered in 3784 m^2/g, and the GA evolves it to a narrower distribution with a peak at 5554 m^2/g after only three generations. It is important to note that even though ChatMOF has access to a dataset of experimental values for MOFs, language model predictions guide their GA, and

no further validation has been made. Lastly, Ansari and Moosavi[504] developed Eunomia, another domain-specific autonomous AI agent that leverages existing knowledge to answer questions about materials. Eunomia[504] can use chemistry tools to access a variety of datasets, scientific papers and unstructured texts to extract and reason about material science information. The authors implemented a CoVe[434] (Consistency Verification) scheme to evaluate the model's answer and minimize hallucination. The authors showed that including CoVe increased the model's precision by $\sim 20\%$ when compared to previous methods such as an agent using ReAct only.[439]

Promoting molecular discovery is a topic with great attention in the literature devoted to it and, as described extensively above, LLMs have leveraged a large amount of unstructured data to accelerate that discovery. Janakarajan et al.[508] discuss the advantages of using LLMs in fields such as *de novo* drug design, reaction chemistry, and property prediction, but they augment the LLM in `IBM ChemChat`, a chatbot with the capability of using common APIs and python packages commonly used daily by a cheminformatics researcher to access molecular information. `ChemChat` has access to tools such as Generative Toolkit for Scientific Discovery (GT4SD),[541] a package with dozens of trained models generative models for science, rxn4chemistry,[542] a package for computing chemistry reactions tasks, HuggingMolecules,[543] a package developed to aggregate molecular property prediction LMs, and RDKit,[544] a package to manipulating molecules. Since `ChemChat` implements an agent in a chat-like environment, users can interactively refine design ideas. Despite being developed to target *de novo* drug design, `ChemChat` nonetheless is a multi-purpose platform that can be more broadly used for molecular discovery.

In addition to the capabilities described above, LLM-based agents can empower users to tackle tasks that typically require extensive technical knowledge. In previous work, Wellawatte and Schwaller[343] and Gandhi and White[290] showed that including natural language explanations (NLE) in explainable AI (XAI) analysis can improve user understanding. More recently, Wellawatte and Schwaller[343] developed XpertAI[343] to seamlessly integrate XAI techniques with LLMs to interpret and explain raw chemical data autonomously. Applying XAI techniques is usually restricted to technical experts but by integrating such techniques with an LLM-based agent to automate the workflow, the authors made XAI accessible to a wider audience.

Their system receives raw data with labels for physicochemical properties. The raw data is used to compute human-interpretable descriptors and then calculate SHAP (or SHapley Addictive exPlanations) values or Z-scores for Local Interpretable Model-agnostic Explanations (LIME). By calculating SHAP values, a value can be assigned to each feature, indicating its contribution to a model's output. LIME interprets a model by making a local approximation, around a particular prediction, to indicate what factors have contributed to that prediction in the model. It may use, for example, a surrogate local linear regression fit to recognized features.[290] In addition to XAI tools, XpertAI can search and leverage scientific literature to provide accessible natural language explanations (NLEs). While ChatGPT provides scientific justifications with similar accuracy, its explanation is often too broad. On the other hand, XpertAI provides data-specific explanations and visual XAI plots to support its explanations.[343] With a similar goal, Zheng et al.[545] prompted the LLM to generate explanatory rules from data.

These developments signify a growing trend in the integration of tools and LLMs in autonomous AI within scientific research. By automating routine tasks, enhancing information retrieval and analysis, and facilitating experimentation, AI is expanding the capabilities of researchers and accelerating the pace of scientific discovery. This review underscores the transformative impact of AI across various scientific domains, heralding a new era of innovation and efficiency in chemical research.

5.5   Agents for Hypothesis Creation

Following the agent's classification proposed by Gao et al.[420], the studies we have discussed previously lie mainly in level 1, i.e. AI agents as a research assistant. Therefore, such agents can support researchers in executing predefined tasks, but they lack the autonomy to propose, test, and refine new scientific hypotheses. New research has been focusing on making agents able to refine scientists' initial hypotheses collaboratively, which is a required skill to achieve level 2 in the Gao et al.[420] classification.

The idea of an "AI scientist" who can generate new, relevant research questions (RQ) has been pursued by groups such as Wang et al.[546], who developed a framework called Scientific Inspiration Machines Optimized for Novelty (SciMON). SciMON uses LLMs to produce new scientific ideas grounded in existing literature. It retrieves inspirations from past papers and iteratively refines generated ideas to optimize novelty by comparing them with prior work. Extending these ideas, Gu and Krenn[547] used LLMs to search over a knowledge graph for inspiration to propose new personalized research ideas. Aligned with this vision, Liu et al.[506] developed CoQuest, partially automating the brainstorming for new the RQ process. This system uses a human-computing interface (HCI) to allow the agent to create new RQs that can be further enhanced by human feedback. They developed two strategies for RQ generation: breadth-first, where the agent generates multiple RQs simultaneously following the original user's prompt, and depth-first, where

multiple RQs are created sequentially, building on the top of the previously generated RQ. For each RQ generation, the agent implements a ReAct[439] framework with tools for literature discovery, hypothesis proposition, refinement, and evaluation. Upon evaluation of 20 HCI doctoral researchers by a post-interaction survey, the breadth-first approach was preferred by 60% of the evaluators. Interestingly, despite the evaluators' report that the breadth-first approach gave them more control and resulted in more trustworthy RQs, the depth-first had better scores for novelty and surprise. This difference might be caused by the fact that the depth-first uses its own RQ to iterate. This process can introduce new keywords that users have not considered.

Focusing on generating and testing hypotheses, ChemReasoner[497] uses a domain-specific reward function and computational chemistry feedback to validate agent responses. The authors combined a Monte Carlo thought search[548] for catalysis with a reward function from atomistic GNNs trained to predict adsorption energy or reaction energy barriers. While the search is responsible for exploiting literature information and allowing the model to propose new materials, the hypothetic material is further tested by the GNN. This framework was applied to suggest materials for adsorbates, biofuel catalysts, and catalysts for CO_2 to methanol conversion. The LLM generated the top five catalysts for each task, with ChemReasoner significantly outperforming GPT-4 based on the reward score.

Similarly, Ma et al.[494] developed Scientific Generative Agent (SGA) to generate hypotheses and iteratively refine them through computational simulations. Initially, the LLM generates a hypothesis. In the use cases considered, it can be a code snippet or a molecule. In the sequence, a search algorithm is used to find a better initial hypothesis for solving the initial query. Finally, this hypothesis — code or molecule — is optimized using a gradient-based algorithm. Lastly, the optimization output serves as feedback to the LLM to iterate. In their molecule design task, the goal was to generate a molecule with a specified HOMO-LUMO gap. The hypothesis is a molecule, that is, a SMILES string and a set of atomic coordinates. The gap is predicted by employing UniMol.[549] They showed that SGA could generate molecules based on quantum mechanical properties, but the results were not validated.

6   Challenges and Opportunities

LLMs hold great potential in chemistry due to their ability to both predict properties, new molecules, and their syntheses and to orchestrate existing computational and experimental tools. These capabilities enhance the accuracy and efficiency of chemical research and open up new avenues for discovery and innovation. By encapsulating AI models, data analysis software, and laboratory equipment within agent-based frameworks, researchers can harness these sophisticated tools through a unified interface. This approach not only simplifies the interaction with complex systems but also democratizes the immense capabilities of modern computational tools, thereby maximizing their utility in advancing chemical research and development. Other publications and reviews also shared their opinions on the challenges for the future of LLMs and LLM-based agents in chemistry.[76,194,423,550–552] Nonetheless, some important challenges and opportunities for progress remain, which we summarize here.

Data Quality and Availability   Quality and availability of data are critical factors that influence the efficacy of LLMs. Indeed, scaling both the model size and the amount of training data used has proven to improve capabilities.[553] However, current AI models are not trained on large amounts of chemical data, which limits their capabilities to reason about advanced chemical concepts.[183]

There are two types of datasets commonly used to train LLMs: unlabeled and labeled datasets. Unlabeled datasets, or pretraining data, are used during the semi-supervised training, which focuses on creating a "prior belief" about a molecule. Currently, we have huge datasets composed of hypothetical and/or theoretical data. When a model is trained on data that is not grounded in real chemical information, this might cause the model to learn a wrong prior belief.[554]

Labeled datasets, often used in benchmarks, also suffer from their inclusion of hypothetical and calculated data. Benchmarks are necessary for quantifying improvements in AI modeling and prediction within a competitive field. However, dominant benchmarks like MoleculeNet,[56] have significant limitations that may restrict the generalizability and applicability of evolving models. In his blog, Walters[188] brings to light numerous errors and inconsistencies within the MoleculeNet data, which substantially impact model performance and reliability.[184–187] Walters also argues that the properties present in these benchmarks do not directly correlate with real chemistry improvement. As such, new benchmarks need to translate to practical chemistry problems directly. For instance, increasing accuracy in predicting LogP is not necessarily mapped to drugs with greater bioavailability. Some promising work has come from the Therapeutic Common Data (TDC),[190,191] which includes data from actual therapeutic essays, providing a more practical foundation for model training.

The community continues to work to organize and curate datasets to prepare data for LLM training and evaluation. Scientific benchmarks,[218,502,555] repositories with curated datasets,[182] and packages for model evaluation[183] have been developed. However, the challenges concerning grounded truth and consistent datasets remain. With advancements in

scientific document processing,[379] there is now the opportunity to obtain new datasets from peer-reviewed scientific papers.[197–200] Due to the multi-modal capabilities of such AI models, these new benchmarks can comprise multiple data types, potentially enhancing the applicability and transferability of these models. The continual curation of new, relevant datasets that represent the complexities of real-world chemical problems will further enhance the robustness and relevance of LLMs in chemistry.

Model Interpretability   Model interpretability is a significant challenge for LLMs due to their "black-box" nature, which obscures the understanding of how predictions are made. However, innovative approaches are being developed to enhance LLMs' interpretability. For instance, Schwaller et al.[556] and Schilter et al.[557] used information from the different multi-attention heads. While Schwaller et al.[556] connected atoms from reactants to atoms in the products, Schilter et al.[557] assigned H-NMR peaks to specific hydrogens in a molecule to indicate how spectra were comprehended, or structures deduced. Additionally, since the LLMs use language, which is intrinsically interpretable, LLMs may be incrementally modified to explain their reasoning processes directly, exemplified with tools like eXpertAI[343] and or simply adjusting prompting.[437,484] These methods address the critical need for transparency in the mechanism of understanding for a good prediction beyond the good prediction itself.

Integration with Domain Knowledge and Cross-disciplinary Applications   While LLMs excel at pattern recognition, integrating explicit chemical rules and domain knowledge into these systems remains challenging. This integration is essential to make predictions that are not only statistically valid but also chemically reasonable. It was shown by Beltagy et al.[220] and Gu et al.[215] that better performance on common NLP tasks can be achieved by developing a vocabulary and pretraining on a domain-specific training corpus. While pretraining with domain-specific datasets that include chemical properties, reaction mechanisms, and experimental results may better capture the nuances of chemistry, but using AI to foster multi-disciplinary research remains a significant challenge. The Galactica LLM[123] also used special tokens for delineating chemical information, to relatively good success on chemistry tasks. Aryal et al.[558] also progress by creating an ensemble of specialist agents with different domains of knowledge, allowing them to interact to better answer the user query. Specifically, Aryal et al.[558] used agents with chemistry, physics, electrochemistry, and materials knowledge.

Tool Development   The effectiveness of a combined LLM/autonomous agent approach hinges significantly on the availability and quality of the tools, as well as on the complexity and diversity of the chemical tasks at hand. Some emphasis should be placed on refining standalone tools, with the confidence that overarching frameworks, like a GPT-4-type wrapper, or "assistant", will eventually integrate these tools seamlessly. Developers should stay informed about existing tools and design their tools to interface effectively with such a wrapper. This ensures that each tool is ready to contribute its unique capabilities to a cohesive agent system.

Reinforcement Learning   RL has been successfully used in LLMs,[103,559,560] with a few applications also proposed for use in agents.[561,562] The next frontier is applying RL to agents directly, to improve their ability on specific tasks. Bou et al.[114] provided a recent framework and example for generative molecular design when viewed as an RL problem (similar to RLHF) and some early success has been seen in applying the RLHF algorithm directly to protein language models where the reward model comes from scientific tasks.[115] Neither of these are direct RL on language model agents, but are a step towards this goal.

Agent Evaluation   Comparing different agent systems is challenging due to the lack of robust benchmarks and evaluation schemes. Consequently, it is difficult to define what constitutes a "superhuman" digital chemist and reach a consensus on the criteria for success.[563,563,564] This issue is similar to the ongoing discussions about defining artificial general intelligence (AGI) and the expected capabilities of cognitive architectures.[565,566] Once a reliable metric for evaluating such AI systems is established, it is crucial for the AI scientific community to set clear guidelines for conducting research. Currently, assessing success is challenging because the goals are not well defined. Building on this, we propose using Bloom's taxonomy[404,405] as a reference point for developing a metric to evaluate more complex reasoning and tool use in autonomous agents. This educational framework categorizes cognitive skills in a hierarchical manner, from basic recall to creative construction, providing a structured approach to assess higher-order thinking and reasoning capabilities in these systems. This adaptation could significantly enhance the evaluation of LLMs and autonomous agents, especially when tackling complex chemical challenges.

Ethical and Safety Concerns   As with all AI technologies, deploying LLMs involves ethical considerations, such as biases in predictions and the potential misuse of AI-generated chemical knowledge. Ruan et al.[567] and Tang et al.[532] highlight the need for multi-level regulation, noting that current alignment methods may be insufficient for ensuring safety and that human evaluation alone is not scalable.

The absence of specialized models for risk control and reliable safety evaluations poses a significant challenge to ensuring the safety of tool-using LLMs. This highlights the urgent need to automate red-teaming strategies to reinforce AI safety protocols. Additionally, the development of safe AI systems should prioritize minimizing harmful hallucinations. While managing dual-use risks is a human responsibility and should be controlled through safety assessments at publication or indirect regulation by the scientific community.

Human-AI Collaboration in Chemical Research   LLMs are poised to transform fields such as drug discovery, materials science, and environmental chemistry due to their ability to predict chemical properties and reactions with remarkable accuracy. Models based on architectures like BERT have demonstrated their capability to achieve state-of-the-art performance in various property prediction tasks.[45,248] Furthermore, studies by Jablonka et al.[142] and Born and Manica[139] have showcased the predictive power of LLMs by reformulating traditional regression and classification tasks as generative tasks, opening up new avenues for chemical modeling. However, as emphasized by Weng[422], maintaining the reliability of LLM outputs is essential, as inaccuracies in formatting, logical reasoning, or content can significantly impede their practical utility. Hallucination is also an intrinsic issue with LLMs.[568] Though agents can deal with hallucinations to some extent by implementing sanity-checking tools, it does not make the response hallucination-proof. A possible approach to address this issue is to use a human-in-the-loop approach, where steps of human-agent interaction are added to the workflow to check if the agent is in the correct pathway to solve the request.[569–571]

The potential of LLMs to design novel molecules and materials was highlighted by the AI-powered robotic lab assistant, A-Lab, which synthesized 41 new materials within just 17 days.[572] Nonetheless, this achievement has sparked debates about the experimental methods and the actual integration of atoms into new crystalline materials, raising questions about the authenticity of the synthesized structures.[573] These controversies underline the necessity for rigorous standards and the critical role of human expertise in validating AI-generated results. Again, the integration of advanced AI tools with the oversight of seasoned chemists is crucial, suggesting that a hybrid approach could significantly enhance both the innovation and integrity of materials science research.

In parallel, we have seen how LLM-based agents are increasingly capable of automating routine tasks in chemical research, which traditionally consume significant time and resources. These models excel in real-time data processing, managing vast datasets, and even conducting comprehensive literature reviews with minimal human intervention. Advances in AI technology now allow agents not only to perform predefined tasks but also to adapt and develop new tools for automating additional processes. For instance, tasks such as data analysis, literature review, and elements of experimental design are now being automated.[47,48,343,502,505,507,574,575] This automation liberates chemists to focus on more innovative and intellectually engaging aspects of their work, and the opportunity is to expand productivity and creativity in their science.

Promotion of Impactful Discoveries   AI technologies offer experimental chemists significant opportunities to streamline repetitive tasks like data collection and analysis, freeing up time for innovation.[417] AI-powered tools can suggest novel experiments and chemical pathways,[533] but the black-box nature of many models raises concerns about trust and transparency. Human expertise remains essential to validate AI-generated results, especially in critical experiments.

A key challenge is translating AI predictions into real-world experiments, where factors like reagent quality and equipment limitations must be considered. To integrate AI effectively in the lab, stronger collaboration between computational and experimental chemists is essential, ensuring AI tools are practical and aligned with lab conditions.[576] Clear communication will help identify the most impactful AI advancements, ensuring tools address the real needs of experimentalists. AI's ability to explore new chemical spaces also offers exciting opportunities for discovery, allowing chemists to harness these insights while maintaining oversight for accuracy and reliability.

AI in Everyday Chemistry   In the near future, AI tools will become integral to the daily workflow of chemists, transforming how routine challenges are approached and resolved. Today's chemist may soon find themselves interacting directly with AI-driven systems, leveraging advanced simulations, literature analyses, and predictive models to accelerate discovery. While this may sound like a glimpse into the future, the reality is that such tools are already emerging, and their widespread adoption has likely already begun. To illustrate this transformation, we propose the following scenario, based largely on prior experience.[577]

A chemist working on synthesizing a challenging target molecule encounters suboptimal yields and an unexpected side product. Despite verifying solvent purity, reaction conditions, and ruling out possible causes such as steric hindrance, or leaving group viability, the issue remains unresolved. The chemist plans a comprehensive systematic study, varying the leaving group and adjusting the length of a a bulky alkyl chain in one of the secondary amines.[577] This would require

weeks of repeated testing and data analysis, creating two lengthy projects for PhD students, diverting 2-3 months of effort. Nonetheless, the starting materials are ordered, and are expected to arrive within a fortnight.

In contrast, another chemist, equipped with methods described here, approaches the problem differently. Through *in-silico* studies, they evaluate the chemical properties of reactants and intermediates using a selection from the chemistry-specific LLMs described above. This strategy allows for rapid hypothesis testing and simulation of reaction conditions. With a human-in-the-loop workflow, the chemist refines the predictions, dismissing implausible pathways and focusing on a promising hypothesis. They use tools like PaperQA2[492] to verify the reaction mechanism against existing literature, ensuring a solid foundation in prior knowledge. This AI-driven workflow enables the chemist to design three targeted experiments, each validating a critical model prediction, thus bypassing the need for a larger methodological studies. Using an automated ChemCrow system,[47] the required starting materials are synthesized overnight. The following day, a PhD chemist performs the reactions, swiftly confirming the AI-derived hypothesis and achieving the desired product within 24 hours. The entire process, from problem identification to successful synthesis, concludes in just one week. Meanwhile, the first group of PhD students continues their extensive exploration of reaction conditions, gaining methodological insights but without directly achieving their original goal.

This comparison underscores how creativity and efficiency in research may benefit from a hybrid approach where there is some computational heavy lifting, along with a team of virtual chemistry experts to help hone and test ideas.

7   Conclusions

Since this review is targeted in part to an audience of chemists, who may not have yet embraced AI technology, we consider it valuable to point out our perspective that AI in chemistry is definitely here to stay. We predict that its use will only grow as a necessary tool that will inevitably lead to more jobs and greater progress. We hope to facilitate the change by connecting the technology to the chemical problems that our readership is already addressing through more traditional methods.

Large Language Models (LLMs) have demonstrated remarkable potential in reshaping chemical research and development workflows. These models have facilitated significant advancements in molecular simulation, reaction prediction, and materials discovery. In this review, we discussed the evolution of LLMs in chemistry and biochemistry. Successful cases where LLMs have proven their potential in promoting scientific discovery were shown with caveats of such models.

Adopting LLM-based autonomous agents in chemistry has enhanced the accuracy and efficiency of traditional research methodologies and introduced innovative approaches to solving complex chemical problems. Looking forward, the continued integration of LLMs promises to accelerate the field's evolution further, driving forward the frontiers of scientific discovery and technological innovation in chemistry. We have shown how agents have been used in chemistry and proposed a framework for thinking about agents as a central LLM followed by interchangeable components.

However, despite the community's astonishing advances in this field, many challenges still require solutions. We identified the main challenges and opportunities that need to be addressed to promote the further development of agents in chemistry. Addressing the challenges related to model transparency, data biases, and computational demands will be crucial for maximizing their utility and ensuring their responsible use in future scientific endeavors.

While there are significant challenges to be addressed, the opportunities presented by LLMs in chemistry are vast and have the potential to fundamentally alter how chemical research and development are conducted. Effectively addressing these challenges will be crucial for realizing the full potential of LLMs in this exciting field. To keep pace with the ever-growing number of relevant publications, we will maintain a repository with an organized structure listing new studies regarding LLMs and LLM-based agents focused on scientific purposes. The repository can be found in `https://github.com/ur-whitelab/LLMs-in-science`

Author contribution

All authors contributed to writing this review article.

Competing interests

The authors have no conflicts to declare.

Acknowledgments

References

[1] Peter Willett. Chemoinformatics: a history. *Wiley Interdisciplinary Reviews: Computational Molecular Science*, 1(1):46–56, 2011.

[2] Edward J. Griffen, Alexander G. Dossetter, and Andrew G. Leach. Chemists: AI Is Here; Unite To Get the Benefits. *Journal of Medicinal Chemistry*, 63(16):8695–8704, August 2020. ISSN 0022-2623. doi:10.1021/acs.jmedchem.0c00163. URL https://doi.org/10.1021/acs.jmedchem.0c00163. Publisher: American Chemical Society.

[3] Zachary J. Baum, Xiang Yu, Philippe Y. Ayala, Yanan Zhao, Steven P. Watkins, and Qiongqiong Zhou. Artificial Intelligence in Chemistry: Current Trends and Future Directions. *Journal of Chemical Information and Modeling*, 61(7):3197–3212, July 2021. ISSN 1549-9596. doi:10.1021/acs.jcim.1c00619. URL https://doi.org/10.1021/acs.jcim.1c00619. Publisher: American Chemical Society.

[4] Lucas B. Ayres, Federico J.V. Gomez, Jeb R. Linton, Maria F. Silva, and Carlos D. Garcia. Taking the leap between analytical chemistry and artificial intelligence: A tutorial review. *Analytica Chimica Acta*, 1161:338403, May 2021. ISSN 00032670. doi:10.1016/j.aca.2021.338403. URL https://linkinghub.elsevier.com/retrieve/pii/S0003267021002294.

[5] Xin Yang, Yifei Wang, Ryan Byrne, Gisbert Schneider, and Shengyong Yang. Concepts of Artificial Intelligence for Computer-Assisted Drug Discovery. *Chemical Reviews*, 119(18):10520–10594, September 2019. ISSN 0009-2665. doi:10.1021/acs.chemrev.8b00728. URL https://doi.org/10.1021/acs.chemrev.8b00728. Publisher: American Chemical Society.

[6] Adam C Mater and Michelle L Coote. Deep learning in chemistry. *Journal of chemical information and modeling*, 59(6):2545–2559, 2019.

[7] Yun-Fei Shi, Zheng-Xin Yang, Sicong Ma, Pei-Lin Kang, Cheng Shang, P Hu, and Zhi-Pan Liu. Machine learning for chemistry: basics and applications. *Engineering*, 2023.

[8] John A Keith, Valentin Vassilev-Galindo, Bingqing Cheng, Stefan Chmiela, Michael Gastegger, Klaus-Robert Muller, and Alexandre Tkatchenko. Combining machine learning and computational chemistry for predictive insights into chemical systems. *Chemical reviews*, 121(16):9816–9872, 2021.

[9] David Kuntz and Angela K Wilson. Machine learning, artificial intelligence, and chemistry: How smart algorithms are reshaping simulation and the laboratory. *Pure and Applied Chemistry*, 94(8):1019–1054, 2022.

[10] Markus Meuwly. Machine learning for chemical reactions. *Chemical Reviews*, 121(16):10218–10239, 2021.

[11] Joshua Lederberg, Georgia L Sutherland, Bruce G Buchanan, Edward A Feigenbaum, Alexander V Robertson, Alan M Duffield, and Carl Djerassi. Applications of artificial intelligence for chemical inference. i. number of possible organic compounds. acyclic structures containing carbon, hydrogen, oxygen, and nitrogen. *Journal of the American Chemical Society*, 91(11):2973–2976, 1969.

[12] Robert K Lindsay, Bruce G Buchanan, Edward A Feigenbaum, and Joshua Lederberg. Dendral: a case study of the first expert system for scientific hypothesis formation. *Artificial intelligence*, 61(2):209–261, 1993.

[13] B. G. Buchanan, D. H. Smith, W. C. White, R. J. Gritter, E. A. Feigenbaum, J. Lederberg, and Carl Djerassi. Applications of artificial intelligence for chemical inference. 22. Automatic rule formation in mass spectrometry by means of the meta-DENDRAL program. *Journal of the American Chemical Society*, 98(20):6168–6178, September 1976. ISSN 0002-7863. doi:10.1021/ja00436a017. URL https://doi.org/10.1021/ja00436a017. Publisher: American Chemical Society.

[14] Corwin Hansch, Peyton P Maloney, Toshio Fujita, and Robert M Muir. Correlation of biological activity of phenoxyacetic acids with hammett substituent constants and partition coefficients. *Nature*, 194(4824):178–180, 1962.

[15] Corwin Hansch and Toshio Fujita. p-σ-π analysis. a method for the correlation of biological activity and chemical structure. *Journal of the American Chemical Society*, 86(8):1616–1626, 1964.

[16] Corwin Hansch and Albert Leo. *Exploring QSAR.: Fundamentals and applications in chemistry and biology*, volume 1. American Chemical Society, 1995.

[17] John G Topliss and Robert J Costello. Chance correlations in structure-activity studies using multiple regression analysis. *Journal of Medicinal Chemistry*, 15(10):1066–1068, 1972.

[18] John N. Weinstein, Kurt W. Kohn, Michael R. Grever, Vellarkad N. Viswanadhan, Lawrence V. Rubinstein, Anne P. Monks, Dominic A. Scudiero, Lester Welch, Antonis D. Koutsoukos, August J. Chiausa, and Kenneth D. Paull. Neural Computing in Cancer Drug Development: Predicting Mechanism of Action. *Science*, 258(5081): 447–451, October 1992. doi:10.1126/science.1411538. URL `https://www.science.org/doi/10.1126/science.1411538`. Publisher: American Association for the Advancement of Science.

[19] William W Van Osdol, Timothy G Myers, Kenneth D Paull, Kurt W Kohn, and John N Weinstein. Use of the kohonen self-organizing map to study the mechanisms of action of chemotherapeutic agents. *JNCI: Journal of the National Cancer Institute*, 86(24):1853–1859, 1994.

[20] Brian B Goldman and W Patrick Walters. Machine learning in computational chemistry. *Annual Reports in Computational Chemistry*, 2:127–140, 2006.

[21] DA Pereira and JA Williams. Origin and evolution of high throughput screening. *British journal of pharmacology*, 152(1):53–61, 2007.

[22] José L. Medina-Franco, Marc A. Giulianotti, Gregory S. Welmaker, and Richard A. Houghten. Shifting from the single to the multitarget paradigm in drug discovery. *Drug Discovery Today*, 18(9):495–501, May 2013. ISSN 1359-6446. doi:10.1016/j.drudis.2013.01.008. URL `https://www.sciencedirect.com/science/article/pii/S1359644613000251`.

[23] Keith T. Butler, Daniel W. Davies, Hugh Cartwright, Olexandr Isayev, and Aron Walsh. Machine learning for molecular and materials science. *Nature*, 559(7715):547–555, July 2018. ISSN 1476-4687. doi:10.1038/s41586-018-0337-2. URL `https://www.nature.com/articles/s41586-018-0337-2`. Publisher: Nature Publishing Group.

[24] Matthias Rupp, Alexandre Tkatchenko, Klaus-Robert Müller, and O. Anatole von Lilienfeld. Fast and Accurate Modeling of Molecular Atomization Energies with Machine Learning. *Physical Review Letters*, 108(5): 058301, January 2012. doi:10.1103/PhysRevLett.108.058301. URL `https://link.aps.org/doi/10.1103/PhysRevLett.108.058301`. Publisher: American Physical Society.

[25] Marcus Olivecrona, Thomas Blaschke, Ola Engkvist, and Hongming Chen. Molecular de-novo design through deep reinforcement learning. *J. Cheminform.*, 9(1):48, September 2017. ISSN 1758-2946,1758-2946. doi:10.1186/s13321-017-0235-x. URL `https://jcheminf.biomedcentral.com/articles/10.1186/s13321-017-0235-x`.

[26] Marwin H S Segler, Thierry Kogej, Christian Tyrchan, and Mark P Waller. Generating focused molecule libraries for drug discovery with recurrent neural networks. *ACS Cent. Sci.*, 4(1):120–131, January 2018. ISSN 2374-7943,2374-7951. doi:10.1021/acscentsci.7b00512. URL `https://pubs.acs.org/doi/abs/10.1021/acscentsci.7b00512`.

[27] Marwin H. S. Segler, Thierry Kogej, Christian Tyrchan, and Mark P. Waller. Generating Focused Molecule Libraries for Drug Discovery with Recurrent Neural Networks. *ACS Central Science*, 4(1):120–131, January 2018. ISSN 2374-7943. doi:10.1021/acscentsci.7b00512. URL `https://doi.org/10.1021/acscentsci.7b00512`. Publisher: American Chemical Society.

[28] Anvita Gupta, Alex T Muller, Berend J H Huisman, Jens A Fuchs, Petra Schneider, and Gisbert Schneider. Generative recurrent networks for de novo drug design. *Mol. Inform.*, 37(1-2):1700111, January 2018. ISSN 1868-1751,1868-1743. doi:10.1002/minf.201700111. URL `https://onlinelibrary.wiley.com/doi/abs/10.1002/minf.201700111`.

[29] Pavel Karpov, Guillaume Godin, and Igor V Tetko. Transformer-CNN: Swiss knife for QSAR modeling and interpretation. *J. Cheminform.*, 12(1):17, March 2020. ISSN 1758-2946,1758-2946. doi:10.1186/s13321-020-00423-w. URL `https://link.springer.com/articles/10.1186/s13321-020-00423-w`.

[30] K. T. Schütt, H. E. Sauceda, P.-J. Kindermans, A. Tkatchenko, and K.-R. Müller. SchNet – A deep learning architecture for molecules and materials. *The Journal of Chemical Physics*, 148(24):241722, March 2018. ISSN 0021-9606. doi:10.1063/1.5019779. URL `https://doi.org/10.1063/1.5019779`.

[31] Maya Hirohara, Yutaka Saito, Yuki Koda, Kengo Sato, and Yasubumi Sakakibara. Convolutional neural network based on SMILES representation of compounds for detecting chemical motif. *BMC Bioinformatics*, 19(Suppl 19):526, December 2018. ISSN 1471-2105,1471-2105. doi:10.1186/s12859-018-2523-5. URL `https://bmcbioinformatics.biomedcentral.com/articles/10.1186/s12859-018-2523-5`.

[32] Connor W Coley, Wengong Jin, Luke Rogers, Timothy F Jamison, Tommi S Jaakkola, William H Green, Regina Barzilay, and Klavs F Jensen. A graph-convolutional neural network model for the prediction of chemical reactivity. *Chem. Sci.*, 10(2):370–377, January 2019. ISSN 2041-6520,2041-6539. doi:10.1039/c8sc04228d. URL https://pubs.rsc.org/en/content/articlelanding/2019/sc/c8sc04228d.

[33] Vijay Prakash Dwivedi, Chaitanya K Joshi, Anh Tuan Luu, Thomas Laurent, Yoshua Bengio, and Xavier Bresson. Benchmarking graph neural networks. *Journal of Machine Learning Research*, 24(43):1–48, 2023.

[34] Benjamin Sanchez-Lengeling, Emily Reif, Adam Pearce, and Alexander B Wiltschko. A gentle introduction to graph neural networks. *Distill*, 6(9):e33, 2021.

[35] Michael M Bronstein, Joan Bruna, Yann LeCun, Arthur Szlam, and Pierre Vandergheynst. Geometric deep learning: going beyond euclidean data. *IEEE Signal Processing Magazine*, 34(4):18–42, 2017.

[36] Zonghan Wu, Shirui Pan, Fengwen Chen, Guodong Long, Chengqi Zhang, and S Yu Philip. A comprehensive survey on graph neural networks. *IEEE transactions on neural networks and learning systems*, 32(1):4–24, 2020.

[37] Justin Gilmer, Samuel S Schoenholz, Patrick F Riley, Oriol Vinyals, and George E Dahl. Neural message passing for quantum chemistry. In *International conference on machine learning*, pages 1263–1272. PMLR, 2017.

[38] Rafael Gómez-Bombarelli, Jennifer N. Wei, David Duvenaud, José Miguel Hernández-Lobato, Benjamín Sánchez-Lengeling, Dennis Sheberla, Jorge Aguilera-Iparraguirre, Timothy D. Hirzel, Ryan P. Adams, and Alán Aspuru-Guzik. Automatic Chemical Design Using a Data-Driven Continuous Representation of Molecules. *ACS Central Science*, 4(2):268–276, February 2018. ISSN 2374-7943. doi:10.1021/acscentsci.7b00572. URL https://doi.org/10.1021/acscentsci.7b00572. Publisher: American Chemical Society.

[39] Thomas Gaudelet, Ben Day, Arian R Jamasb, Jyothish Soman, Cristian Regep, Gertrude Liu, Jeremy BR Hayter, Richard Vickers, Charles Roberts, Jian Tang, et al. Utilizing graph machine learning within drug discovery and development. *Briefings in bioinformatics*, 22(6):bbab159, 2021.

[40] Kamal Choudhary, Brian DeCost, Chi Chen, Anubhav Jain, Francesca Tavazza, Ryan Cohn, Cheol Woo Park, Alok Choudhary, Ankit Agrawal, Simon J. L. Billinge, Elizabeth Holm, Shyue Ping Ong, and Chris Wolverton. Recent advances and applications of deep learning methods in materials science. *npj Computational Materials*, 8 (1):1–26, April 2022. ISSN 2057-3960. doi:10.1038/s41524-022-00734-6. URL https://www.nature.com/articles/s41524-022-00734-6. Publisher: Nature Publishing Group.

[41] Victor Fung, Jiaxin Zhang, Eric Juarez, and Bobby G Sumpter. Benchmarking graph neural networks for materials chemistry. *Npj Comput. Mater.*, 7(1):1–8, June 2021. ISSN 2057-3960,2057-3960. doi:10.1038/s41524-021-00554-0. URL https://www.nature.com/articles/s41524-021-00554-0.

[42] Patrick Reiser, Marlen Neubert, André Eberhard, Luca Torresi, Chen Zhou, Chen Shao, Houssam Metni, Clint van Hoesel, Henrik Schopmans, Timo Sommer, and Pascal Friederich. Graph neural networks for materials science and chemistry. *Commun. Mater.*, 3(1):93, November 2022. ISSN 2662-4443,2662-4443. doi:10.1038/s43246-022-00315-6. URL https://www.nature.com/articles/s43246-022-00315-6.

[43] David Weininger. SMILES, a chemical language and information system. 1. Introduction to methodology and encoding rules. *Journal of Chemical Information and Computer Sciences*, 28(1):31–36, February 1988. ISSN 0095-2338. doi:10.1021/ci00057a005. URL https://doi.org/10.1021/ci00057a005. Publisher: American Chemical Society.

[44] Seyone Chithrananda, Gabe Grand, and Bharath Ramsundar. ChemBERTa: Large-scale self-supervised pretraining for molecular property prediction. *arXiv [cs.LG]*, October 2020. doi:10.48550/arXiv.2010.09885. URL http://arxiv.org/abs/2010.09885.

[45] Juncai Li and Xiaofei Jiang. Mol-BERT: An Effective Molecular Representation with BERT for Molecular Property Prediction. *WIRELESS COMMUNICATIONS & MOBILE COMPUTING*, 2021, September 2021. ISSN 1530-8669. doi:10.1155/2021/7181815.

[46] Ye Wang, Honggang Zhao, Simone Sciabola, and Wenlu Wang. cMolGPT: A conditional generative Pre-Trained transformer for Target-Specific de novo molecular generation. *Molecules*, 28(11), May 2023. ISSN 1420-3049. doi:10.3390/molecules28114430. URL http://dx.doi.org/10.3390/molecules28114430.

[47] Andres M. Bran, Sam Cox, Oliver Schilter, Carlo Baldassari, Andrew D White, and Philippe Schwaller. Augmenting large language models with chemistry tools. *Nature Machine Intelligence*, pages 1–11, May 2024. ISSN 2522-5839, 2522-5839. doi:10.1038/s42256-024-00832-8. URL https://www.nature.com/articles/s42256-024-00832-8.

[48] Daniil A Boiko, Robert MacKnight, Ben Kline, and Gabe Gomes. Autonomous chemical research with large language models. *Nature*, 624(7992):570–578, December 2023. ISSN 0028-0836,1476-4687. doi:10.1038/s41586-023-06792-0. URL https://www.nature.com/articles/s41586-023-06792-0.

[49] Andrew D White. The future of chemistry is language. *Nature Reviews Chemistry*, 7(7):457–458, July 2023. ISSN 2397-3358. doi:10.1038/s41570-023-00502-0. URL `https://www.nature.com/articles/s41570-023-00502-0`.

[50] Christopher J. Collison, Marc J. O'Donnell, and Jessica L. Alexander. Complexation between Rhodamine 101 and Single-Walled Carbon Nanotubes Indicative of Solvent-Nanotube Interaction Strength. *The Journal of Physical Chemistry C*, 112(39):15144–15150, October 2008. ISSN 1932-7447. doi:10.1021/jp804359j. URL `https://doi.org/10.1021/jp804359j`. Publisher: American Chemical Society.

[51] Tyler J. Wiegand, Juan S. Sandoval, Jeremy A. Cody, David W. McCamant, and Christopher J. Collison. Directional Exciton Diffusion, Measured by Subpicosecond Transient Absorption as an Explanation for Squaraine Solar Cell Performance. *The Journal of Physical Chemistry C*, 128(11):4616–4630, March 2024. ISSN 1932-7447, 1932-7455. doi:10.1021/acs.jpcc.3c06361. URL `https://pubs.acs.org/doi/10.1021/acs.jpcc.3c06361`.

[52] Rashad Ahmadov, Shane S. Michtavy, and Marc D. Porosoff. Dual Functional Materials: At the Interface of Catalysis and Separations. *Langmuir*, 40(19):9833–9841, May 2024. ISSN 0743-7463. doi:10.1021/acs.langmuir.3c03888. URL `https://doi.org/10.1021/acs.langmuir.3c03888`. Publisher: American Chemical Society.

[53] Thomas Fischer, Silvia Gazzola, and Rainer Riedl. Approaching Target Selectivity by De Novo Drug Design. *Expert Opinion on Drug Discovery*, 14(8):791–803, August 2019. ISSN 1746-0441. doi:10.1080/17460441.2019.1615435. URL `https://doi.org/10.1080/17460441.2019.1615435`. Publisher: Taylor & Francis _eprint: https://doi.org/10.1080/17460441.2019.1615435.

[54] Zhuo Wang, Zhehao Sun, Hang Yin, Xinghui Liu, Jinlan Wang, Haitao Zhao, Cheng Heng Pang, Tao Wu, Shuzhou Li, Zongyou Yin, and Xue-Feng Yu. Data-Driven Materials Innovation and Applications. *Advanced Materials*, 34(36):2104113, 2022. ISSN 1521-4095. doi:10.1002/adma.202104113. URL `https://onlinelibrary.wiley.com/doi/abs/10.1002/adma.202104113`. _eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/adma.202104113.

[55] Bhuvanesh Sridharan, Manan Goel, and U. Deva Priyakumar. Modern machine learning for tackling inverse problems in chemistry: molecular design to realization. *Chemical Communications*, 58(35):5316–5331, April 2022. ISSN 1364-548X. doi:10.1039/D1CC07035E. URL `https://pubs.rsc.org/en/content/articlelanding/2022/cc/d1cc07035e`. Publisher: The Royal Society of Chemistry.

[56] Zhenqin Wu, Bharath Ramsundar, Evan N. Feinberg, Joseph Gomes, Caleb Geniesse, Aneesh S. Pappu, Karl Leswing, and Vijay Pande. MoleculeNet: a benchmark for molecular machine learning. *Chemical Science*, 9(2):513–530, 2018. ISSN 2041-6520, 2041-6539. doi:10.1039/C7SC02664A. URL `http://xlink.rsc.org/?DOI=C7SC02664A`.

[57] Guillermo Restrepo. Chemical space: limits, evolution and modelling of an object bigger than our universal library. *Digital Discovery*, 1(5):568–585, October 2022. ISSN 2635-098X. doi:10.1039/D2DD00030J. URL `https://pubs.rsc.org/en/content/articlelanding/2022/dd/d2dd00030j`. Publisher: RSC.

[58] Peter Kirkpatrick and Clare Ellis. Chemical space. *Nature*, 432(7019):823–824, 2004.

[59] Asher Mullard et al. The drug-maker's guide to the galaxy. *Nature*, 549(7673):445–447, 2017.

[60] Eugenio J Llanos, Wilmer Leal, Duc H Luu, Jürgen Jost, Peter F Stadler, and Guillermo Restrepo. Exploration of the chemical space and its three historical regimes. *Proceedings of the National Academy of Sciences*, 116(26):12660–12665, 2019.

[61] Joshua Schrier, Alexander J. Norquist, Tonio Buonassisi, and Jakoah Brgoch. In Pursuit of the Exceptional: Research Directions for Machine Learning in Chemical and Materials Science. *Journal of the American Chemical Society*, 145(40):21699–21716, October 2023. ISSN 0002-7863. doi:10.1021/jacs.3c04783. URL `https://doi.org/10.1021/jacs.3c04783`. Publisher: American Chemical Society.

[62] Piotr S Gromski, Alon B Henson, Jarosław M Granda, and Leroy Cronin. How to explore chemical space using algorithms and automation. *Nature Reviews Chemistry*, 3(2):119–128, 2019.

[63] Sebastian Steiner, Jakob Wolf, Stefan Glatzel, Anna Andreou, Jarosław M Granda, Graham Keenan, Trevor Hinkley, Gerardo Aragon-Camarasa, Philip J Kitson, Davide Angelone, et al. Organic synthesis in a modular robotic system driven by a chemical programming language. *Science*, 363(6423):eaav2211, 2019.

[64] Benjamin Burger, Phillip M Maffettone, Vladimir V Gusev, Catherine M Aitchison, Yang Bai, Xiaoyan Wang, Xiaobo Li, Ben M Alston, Buyi Li, Rob Clowes, et al. A mobile robotic chemist. *Nature*, 583(7815):237–241, 2020.

[65] Benjamin P MacLeod, Fraser GL Parlane, Thomas D Morrissey, Florian Häse, Loïc M Roch, Kevan E Dettelbach, Raphaell Moreira, Lars PE Yunker, Michael B Rooney, Joseph R Deeth, et al. Self-driving laboratory for accelerated discovery of thin-film materials. *Science Advances*, 6(20):eaaz8867, 2020.

[66] Qiang Zhang, Keyang Ding, Tianwen Lyv, Xinda Wang, Qingyu Yin, Yiwen Zhang, Jing Yu, Yuhao Wang, Xiaotong Li, Zhuoyi Xiang, Xiang Zhuang, Zeyuan Wang, Ming Qin, Mengyao Zhang, Jinlu Zhang, Jiyu Cui, Renjun Xu, Hongyang Chen, Xiaohui Fan, Huabin Xing, and Huajun Chen. Scientific large language models: A survey on biological & chemical domains. *arXiv [cs.CL]*, January 2024. URL http://arxiv.org/abs/2401.14656.

[67] David E Rumelhart, Geoffrey E Hinton, and Ronald J Williams. Learning internal representations by error propagation, parallel distributed processing, explorations in the microstructure of cognition, ed. de rumelhart and j. mcclelland. vol. 1. 1986. *Biometrika*, 71:599–607, 1986.

[68] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.

[69] Antônio H. Ribeiro, Koen Tiels, Luis A. Aguirre, and Thomas B. Schön. Beyond exploding and vanishing gradients: analysing RNN training using attractors and smoothness, March 2020. URL http://arxiv.org/abs/1906.08482. arXiv:1906.08482 [cs, math, stat].

[70] Dr Barak Or. The Exploding and Vanishing Gradients Problem in Time Series, December 2023. URL https://medium.com/metaor-artificial-intelligence/the-exploding-and-vanishing-gradients-problem-in-time-series-6b87d558d22.

[71] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. *arXiv [cs.CL]*, June 2017. URL http://arxiv.org/abs/1706.03762.

[72] Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv [cs.LG]*, December 2023. URL http://arxiv.org/abs/2312.00752.

[73] Samy Jelassi, David Brandfonbrener, Sham M Kakade, and Eran Malach. Repeat after me: Transformers are better than state space models at copying. *arXiv [cs.LG]*, February 2024. URL http://arxiv.org/abs/2402.01032.

[74] Bo Peng, Eric Alcaide, Quentin Anthony, Alon Albalak, Samuel Arcadinho, Huanqi Cao, Xin Cheng, Michael Chung, Matteo Grella, Kranthi Kiran GV, et al. Rwkv: Reinventing rnns for the transformer era. *arXiv preprint arXiv:2305.13048*, 2023.

[75] Maximilian Beck, Korbinian Pöppel, Markus Spanring, Andreas Auer, Oleksandra Prudnikova, Michael Kopp, Günter Klambauer, Johannes Brandstetter, and Sepp Hochreiter. xLSTM: Extended Long Short-Term Memory, May 2024. URL http://arxiv.org/abs/2405.04517. arXiv:2405.04517 [cs, stat].

[76] Shervin Minaee, Tomas Mikolov, Narjes Nikzad, Meysam Chenaghlu, Richard Socher, Xavier Amatriain, and Jianfeng Gao. Large language models: A survey. *arXiv [cs.CL]*, February 2024. URL http://arxiv.org/abs/2402.06196.

[77] The Annotated Transformer, 2022. URL https://nlp.seas.harvard.edu/annotated-transformer/.

[78] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural Machine Translation by Jointly Learning to Align and Translate, May 2016. URL http://arxiv.org/abs/1409.0473. arXiv:1409.0473 [cs, stat].

[79] Tianle Li, Ge Zhang, Quy Duc Do, Xiang Yue, and Wenhu Chen. Long-context LLMs struggle with long in-context learning. *arXiv [cs.CL]*, April 2024. URL http://arxiv.org/abs/2404.02060.

[80] Yusen Zhang, Ruoxi Sun, Yanfei Chen, Tomas Pfister, Rui Zhang, and Sercan Ö Arik. Chain of agents: Large language models collaborating on long-context tasks. *arXiv [cs.CL]*, June 2024. URL http://arxiv.org/abs/2406.02818.

[81] Taku Kudo. Subword regularization: Improving neural network translation models with multiple subword candidates. *arXiv [cs.CL]*, April 2018. URL http://arxiv.org/abs/1804.10959.

[82] Taku Kudo and John Richardson. SentencePiece: A simple and language independent subword tokenizer and detokenizer for neural text processing. *arXiv [cs.CL]*, August 2018. URL http://arxiv.org/abs/1808.06226.

[83] Xinying Song, Alex Salcianu, Yang Song, Dave Dopson, and Denny Zhou. Fast WordPiece tokenization. *arXiv [cs.CL]*, December 2020. URL http://arxiv.org/abs/2012.15524.

[84] Phillip Rust, Jonas Pfeiffer, Ivan Vulić, Sebastian Ruder, and Iryna Gurevych. How good is your tokenizer? on the monolingual performance of multilingual language models. *arXiv [cs.CL]*, December 2020. URL http://arxiv.org/abs/2012.15613.

[85] Martin Berglund and Brink van der Merwe. Formalizing BPE tokenization. *arXiv [cs.FL]*, September 2023. URL http://arxiv.org/abs/2309.08715.

[86] Jonas Gehring, Michael Auli, David Grangier, Denis Yarats, and Yann N Dauphin. Convolutional sequence to sequence learning. *arXiv [cs.CL]*, pages 1243–1252, May 2017. URL https://proceedings.mlr.press/v70/gehring17a.html.

[87] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional Transformers for language understanding. *arXiv [cs.CL]*, October 2018. doi:10.48550/arXiv.1810.04805. URL http://arxiv.org/abs/1810.04805.

[88] Vinod Nair and Geoffrey E Hinton. Rectified linear units improve restricted boltzmann machines. In *Proceedings of the 27th international conference on machine learning (ICML-10)*, pages 807–814, 2010.

[89] Noam Shazeer. GLU Variants Improve Transformer, February 2020. URL http://arxiv.org/abs/2002.05202. arXiv:2002.05202 [cs, stat].

[90] Dan Hendrycks and Kevin Gimpel. Gaussian Error Linear Units (GELUs), June 2023. URL http://arxiv.org/abs/1606.08415. arXiv:1606.08415 [cs].

[91] Siru Ouyang, Zhuosheng Zhang, Bing Yan, Xuan Liu, Jiawei Han, and Lianhui Qin. Structured chemistry reasoning with large language models. *arXiv [cs.CL]*, November 2023. URL http://arxiv.org/abs/2311.09656.

[92] Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F Christiano. Learning to summarize with human feedback. *Advances in Neural Information Processing Systems*, 33:3008–3021, 2020.

[93] Tom B Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language Models are Few-Shot Learners. *arXiv [cs.CL]*, May 2020. URL http://arxiv.org/abs/2005.14165.

[94] Federico Errica, Giuseppe Siracusano, Davide Sanvito, and Roberto Bifulco. What did I do wrong? quantifying LLMs' sensitivity and consistency to prompt engineering. *arXiv [cs.LG]*, June 2024. URL http://arxiv.org/abs/2406.12334.

[95] Tianhao Shen, Renren Jin, Yufei Huang, Chuang Liu, Weilong Dong, Zishan Guo, Xinwei Wu, Yan Liu, and Deyi Xiong. Large language model alignment: A survey. *arXiv [cs.CL]*, September 2023. URL http://arxiv.org/abs/2309.15025.

[96] Changho Lee, Janghoon Han, Seonghyeon Ye, Stanley Jungkyu Choi, Honglak Lee, and Kyunghoon Bae. Instruction matters, a simple yet effective task selection approach in instruction tuning for specific tasks. *arXiv [cs.CL]*, April 2024. URL http://arxiv.org/abs/2404.16418.

[97] John Hewitt, Nelson F Liu, Percy Liang, and Christopher D Manning. Instruction following without instruction tuning. *arXiv [cs.CL]*, September 2024. URL http://arxiv.org/abs/2409.14254.

[98] Shengyu Zhang, Linfeng Dong, Xiaoya Li, Sen Zhang, Xiaofei Sun, Shuhe Wang, Jiwei Li, Runyi Hu, Tianwei Zhang, Fei Wu, and Guoyin Wang. Instruction tuning for large language models: A survey. *arXiv [cs.CL]*, August 2023. URL http://arxiv.org/abs/2308.10792.

[99] Yan Duan, John Schulman, Xi Chen, Peter L Bartlett, Ilya Sutskever, and Pieter Abbeel. RL^2: Fast reinforcement learning via slow reinforcement learning. *arXiv [cs.AI]*, November 2016. URL http://arxiv.org/abs/1611.02779.

[100] Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. Fine-tuning language models from human preferences. *arXiv [cs.CL]*, September 2019. URL http://arxiv.org/abs/1909.08593.

[101] Eyal Mazuz, Guy Shtar, Bracha Shapira, and Lior Rokach. Molecule generation using transformers and policy gradient reinforcement learning. *Sci. Rep.*, 13(1):8799, May 2023. ISSN 2045-2322. doi:10.1038/s41598-023-35648-w. URL http://dx.doi.org/10.1038/s41598-023-35648-w.

[102] Michael Laskin, Luyu Wang, Junhyuk Oh, Emilio Parisotto, Stephen Spencer, Richie Steigerwald, D J Strouse, Steven Hansen, Angelos Filos, Ethan Brooks, Maxime Gazeau, Himanshu Sahni, Satinder Singh, and Volodymyr Mnih. In-context reinforcement learning with algorithm distillation. *arXiv [cs.LG]*, October 2022. URL `http://arxiv.org/abs/2210.14215`.

[103] Aligning language models to follow instructions. `https://openai.com/index/instruction-following`, 2022. URL `https://openai.com/index/instruction-following`. Accessed: 2024-5-1.

[104] Sungdong Kim, Sanghwan Bae, Jamin Shin, Soyoung Kang, Donghyun Kwak, Kang Min Yoo, and Minjoon Seo. Aligning large language models through synthetic feedback. *arXiv [cs.CL]*, May 2023. URL `http://arxiv.org/abs/2305.13735`.

[105] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv [cs.LG]*, July 2017. URL `http://arxiv.org/abs/1707.06347`.

[106] Junzi Zhang, Jongho Kim, Brendan O'Donoghue, and Stephen Boyd. Sample efficient reinforcement learning with REINFORCE. *arXiv [cs.LG]*, October 2020. URL `http://arxiv.org/abs/2010.11364`.

[107] Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *arXiv [cs.AI]*, March 2023. URL `http://arxiv.org/abs/2303.11366`.

[108] Afra Feyza Akyurek, Ekin Akyurek, Aman Madaan, Ashwin Kalyan, Peter Clark, Derry Wijaya, and Niket Tandon. RL4F: Generating natural language feedback with reinforcement learning for repairing model outputs. *arXiv [cs.CL]*, May 2023. URL `http://arxiv.org/abs/2305.08844`.

[109] Yuji Cao, Huan Zhao, Yuheng Cheng, Ting Shu, Guolong Liu, Gaoqi Liang, Junhua Zhao, and Yun Li. Survey on large language model-enhanced reinforcement learning: Concept, taxonomy, and methods. *arXiv [cs.LG]*, March 2024. URL `http://arxiv.org/abs/2404.00282`.

[110] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *arXiv [cs.LG]*, May 2023. URL `http://arxiv.org/abs/2305.18290`.

[111] Zheng Yuan, Hongyi Yuan, Chuanqi Tan, Wei Wang, Songfang Huang, and Fei Huang. RRHF: Rank responses to align language models with human feedback without tears. *arXiv [cs.CL]*, April 2023. URL `http://arxiv.org/abs/2304.05302`.

[112] Feifan Song, Bowen Yu, Minghao Li, Haiyang Yu, Fei Huang, Yongbin Li, and Houfeng Wang. Preference ranking optimization for human alignment. *arXiv [cs.CL]*, June 2023. URL `http://arxiv.org/abs/2306.17492`.

[113] Shusheng Xu, Wei Fu, Jiaxuan Gao, Wenjie Ye, Weilin Liu, Zhiyu Mei, Guangju Wang, Chao Yu, and Yi Wu. Is DPO superior to PPO for LLM alignment? a comprehensive study. *arXiv [cs.CL]*, April 2024. URL `http://arxiv.org/abs/2404.10719`.

[114] Albert Bou, Morgan Thomas, Sebastian Dittert, Carles Navarro Ramírez, Maciej Majewski, Ye Wang, Shivam Patel, Gary Tresadern, Mazen Ahmad, Vincent Moens, et al. Acegen: Reinforcement learning of generative chemical agents for drug discovery. *arXiv preprint arXiv:2405.04657*, 2024.

[115] Thomas Hayes, Roshan Rao, Halil Akin, Nicholas James Sofroniew, Deniz Oktay, Zeming Lin, Robert Verkuil, Vincent Quy Tran, Jonathan Deaton, Marius Wiggert, Rohil Badkundri, Irhum Shafkat, Jun Gong, Alexander Derry, Raul Santiago Molina, Neil Thomas, Yousuf Khan, Chetan Mishra, Carolyn Kim, Liam J. Bartie, Patrick D. Hsu, Tom Sercu, Salvatore Candido, and Alexander Rives. Simulating 500 million years of evolution with a language model. *self-hosted preprint*, 2024. Preprint.

[116] Qizhi Pei, Wei Zhang, Jinhua Zhu, Kehan Wu, Kaiyuan Gao, Lijun Wu, Yingce Xia, and Rui Yan. BioT5: Enriching cross-modal integration in biology with chemical knowledge and natural language associations. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 1102–1123, Stroudsburg, PA, USA, December 2023. Association for Computational Linguistics. doi:10.18653/v1/2023.emnlp-main.70. URL `https://aclanthology.org/2023.emnlp-main.70`.

[117] Qizhi Pei, Lijun Wu, Kaiyuan Gao, Xiaozhuan Liang, Yin Fang, Jinhua Zhu, Shufang Xie, Tao Qin, and Rui Yan. BioT5+: Towards generalized biological understanding with IUPAC integration and multi-task tuning. *arXiv [q-bio.QM]*, February 2024. URL `http://arxiv.org/abs/2402.17810`.

[118] Juncai Li and Xiaofei Jiang. Mol-BERT: An effective molecular representation with BERT for molecular property prediction. *Proc. Int. Wirel. Commun. Mob. Comput. Conf.*, 2021, September 2021. ISSN 1530-8669, 1530-8669. doi:10.1155/2021/7181815. URL `https://www.hindawi.com/journals/wcmc/2021/7181815/`.

[119] Chen Qian, Huayi Tang, Zhirui Yang, Hong Liang, and Yong Liu. Can large language models empower molecular property prediction? July 2023. URL `http://arxiv.org/abs/2307.07443`.

[120] Thanh-Hoang Nguyen-Vo, Quang H. Trinh, Loc Nguyen, Trang T. T. Do, Matthew Chin Heng Chua, and Binh P. Nguyen. Predicting Antimalarial Activity in Natural Products Using Pretrained Bidirectional Encoder Representations from Transformers. *JOURNAL OF CHEMICAL INFORMATION AND MODELING*, 62(21): 5050–5058, November 2022. ISSN 1549-9596. doi:10.1021/acs.jcim.1c00584.

[121] Qiao Jin, Bhuwan Dhingra, Zhengping Liu, William W Cohen, and Xinghua Lu. PubMedQA: A dataset for biomedical research question answering. *arXiv [cs.CL]*, September 2019. URL `http://arxiv.org/abs/1909.06146`.

[122] Walid Ahmad, Elana Simon, Seyone Chithrananda, Gabriel Grand, and Bharath Ramsundar. ChemBERTa-2: Towards Chemical Foundation Models, September 2022. URL `http://arxiv.org/abs/2209.01712`. arXiv:2209.01712 [cs, q-bio].

[123] Ross Taylor, Marcin Kardas, Guillem Cucurull, Thomas Scialom, Anthony Hartshorn, Elvis Saravia, Andrew Poulton, Viktor Kerkez, and Robert Stojnic. Galactica: A large language model for science. *arXiv [cs.CL]*, November 2022. URL `http://arxiv.org/abs/2211.09085`.

[124] Wilson L. Taylor. "Cloze Procedure": A New Tool for Measuring Readability. *Journalism Quarterly*, 30(4): 415–433, September 1953. ISSN 0022-5533. doi:10.1177/107769905303000401. URL `https://doi.org/10.1177/107769905303000401`. Publisher: SAGE Publications.

[125] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. RoBERTa: A Robustly Optimized BERT Pretraining Approach, July 2019. URL `http://arxiv.org/abs/1907.11692`. arXiv:1907.11692 [cs].

[126] Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. Improving language understanding with unsupervised learning. 2018.

[127] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Ves Stoyanov, and Luke Zettlemoyer. BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension, October 2019. URL `http://arxiv.org/abs/1910.13461`. arXiv:1910.13461 [cs, stat].

[128] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *arXiv [cs.LG]*, October 2019. URL `http://arxiv.org/abs/1910.10683`.

[129] Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Yunxuan Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, Albert Webson, Shixiang Shane Gu, Zhuyun Dai, Mirac Suzgun, Xinyun Chen, Aakanksha Chowdhery, Alex Castro-Ros, Marie Pellat, Kevin Robinson, Dasha Valter, Sharan Narang, Gaurav Mishra, Adams Yu, Vincent Zhao, Yanping Huang, Andrew Dai, Hongkun Yu, Slav Petrov, Ed H Chi, Jeff Dean, Jacob Devlin, Adam Roberts, Denny Zhou, Quoc V Le, and Jason Wei. Scaling instruction-finetuned language models. *arXiv [cs.LG]*, October 2022. URL `http://arxiv.org/abs/2210.11416`.

[130] Bowen Tan, Yun Zhu, Lijuan Liu, Eric Xing, Zhiting Hu, and Jindong Chen. Cappy: Outperforming and boosting large multi-task LMs with a small scorer. *arXiv [cs.LG]*, November 2023. doi:10.48550/arXiv.2311.06720. URL `https://proceedings.neurips.cc/paper_files/paper/2023/file/b860c0c546f4a3a786f9c9468228c99f-Paper-Conference.pdf`.

[131] Junhong Shen, Neil Tenenholtz, James Brian Hall, David Alvarez-Melis, and Nicolo Fusi. Tag-LLM: Re-purposing general-purpose LLMs for specialized domains. *arXiv [cs.LG]*, February 2024. URL `http://arxiv.org/abs/2402.05140`.

[132] Guijin Son, Sangwon Baek, Sangdae Nam, Ilgyun Jeong, and Seungone Kim. Multi-task inference: Can large language models follow multiple instructions at once? *arXiv [cs.CL]*, February 2024. URL `http://arxiv.org/abs/2402.11597`.

[133] Wenfeng Feng, Chuzhan Hao, Yuewei Zhang, Yu Han, and Hao Wang. Mixture-of-LoRAs: An efficient multitask tuning for large language models. *arXiv [cs.CL]*, March 2024. URL `http://arxiv.org/abs/2403.03432`.

[134] Fuyu-8B: A multimodal architecture for AI agents. `https://www.adept.ai/blog/fuyu-8b`, 2023. URL `https://www.adept.ai/blog/fuyu-8b`. Accessed: 2023-11-8.

[135] Shengqiong Wu, Hao Fei, Leigang Qu, Wei Ji, and Tat-Seng Chua. NExT-GPT: Any-to-any multimodal LLM. *arXiv [cs.AI]*, September 2023. URL `http://arxiv.org/abs/2309.05519`.

[136] Debjyoti Bhattacharya, Harrison Cassady, Michael Hickner, and Wesley Reinhart. Large language models as molecular design engines. *ChemRxiv*, May 2024. doi:10.26434/chemrxiv-2024-n0l8q. URL https://chemrxiv.org/engage/api-gateway/chemrxiv/assets/orp/resource/item/664b795e418a5379b0d12460/original/large-language-models-as-molecular-design-engines.pdf.

[137] Mantas Vaškevičius, Jurgita Kapočiūtė-Dzikienė, and Liudas Šlepikas. Generative LLMs in organic chemistry: Transforming esterification reactions into natural language procedures. *Appl. Sci. (Basel)*, 13(24):13140, December 2023. ISSN 2076-3417,2076-3417. doi:10.3390/app132413140. URL https://www.mdpi.com/2076-3417/13/24/13140.

[138] Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust speech recognition via large-scale weak supervision. *arXiv [eess.AS]*, December 2022. URL http://arxiv.org/abs/2212.04356.

[139] Jannis Born and Matteo Manica. Regression Transformer enables concurrent sequence regression and generation for molecular language modelling. *Nat. Mach. Intell.*, 5(4):432–444, April 2023. ISSN 2522-5839,2522-5839. doi:10.1038/s42256-023-00639-z. URL https://www.nature.com/articles/s42256-023-00639-z.

[140] Jiashun Mao, Jianmin Wang, Kwang-Hwi Cho, and Kyoung Tai No. iupacGPT: IUPAC-based large-scale molecular pre-trained model for property prediction and molecule generation. *ChemRxiv*, pages 1–13, May 2023. doi:10.26434/chemrxiv-2023-5kjvh. URL https://chemrxiv.org/engage/chemrxiv/public-dashboard.

[141] Nima Shoghi, Adeesh Kolluru, John R. Kitchin, Zachary W. Ulissi, C. Lawrence Zitnick, and Brandon M. Wood. From Molecules to Materials: Pre-training Large Generalizable Models for Atomic Property Prediction, May 2024. URL http://arxiv.org/abs/2310.16802. arXiv:2310.16802 [cs].

[142] Kevin Maik Jablonka, Philippe Schwaller, Andres Ortega-Guerrero, and Berend Smit. Leveraging large language models for predictive chemistry. *Nat. Mach. Intell.*, 6(2):161–169, February 2024. ISSN 2522-5839,2522-5839. doi:10.1038/s42256-023-00788-1. URL https://www.nature.com/articles/s42256-023-00788-1.

[143] Ryan Jacobs, Maciej P Polak, Lane E Schultz, Hamed Mahdavi, Vasant Honavar, and Dane Morgan. Regression with large language models for materials and molecular property prediction. *arXiv [cond-mat.mtrl-sci]*, September 2024. URL http://arxiv.org/abs/2409.06080.

[144] Yu-Chen Lo, Stefano E Rensi, Wen Torng, and Russ B Altman. Machine learning in chemoinformatics and drug discovery. *Drug discovery today*, 23(8):1538–1546, 2018.

[145] Laurianne David, Amol Thakkar, Rocío Mercado, and Ola Engkvist. Molecular representations in AI-driven drug discovery: a review and practical guide. *Journal of Cheminformatics*, 12(1):56, September 2020. ISSN 1758-2946. doi:10.1186/s13321-020-00460-5. URL https://doi.org/10.1186/s13321-020-00460-5.

[146] Kenneth Atz, Francesca Grisoni, and Gisbert Schneider. Geometric deep learning on molecular representations. *arXiv [physics.chem-ph]*, July 2021. URL https://www.nature.com/articles/s42256-021-00418-8.

[147] W Patrick Walters and Regina Barzilay. Applications of deep learning in molecule generation and molecular property prediction. *Acc. Chem. Res.*, 54(2):263–270, January 2021. ISSN 0001-4842,1520-4898. doi:10.1021/acs.accounts.0c00699. URL https://pubs.acs.org/doi/abs/10.1021/acs.accounts.0c00699.

[148] Akshaya Karthikeyan and U Deva Priyakumar. Artificial intelligence: machine learning for chemical sciences. *J. Chem. Sci. (Bangalore)*, 134(1):2, 2022. ISSN 0974-3626,0973-7103. doi:10.1007/s12039-021-01995-2. URL https://link.springer.com/article/10.1007/s12039-021-01995-2.

[149] Zhen Li, Mingjian Jiang, Shuang Wang, and Shugang Zhang. Deep learning methods for molecular representation and property prediction. *Drug Discov. Today*, 27(12):103373, December 2022. ISSN 1359-6446,1878-5832. doi:10.1016/j.drudis.2022.103373. URL http://dx.doi.org/10.1016/j.drudis.2022.103373.

[150] Chi Chen, Weike Ye, Yunxing Zuo, Chen Zheng, and Shyue Ping Ong. Graph networks as a universal machine learning framework for molecules and crystals. *Chemistry of Materials*, 31(9):3564–3572, 2019.

[151] Weihua Hu, Matthias Fey, Marinka Zitnik, Yuxiao Dong, Hongyu Ren, Bowen Liu, Michele Catasta, and Jure Leskovec. Open graph benchmark: Datasets for machine learning on graphs. *Advances in neural information processing systems*, 33:22118–22133, 2020.

[152] Steven Kearnes, Kevin McCloskey, Marc Berndl, Vijay Pande, and Patrick Riley. Molecular graph convolutions: moving beyond fingerprints. *Journal of computer-aided molecular design*, 30:595–608, 2016.

[153] Yeji Wang, Shuo Wu, Yanwen Duan, and Yong Huang. A point cloud-based deep learning strategy for protein–ligand binding affinity prediction. *Briefings in Bioinformatics*, 23(1):bbab474, 2022.

[154] Nathaniel Thomas, Tess Smidt, Steven Kearnes, Lusann Yang, Li Li, Kai Kohlhoff, and Patrick Riley. Tensor field networks: Rotation-and translation-equivariant neural networks for 3d point clouds. *arXiv preprint arXiv:1802.08219*, 2018.

[155] Liguo Wang, Lu Zhao, Xian Liu, Jianjie Fu, and Aiqian Zhang. Seppcnet: deeping learning on a 3d surface electrostatic potential point cloud for enhanced toxicity classification and its application to suspected environmental estrogens. *Environmental Science & Technology*, 55(14):9958–9967, 2021.

[156] Soroush Ahmadi, Mohammad Amin Ghanavati, and Sohrab Rohani. Machine learning-guided prediction of cocrystals using point cloud-based molecular representation. *Chemistry of Materials*, 2024.

[157] Sukriti Singh and Raghavan B. Sunoj. Molecular Machine Learning for Chemical Catalysis: Prospects and Challenges. *Accounts of Chemical Research*, 56(3):402–412, February 2023. ISSN 0001-4842. doi:10.1021/acs.accounts.2c00801. URL https://doi.org/10.1021/acs.accounts.2c00801. Publisher: American Chemical Society.

[158] Andres M Bran and Philippe Schwaller. Transformers and large language models for chemistry and drug discovery. *arXiv [cs.LG]*, October 2023. URL http://arxiv.org/abs/2310.06083.

[159] Shilpa Shilpa, Gargee Kashyap, and Raghavan B. Sunoj. Recent Applications of Machine Learning in Molecular Property and Chemical Reaction Outcome Predictions. *The Journal of Physical Chemistry A*, 127(40):8253–8271, October 2023. ISSN 1089-5639. doi:10.1021/acs.jpca.3c04779. URL https://doi.org/10.1021/acs.jpca.3c04779. Publisher: American Chemical Society.

[160] Daniel S Wigh, Jonathan M Goodman, and Alexei A Lapkin. A review of molecular representation in the age of machine learning. *Wiley Interdiscip. Rev. Comput. Mol. Sci.*, 12(5), September 2022. ISSN 1759-0876,1759-0884. doi:10.1002/wcms.1603. URL https://wires.onlinelibrary.wiley.com/doi/10.1002/wcms.1603.

[161] Noel O'Boyle and Andrew Dalke. DeepSMILES: An Adaptation of SMILES for Use in Machine-Learning of Chemical Structures, September 2018. URL https://chemrxiv.org/engage/chemrxiv/article-details/60c73ed6567dfe7e5fec388d.

[162] Mario Krenn, Florian Häse, AkshatKumar Nigam, Pascal Friederich, and Alán Aspuru-Guzik. Self-Referencing Embedded Strings (SELFIES): A 100% robust molecular string representation. *Machine Learning: Science and Technology*, 1(4):045024, December 2020. ISSN 2632-2153. doi:10.1088/2632-2153/aba947. URL http://arxiv.org/abs/1905.13741. arXiv:1905.13741 [physics, physics:quant-ph, stat].

[163] Stephen R. Heller, Alan McNaught, Igor Pletnev, Stephen Stein, and Dmitrii Tchekhovskoi. InChI, the IUPAC International Chemical Identifier. *Journal of Cheminformatics*, 7(1):23, May 2015. ISSN 1758-2946. doi:10.1186/s13321-015-0068-4. URL https://doi.org/10.1186/s13321-015-0068-4.

[164] Manajit Das, Ankit Ghosh, and Raghavan B Sunoj. Advances in machine learning with chemical language models in molecular property and reaction outcome predictions. *J. Comput. Chem.*, February 2024. ISSN 0192-8651, 1096-987X. doi:10.1002/jcc.27315. URL http://dx.doi.org/10.1002/jcc.27315.

[165] John J Irwin, Teague Sterling, Michael M Mysinger, Erin S Bolstad, and Ryan G Coleman. Zinc: a free tool to discover chemistry for biology. *Journal of chemical information and modeling*, 52(7):1757–1768, 2012.

[166] Sunghwan Kim, Paul A Thiessen, Evan E Bolton, Jie Chen, Gang Fu, Asta Gindulyte, Lianyi Han, Jane He, Siqian He, Benjamin A Shoemaker, et al. Pubchem substance and compound databases. *Nucleic acids research*, 44(D1):D1202–D1213, 2016.

[167] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.

[168] Anna Gaulton, Louisa J Bellis, A Patricia Bento, Jon Chambers, Mark Davies, Anne Hersey, Yvonne Light, Shaun McGlinchey, David Michalovich, Bissan Al-Lazikani, and John P Overington. ChEMBL: a large-scale bioactivity database for drug discovery. *Nucleic Acids Res.*, 40(Database issue):D1100–7, January 2012. ISSN 0305-1048,1362-4962. doi:10.1093/nar/gkr777. URL http://dx.doi.org/10.1093/nar/gkr777.

[169] Rodney Kinney, Chloe Anastasiades, Russell Authur, Iz Beltagy, Jonathan Bragg, Alexandra Buraczynski, Isabel Cachola, Stefan Candra, Yoganand Chandrasekhar, Arman Cohan, Miles Crawford, Doug Downey, Jason Dunkelberger, Oren Etzioni, Rob Evans, Sergey Feldman, Joseph Gorney, David Graham, Fangzhou Hu, Regan Huff, Daniel King, Sebastian Kohlmeier, Bailey Kuehl, Michael Langan, Daniel Lin, Haokun Liu, Kyle Lo, Jaron Lochner, Kelsey MacMillan, Tyler Murray, Chris Newell, Smita Rao, Shaurya Rohatgi, Paul Sayre,

Zejiang Shen, Amanpreet Singh, Luca Soldaini, Shivashankar Subramanian, Amber Tanaka, Alex D Wade, Linda Wagner, Lucy Lu Wang, Chris Wilhelm, Caroline Wu, Jiangjiang Yang, Angele Zamarron, Madeleine Van Zuylen, and Daniel S Weld. The semantic scholar open data platform. *arXiv [cs.DL]*, January 2023. URL http://arxiv.org/abs/2301.10140.

[170] Yin Fang, Xiaozhuan Liang, Ningyu Zhang, Kangwei Liu, Rui Huang, Zhuo Chen, Xiaohui Fan, and Huajun Chen. Mol-instructions: A large-scale biomolecular instruction dataset for large language models. *arXiv [q-bio.QM]*, June 2023. URL http://arxiv.org/abs/2306.08018.

[171] Jin-Mao Wei, Xiao-Jie Yuan, Qing-Hua Hu, and Shu-Qin Wang. A novel measure for evaluating classifiers. *Expert Syst. Appl.*, 37(5):3799–3809, May 2010. ISSN 0957-4174,1873-6793. doi:10.1016/j.eswa.2009.11.040. URL http://dx.doi.org/10.1016/j.eswa.2009.11.040.

[172] Martin Krallinger, Obdulia Rabal, Florian Leitner, Miguel Vazquez, David Salgado, Zhiyong Lu, Robert Leaman, Yanan Lu, Donghong Ji, Daniel M Lowe, Roger A Sayle, Riza Theresa Batista-Navarro, Rafal Rak, Torsten Huber, Tim Rocktäschel, Sérgio Matos, David Campos, Buzhou Tang, Hua Xu, Tsendsuren Munkhdalai, Keun Ho Ryu, S V Ramanan, Senthil Nathan, Slavko Žitnik, Marko Bajec, Lutz Weber, Matthias Irmer, Saber A Akhondi, Jan A Kors, Shuo Xu, Xin An, Utpal Kumar Sikdar, Asif Ekbal, Masaharu Yoshioka, Thaer M Dieb, Miji Choi, Karin Verspoor, Madian Khabsa, C Lee Giles, Hongfang Liu, Komandur Elayavilli Ravikumar, Andre Lamurias, Francisco M Couto, Hong-Jie Dai, Richard Tzong-Han Tsai, Caglar Ata, Tolga Can, Anabel Usié, Rui Alves, Isabel Segura-Bedmar, Paloma Martínez, Julen Oyarzabal, and Alfonso Valencia. The CHEMDNER corpus of chemicals and drugs and its annotation principles. *J. Cheminform.*, 7(Suppl 1 Text mining for chemistry and the CHEMDNER track):S2, January 2015. ISSN 1758-2946,1758-2946. doi:10.1186/1758-2946-7-S1-S2. URL https://jcheminf.biomedcentral.com/articles/10.1186/1758-2946-7-S1-S2.

[173] Jiao Li, Yueping Sun, Robin J Johnson, Daniela Sciaky, Chih-Hsuan Wei, Robert Leaman, Allan Peter Davis, Carolyn J Mattingly, Thomas C Wiegers, and Zhiyong Lu. BioCreative V CDR task corpus: a resource for chemical disease relation extraction. *Database (Oxford)*, 2016:baw068, May 2016. ISSN 1758-0463. doi:10.1093/database/baw068. URL https://academic.oup.com/database/article-pdf/doi/10.1093/database/baw068/8224483/baw068.pdf.

[174] Rezarta Islamaj Dogan, Sun Kim, Andrew Chatr-Aryamontri, Chih-Hsuan Wei, Donald C Comeau, Rui Antunes, Sérgio Matos, Qingyu Chen, Aparna Elangovan, Nagesh C Panyam, Karin Verspoor, Hongfang Liu, Yanshan Wang, Zhuang Liu, Berna Altinel, Zehra Melce Hüsünbeyi, Arzucan Özgür, Aris Fergadis, Chen-Kai Wang, Hong-Jie Dai, Tung Tran, Ramakanth Kavuluru, Ling Luo, Albert Steppi, Jinfeng Zhang, Jinchan Qu, and Zhiyong Lu. Overview of the BioCreative VI precision medicine track: mining protein interactions and mutations for precision medicine. *Database (Oxford)*, 2019:bay147, January 2019. ISSN 1758-0463. doi:10.1093/database/bay147. URL https://academic.oup.com/database/article-pdf/doi/10.1093/database/bay147/27617946/bay147.pdf.

[175] M Ashburner, C A Ball, J A Blake, D Botstein, H Butler, J M Cherry, A P Davis, K Dolinski, S S Dwight, J T Eppig, M A Harris, D P Hill, L Issel-Tarver, A Kasarskis, S Lewis, J C Matese, J E Richardson, M Ringwald, G M Rubin, and G Sherlock. Gene ontology: tool for the unification of biology. the gene ontology consortium. *Nat. Genet.*, 25(1):25–29, May 2000. ISSN 1061-4036,1546-1718. doi:10.1038/75556. URL https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3037419/.

[176] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *International Conference on Learning Representations*, October 2020. URL https://openreview.net/pdf?id=d7KBjmI3GmQ.

[177] Sunghwan Kim, Jie Chen, Tiejun Cheng, Asta Gindulyte, Jia He, Siqian He, Qingliang Li, Benjamin A Shoemaker, Paul A Thiessen, Bo Yu, Leonid Zaslavsky, Jian Zhang, and Evan E Bolton. PubChem in 2021: new data content and improved web interfaces. *Nucleic Acids Res.*, 49(D1):D1388–D1395, January 2021. ISSN 0305-1048,1362-4962. doi:10.1093/nar/gkaa971. URL https://academic.oup.com/nar/article-pdf/49/D1/D1388/35363961/gkaa971.pdf.

[178] Ankit Pal, Logesh Kumar Umapathi, and Malaikannan Sankarasubbu. MedMCQA: A large-scale multi-subject multi-choice dataset for medical domain question answering. In *Conference on Health, Inference, and Learning*, pages 248–260. PMLR, April 2022. URL https://proceedings.mlr.press/v174/pal22a.html.

[179] Jieyu Lu and Yingkai Zhang. Unified Deep Learning Model for Multitask Reaction Predictions with Explanation. *J. Chem. Inf. Model.*, 62(6):1376–1387, March 2022. ISSN 1549-9596, 1549-960X. doi:10.1021/acs.jcim.1c01467. URL https://doi.org/10.1021/acs.jcim.1c01467.

[180] UniProt Consortium. UniProt: The universal protein knowledgebase in 2023. *Nucleic Acids Res.*, 51(D1): D523–D531, January 2023. ISSN 0305-1048,1362-4962. doi:10.1093/nar/gkac1052. URL https://academic.oup.com/nar/article-pdf/51/D1/D523/48441158/gkac1052.pdf.

[181] Marwin HS Segler, Mike Preuss, and Mark P Waller. Planning chemical syntheses with deep neural networks and symbolic ai. *Nature*, 555(7698):604–610, 2018.

[182] awesome-chemistry-datasets/code_of_conduct.md at main · kjappelbaum/awesome-chemistry-datasets, 2023. URL `https://github.com/kjappelbaum/awesome-chemistry-datasets/blob/main/code_of_conduct.md`.

[183] Adrian Mirza, Nawaf Alampara, Sreekanth Kunchapu, Benedict Emoekabu, Aswanth Krishnan, Mara Wilhelmi, Macjonathan Okereke, Juliane Eberhardt, Amir Mohammad Elahi, Maximilian Greiner, Caroline T Holick, Tanya Gupta, Mehrdad Asgari, Christina Glaubitz, Lea C Klepsch, Yannik Koster, Jakob Meyer, Santiago Miret, Tim Hoffmann, Fabian Alexander Kreth, Michael Ringleb, Nicole Roesner, Ulrich S Schubert, Leanne M Stafast, Dinga Wonanke, Michael Pieler, Philippe Schwaller, and Kevin Maik Jablonka. Are large language models superhuman chemists? *arXiv [cs.LG]*, April 2024. URL `http://arxiv.org/abs/2404.01475`.

[184] David E Gloriam. Bigger is better in virtual drug screens. *Nature*, 566(7743):193–194, February 2019. ISSN 0028-0836,1476-4687. doi:10.1038/d41586-019-00145-6. URL `http://dx.doi.org/10.1038/d41586-019-00145-6`.

[185] Ross Irwin, Spyridon Dimitriadis, Jiazhen He, and Esben Jannik Bjerrum. Chemformer: a pre-trained transformer for computational chemistry. *Machine Learning: Science and Technology*, 3(1):015022, January 2022. ISSN 2632-2153. doi:10.1088/2632-2153/ac3ffb. URL `https://dx.doi.org/10.1088/2632-2153/ac3ffb`. Publisher: IOP Publishing.

[186] Shengchao Liu, Weili Nie, Chengpeng Wang, Jiarui Lu, Zhuoran Qiao, Ling Liu, Jian Tang, Chaowei Xiao, and Animashree Anandkumar. Multi-modal molecule structure–text model for text-based retrieval and editing. *Nature Machine Intelligence*, 5(12):1447–1457, December 2023. ISSN 2522-5839. doi:10.1038/s42256-023-00759-6. URL `https://www.nature.com/articles/s42256-023-00759-6`. Publisher: Nature Publishing Group.

[187] Micha Livne, Zulfat Miftahutdinov, Elena Tutubalina, Maksim Kuznetsov, Daniil Polykovskiy, Annika Brundyn, Aastha Jhunjhunwala, Anthony Costa, Alex Aliper, Alán Aspuru-Guzik, and Alex Zhavoronkov. Nach0: Multimodal natural and chemical languages foundation model. *arXiv [cs.CL]*, November 2023. URL `http://arxiv.org/abs/2311.12410`.

[188] Pat Walters. We need better benchmarks for machine learning in drug discovery, 2023. URL `http://practicalcheminformatics.blogspot.com/2023/08/we-need-better-benchmarks-for-machine.html`.

[189] Cheng Fang, Ye Wang, Richard Grater, Sudarshan Kapadnis, Cheryl Black, Patrick Trapa, and Simone Sciabola. Prospective validation of machine learning algorithms for absorption, distribution, metabolism, and excretion prediction: An industrial perspective. *J. Chem. Inf. Model.*, 63(11):3263–3274, June 2023. ISSN 1549-9596,1549-960X. doi:10.1021/acs.jcim.3c00160. URL `https://pubs.acs.org/doi/full/10.1021/acs.jcim.3c00160`.

[190] Therapeutics data commons. `https://tdcommons.ai/`. Accessed: 2024-6-13.

[191] Kexin Huang, Tianfan Fu, Wenhao Gao, Yue Zhao, Yusuf Roohani, Jure Leskovec, Connor W Coley, Cao Xiao, Jimeng Sun, and Marinka Zitnik. Therapeutics data commons: Machine learning datasets and tasks for drug discovery and development. *arXiv [cs.LG]*, February 2021. URL `http://arxiv.org/abs/2102.09548`.

[192] Alejandro Velez-Arce, Kexin Huang, Michelle M Li, Xiang Lin, Wenhao Gao, Tianfan Fu, Manolis Kellis, Bradley L Pentelute, and Marinka Zitnik. TDC-2: Multimodal foundation for therapeutic science. *bioRxiv*, page 2024.06.12.598655, June 2024. doi:10.1101/2024.06.12.598655. URL `https://www.biorxiv.org/content/10.1101/2024.06.12.598655v2.abstract`.

[193] Alex Rich and Ben Birnbaum. Building adme benchmark datasets that drive impact, 2023. URL `https://www.inductive.bio/blog/building-better-benchmarks-for-adme-optimization`.

[194] Kausik Hira, Mohd Zaki, Dhruvil Sheth, Mausam, and N M Anoop Krishnan. Reconstructing the materials tetrahedron: challenges in materials information extraction. *Digit. Discov.*, 3(5):1021–1037, 2024. ISSN 2635-098X. doi:10.1039/d4dd00032c. URL `https://pubs.rsc.org/en/content/articlelanding/2024/dd/d4dd00032c`.

[195] John Dagdelen, Alexander Dunn, Sanghoon Lee, Nicholas Walker, Andrew S Rosen, Gerbrand Ceder, Kristin A Persson, and Anubhav Jain. Structured information extraction from scientific text with large language models. *Nat. Commun.*, 15(1):1418, February 2024. ISSN 2041-1723,2041-1723. doi:10.1038/s41467-024-45563-x. URL `https://www.nature.com/articles/s41467-024-45563-x`.

[196] Defne Circi, Ghazal Khalighinejad, Anlan Chen, Bhuwan Dhingra, and L Catherine Brinson. How well do large language models understand tables in materials science? *Integr. Mater. Manuf. Innov.*, 13(3):669–687, September

2024. ISSN 2193-9764,2193-9772. doi:10.1007/s40192-024-00362-6. URL `https://link.springer.com/article/10.1007/s40192-024-00362-6`.

[197] Jon M Laurent, Joseph D Janizek, Michael Ruzo, Michaela M Hinks, Michael J Hammerling, Siddharth Narayanan, Manvitha Ponnapati, Andrew D White, and Samuel G Rodriques. LAB-bench: Measuring capabilities of language models for biology research. *arXiv [cs.AI]*, July 2024. URL `http://arxiv.org/abs/2407.10362`.

[198] Nawaf Alampara, Santiago Miret, and Kevin Maik Jablonka. MatText: Do language models need more than text & scale for materials modeling? *arXiv [cond-mat.mtrl-sci]*, June 2024. URL `http://arxiv.org/abs/2406.17295`.

[199] Yu Song, Santiago Miret, and Bang Liu. MatSci-NLP: Evaluating scientific language models on materials science language tasks using text-to-schema modeling. *arXiv [cs.CL]*, May 2023. URL `http://arxiv.org/abs/2305.08264`.

[200] Mohd Zaki, Jayadeva, Mausam, and N M Anoop Krishnan. MaScQA: investigating materials science knowledge of large language models. *Digital Discovery*, 3(2):313–327, 2024. doi:10.1039/D3DD00188A. URL `https://pubs.rsc.org/en/content/articlehtml/2024/dd/d3dd00188a`.

[201] Afnan Sultan, Jochen Sieg, Miriam Mathea, and Andrea Volkamer. Transformers for molecular property prediction: Lessons learned from the past five years. *arXiv [cs.LG]*, April 2024. URL `http://arxiv.org/abs/2404.03969`.

[202] Janghoon Ock, Chakradhar Guntuboina, and Amir Barati Farimani. Catalyst Energy Prediction with CatBERTa: Unveiling Feature Exploration Strategies through Large Language Models. *ACS CATALYSIS*, 13(24):16032–16044, November 2023. ISSN 2155-5435. doi:10.1021/acscatal.3c04956.

[203] Atakan Yuksel, Erva Ulusoy, Atabey Ünlü, and Tunca Doğan. SELFormer: Molecular representation learning via SELFIES language models. *arXiv [q-bio.QM]*, April 2023. URL `http://arxiv.org/abs/2304.04662`.

[204] Michiko Yoshitake, Fumitaka Sato, Hiroyuki Kawano, and Hiroshi Teraoka. MaterialBERT for natural language processing of materials science texts. *Science and Technology of Advanced Materials: Methods*, 2(1):372–380, December 2022. ISSN 2766-0400. doi:10.1080/27660400.2022.2124831. URL `https://doi.org/10.1080/27660400.2022.2124831`.

[205] Jiahui Yu, Chengwei Zhang, Yingying Cheng, Yun-Fang Yang, Yuan-Bin She, Fengfan Liu, Weike Su, and An Su. SolvBERT for solvation free energy and solubility prediction: a demonstration of an NLP model for predicting the properties of molecular complexes. July 2022. URL `https://chemrxiv.org/engage/chemrxiv/article-details/62df4881a8e4dcc8f41cbadf`.

[206] Samuel Boobier, David RJ Hose, A John Blacker, and Bao N Nguyen. Machine learning with physicochemical relationships: solubility prediction in organic solvents and water. *Nature communications*, 11(1):5753, 2020.

[207] Zhi Hong, Aswathy Ajith, Gregory Pauloski, Eamon Duede, Kyle Chard, and Ian Foster. The diminishing returns of masked language models to science. *arXiv [cs.CL]*, May 2022. URL `http://arxiv.org/abs/2205.11342`.

[208] Shu Huang and Jacqueline M. Cole. BatteryBERT: A Pretrained Language Model for Battery Database Enhancement. *Journal of Chemical Information and Modeling*, 62(24):6365–6377, December 2022. ISSN 1549-9596. doi:10.1021/acs.jcim.2c00035. URL `https://doi.org/10.1021/acs.jcim.2c00035`. Publisher: American Chemical Society.

[209] Amalie Trewartha, Nicholas Walker, Haoyan Huo, Sanghoon Lee, Kevin Cruse, John Dagdelen, Alexander Dunn, Kristin A Persson, Gerbrand Ceder, and Anubhav Jain. Quantifying the advantage of domain-specific pre-training on named entity recognition tasks in materials science. *Patterns (N Y)*, 3(4):100488, April 2022. ISSN 2666-3899. doi:10.1016/j.patter.2022.100488. URL `http://dx.doi.org/10.1016/j.patter.2022.100488`.

[210] Tanishq Gupta, Mohd Zaki, N M Anoop Krishnan, and Mausam. MatSciBERT: A materials domain language model for text mining and information extraction. *npj Computational Materials*, 8(1):1–11, May 2022. ISSN 2057-3960, 2057-3960. doi:10.1038/s41524-022-00784-w. URL `https://www.nature.com/articles/s41524-022-00784-w`.

[211] Jerret Ross, Brian Belgodere, Vijil Chenthamarakshan, Inkit Padhi, Youssef Mroueh, and Payel Das. Large-scale chemical language representations capture molecular structure and properties. *Nat. Mach. Intell.*, 4(12):1256–1264, December 2022. ISSN 2522-5839,2522-5839. doi:10.1038/s42256-022-00580-7. URL `http://dx.doi.org/10.1038/s42256-022-00580-7`.

[212] Jiang Guo, A Santiago Ibanez-Lopez, Hanyu Gao, Victor Quach, Connor W Coley, Klavs F Jensen, and Regina Barzilay. Automated chemical reaction extraction from scientific literature. *J. Chem. Inf. Model.*,

62(9):2035–2045, May 2022. ISSN 1549-9596, 1549-960X. doi:10.1021/acs.jcim.1c00284. URL `http://dx.doi.org/10.1021/acs.jcim.1c00284`.

[213] Benedek Fabian, Thomas Edlich, Héléna Gaspar, Marwin Segler, Joshua Meyers, Marco Fiscato, and Mohamed Ahmed. Molecular representation learning with language models and domain-relevant auxiliary tasks. *arXiv [cs.LG]*, November 2020. URL `http://arxiv.org/abs/2011.13230`.

[214] Hoo-Chang Shin, Yang Zhang, Evelina Bakhturina, Raul Puri, Mostofa Patwary, Mohammad Shoeybi, and Raghav Mani. BioMegatron: Larger biomedical domain language model. *arXiv [cs.CL]*, October 2020. URL `http://arxiv.org/abs/2010.06060`.

[215] Yu Gu, Robert Tinn, Hao Cheng, Michael Lucas, Naoto Usuyama, Xiaodong Liu, Tristan Naumann, Jianfeng Gao, and Hoifung Poon. Domain-Specific language model pretraining for biomedical natural language processing. *ACM Trans. Comput. Healthcare*, 3(1):1–23, October 2021. doi:10.1145/3458754. URL `https://doi.org/10.1145/3458754`.

[216] Łukasz Maziarka, Tomasz Danel, Sławomir Mucha, Krzysztof Rataj, Jacek Tabor, and Stanisław Jastrzębski. Molecule attention transformer. *arXiv [cs.LG]*, February 2020. URL `http://arxiv.org/abs/2002.08264`.

[217] Sheng Wang, Yuzhi Guo, Yuhong Wang, Hongmao Sun, and Junzhou Huang. SMILES-BERT: Large Scale Unsupervised Pre-Training for Molecular Property Prediction. In *Proceedings of the 10th ACM International Conference on Bioinformatics, Computational Biology and Health Informatics*, BCB '19, pages 429–436, Niagara Falls NY USA, September 2019. Association for Computing Machinery. ISBN 9781450366663. doi:10.1145/3307339.3342186. URL `https://dl.acm.org/doi/10.1145/3307339.3342186`.

[218] Yifan Peng, Shankai Yan, and Zhiyong Lu. Transfer learning in biomedical natural language processing: An evaluation of BERT and ELMo on ten benchmarking datasets. *arXiv [cs.CL]*, June 2019. URL `http://arxiv.org/abs/1906.05474`.

[219] Kexin Huang, Jaan Altosaar, and Rajesh Ranganath. ClinicalBERT: Modeling clinical notes and predicting hospital readmission. *arXiv [cs.CL]*, April 2019. URL `http://arxiv.org/abs/1904.05342`.

[220] Iz Beltagy, Kyle Lo, and Arman Cohan. SciBERT: A pretrained language model for scientific text. *arXiv [cs.CL]*, March 2019. URL `http://arxiv.org/abs/1903.10676`.

[221] Jinhyuk Lee, Wonjin Yoon, Sungdong Kim, Donghyeon Kim, Sunkyu Kim, Chan Ho So, and Jaewoo Kang. BioBERT: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4):1234–1240, February 2020. ISSN 1367-4803, 1367-4811. doi:10.1093/bioinformatics/btz682. URL `http://dx.doi.org/10.1093/bioinformatics/btz682`.

[222] Di Wu, Qi Chen, Xiaojie Chen, Feng Han, Zhong Chen, and Yi Wang. The blood–brain barrier: structure, regulation, and drug delivery. *Signal Transduction and Targeted Therapy*, 8(1):1–27, May 2023. ISSN 2059-3635. doi:10.1038/s41392-023-01481-w. URL `https://www.nature.com/articles/s41392-023-01481-w`. Publisher: Nature Publishing Group.

[223] Caterina Bissantz, Bernd Kuhn, and Martin Stahl. A medicinal chemist's guide to molecular interactions. *Journal of medicinal chemistry*, 53(14):5061–5084, 2010.

[224] Stephen D Roughley and Allan M Jordan. The medicinal chemist's toolbox: an analysis of reactions used in the pursuit of drug candidates. *Journal of medicinal chemistry*, 54(10):3451–3479, 2011.

[225] Irini Doytchinova. Drug Design—Past, Present, Future. *Molecules*, 27(5):1496, February 2022. ISSN 1420-3049. doi:10.3390/molecules27051496. URL `https://www.ncbi.nlm.nih.gov/pmc/articles/PMC8911833/`.

[226] Introduction to Physical Polymer Science, 4th Edition | Wiley. URL `https://www.wiley.com/en-us/Introduction+to+Physical+Polymer+Science%2C+4th+Edition-p-9780471706069`.

[227] David J. Newman and Gordon M. Cragg. Natural Products as Sources of New Drugs from 1981 to 2014. *Journal of Natural Products*, 79(3):629–661, March 2016. ISSN 0163-3864. doi:10.1021/acs.jnatprod.5b01055. URL `https://doi.org/10.1021/acs.jnatprod.5b01055`. Publisher: American Chemical Society.

[228] Paulo Michel Pinheiro Ferreira, Daniel Dias Rufino Arcanjo, and Ana Paula Peron. Drug development, Brazilian biodiversity and political choices: Where are we heading? *Journal of Toxicology and Environmental Health, Part B*, 26(5):257–274, July 2023. ISSN 1093-7404. doi:10.1080/10937404.2023.2193762. URL `https://doi.org/10.1080/10937404.2023.2193762`. Publisher: Taylor & Francis _eprint: https://doi.org/10.1080/10937404.2023.2193762.

[229] Hartmuth C Kolb and K. Barry Sharpless. The growing impact of click chemistry on drug discovery. *Drug Discovery Today*, 8(24):1128–1137, December 2003. ISSN 1359-6446. doi:10.1016/S1359-6446(03)02933-7. URL `https://www.sciencedirect.com/science/article/pii/S1359644603029337`.

[230] Nathan J. Castellino, Andrew P. Montgomery, Jonathan J. Danon, and Michael Kassiou. Late-stage Functional-ization for Improving Drug-like Molecular Properties. *Chemical Reviews*, 123(13):8127–8153, July 2023. ISSN 0009-2665. doi:10.1021/acs.chemrev.2c00797. URL https://doi.org/10.1021/acs.chemrev.2c00797. Publisher: American Chemical Society.

[231] Komal Sharma, Krishna K. Sharma, Anku Sharma, and Rahul Jain. Peptide-based drug discovery: Cur-rent status and recent advances. *Drug Discovery Today*, 28(2):103464, February 2023. ISSN 1359-6446. doi:10.1016/j.drudis.2022.103464. URL https://www.sciencedirect.com/science/article/pii/S1359644622004573.

[232] Brandon J. Reizman and Klavs F. Jensen. Feedback in Flow for Accelerated Reaction Development. *Accounts of Chemical Research*, 49(9):1786–1796, September 2016. ISSN 0001-4842. doi:10.1021/acs.accounts.6b00261. URL https://doi.org/10.1021/acs.accounts.6b00261. Publisher: American Chemical Society.

[233] Joseph A. DiMasi, Henry G. Grabowski, and Ronald W. Hansen. Innovation in the pharmaceutical indus-try: New estimates of R&D costs. *Journal of Health Economics*, 47:20–33, May 2016. ISSN 1879-1646. doi:10.1016/j.jhealeco.2016.01.012.

[234] Errol G Lewars. *Computational chemistry: Introduction to the theory and applications of molecular and quantum mechanics*. Springer International Publishing, Cham, Switzerland, 4 edition, April 2024. ISBN 9783031514425,9783031514432. doi:10.1007/978-3-031-51443-2. URL https://link.springer.com/book/10.1007/978-3-031-51443-2.

[235] Xavier Bidault and Santanu Chaudhuri. How Accurate Can Crystal Structure Predictions Be for High-Energy Molecular Crystals? *Molecules*, 28(11):4471, January 2023. ISSN 1420-3049. doi:10.3390/molecules28114471. URL https://www.mdpi.com/1420-3049/28/11/4471. Number: 11 Publisher: Multidisciplinary Digital Publishing Institute.

[236] Edward O. Pyzer-Knapp, Linjiang Chen, Graeme M. Day, and Andrew I. Cooper. Accelerating computational discovery of porous solids through improved navigation of energy-structure-function maps. *Science Advances*, 7 (33):eabi4763, August 2021. doi:10.1126/sciadv.abi4763. URL https://www.science.org/doi/10.1126/sciadv.abi4763. Publisher: American Association for the Advancement of Science.

[237] Scott Fredericks, Kevin Parrish, Dean Sayre, and Qiang Zhu. PyXtal: A Python library for crystal structure generation and symmetry analysis. *Computer Physics Communications*, 261:107810, April 2021. ISSN 0010-4655. doi:10.1016/j.cpc.2020.107810. URL https://www.sciencedirect.com/science/article/pii/S0010465520304057.

[238] David H. Case, Josh E. Campbell, Peter J. Bygrave, and Graeme M. Day. Convergence Properties of Crystal Structure Prediction by Quasi-Random Sampling. *Journal of Chemical Theory and Computation*, 12(2):910–924, February 2016. ISSN 1549-9618. doi:10.1021/acs.jctc.5b01112. URL https://doi.org/10.1021/acs.jctc.5b01112. Publisher: American Chemical Society.

[239] A. V. Kazantsev, P. G. Karamertzanis, C. S. Adjiman, and C. C. Pantelides. Efficient Handling of Molecular Flexibility in Lattice Energy Minimization of Organic Crystals. *Journal of Chemical Theory and Computation*, 7(6):1998–2016, June 2011. ISSN 1549-9618. doi:10.1021/ct100597e. URL https://doi.org/10.1021/ct100597e. Publisher: American Chemical Society.

[240] Guannan Huang, Yani Guo, Ye Chen, and Zhengwei Nie. Application of Machine Learning in Material Synthesis and Property Prediction. *Materials*, 16(17):5977, January 2023. ISSN 1996-1944. doi:10.3390/ma16175977. URL https://www.mdpi.com/1996-1944/16/17/5977. Number: 17 Publisher: Multidisciplinary Digital Publishing Institute.

[241] Karina Martinez-Mayorga, José G. Rosas-Jiménez, Karla Gonzalez-Ponce, Edgar López-López, Antonio Neme, and José L. Medina-Franco. The pursuit of accurate predictive models of the bioactivity of small molecules. *Chemical Science*, 15(6):1938–1952, February 2024. ISSN 2041-6539. doi:10.1039/D3SC05534E. URL https://pubs.rsc.org/en/content/articlelanding/2024/sc/d3sc05534e. Publisher: The Royal Society of Chemistry.

[242] Geemi P Wellawatte, Heta A Gandhi, Aditi Seshadri, and Andrew D White. A PERSPECTIVE ON EXPLA-NATIONS OF MOLECULAR PREDICTION MODELS. *J. Chem. Theory Comput.*, 19(8):2149–2160, April 2023. ISSN 1549-9618, 1549-9626. doi:10.1021/acs.jctc.2c01235. URL http://dx.doi.org/10.1021/acs.jctc.2c01235.

[243] Xiang Deng, Vasilisa Bashlovkina, Feng Han, Simon Baumgartner, and Michael Bendersky. LLMs to the Moon? Reddit Market Sentiment Analysis with Large Language Models. In *Companion Proceedings of the ACM Web Conference 2023*, WWW '23 Companion, pages 1014–1019, New York, NY, USA, April 2023.

Association for Computing Machinery. ISBN 978-1-4503-9419-2. doi:10.1145/3543873.3587605. URL https://dl.acm.org/doi/10.1145/3543873.3587605.

[244] P Schwaller, Daniel Probst, A Vaucher, Vishnu H Nair, D Kreutter, T Laino, and J Reymond. Mapping the space of chemical reactions using attention-based neural networks. *Nat. Mach. Intell.*, 3(2):144–152, August 2020. ISSN 2522-5839,2522-5839. doi:10.1038/s42256-020-00284-w. URL https://www.nature.com/articles/s42256-020-00284-w.

[245] Alessandra Toniato, Alain C Vaucher, Philippe Schwaller, and Teodoro Laino. Enhancing diversity in language based models for single-step retrosynthesis. *Digital Discovery*, 2(2):489–501, April 2023. ISSN 2635-098X. doi:10.1039/D2DD00110A. URL https://pubs.rsc.org/en/content/articlelanding/2023/dd/d2dd00110a.

[246] Philippe Schwaller, Benjamin Hoover, Jean-Louis Reymond, Hendrik Strobelt, and Teodoro Laino. Extraction of organic chemistry grammar from unsupervised learning of chemical reactions. *Science Advances*, 7 (15):eabe4166, April 2021. doi:10.1126/sciadv.abe4166. URL https://www.science.org/doi/10.1126/sciadv.abe4166. Publisher: American Association for the Advancement of Science.

[247] Philippe Schwaller, Alain C Vaucher, Teodoro Laino, and Jean-Louis Reymond. Prediction of chemical reaction yields using deep learning. *Mach. Learn. Sci. Technol.*, 2(1):015016, March 2021. ISSN 2632-2153. doi:10.1088/2632-2153/abc81d. URL https://iopscience.iop.org/article/10.1088/2632-2153/abc81d/meta.

[248] Sheng Wang, Yuzhi Guo, Yuhong Wang, Hongmao Sun, and Junzhou Huang. SMILES-BERT: Large Scale Unsupervised Pre-Training for Molecular Property Prediction. In *Proceedings of the 10th ACM International Conference on Bioinformatics, Computational Biology and Health Informatics*, BCB '19, pages 429–436, New York, NY, USA, September 2019. Association for Computing Machinery. ISBN 978-1-4503-6666-3. doi:10.1145/3307339.3342186. URL https://doi.org/10.1145/3307339.3342186.

[249] Xiao-Chen Zhang, Cheng-Kun Wu, Jia-Cai Yi, Xiang-Xiang Zeng, Can-Qun Yang, Ai-Ping Lu, Hou T-j, Ting-Jun Hou, and Dong-Sheng Cao. Pushing the Boundaries of Molecular Property Prediction for Drug Discovery with Multitask Learning BERT Enhanced by SMILES Enumeration. *Research*, 2022:0004, December 2022. ISSN 2096-5168. doi:10.34133/research.0004. URL http://dx.doi.org/10.34133/research.0004.

[250] Guoli Xiong, Zhenxing Wu, Jiacai Yi, Li Fu, Zhijiang Yang, Changyu Hsieh, Mingzhu Yin, Xiangxiang Zeng, Chengkun Wu, Aiping Lu, Xiang Chen, Tingjun Hou, and Dongsheng Cao. ADMETlab 2.0: an integrated online platform for accurate and comprehensive predictions of ADMET properties. *Nucleic Acids Res.*, 49(W1):W5–W14, July 2021. ISSN 0305-1048, 1362-4962. doi:10.1093/nar/gkab255. URL http://dx.doi.org/10.1093/nar/gkab255.

[251] Teague Sterling and John J. Irwin. ZINC 15 – Ligand Discovery for Everyone. *Journal of Chemical Information and Modeling*, 55(11):2324–2337, November 2015. ISSN 1549-9596. doi:10.1021/acs.jcim.5b00559. URL https://doi.org/10.1021/acs.jcim.5b00559. Publisher: American Chemical Society.

[252] Jerret Ross, Brian Belgodere, Vijil Chenthamarakshan, Inkit Padhi, Youssef Mroueh, and Payel Das. Large-scale chemical language representations capture molecular structure and properties. *Nature Machine Intelligence*, 4(12):1256–1264, December 2022. ISSN 2522-5839, 2522-5839. doi:10.1038/s42256-022-00580-7. URL https://www.nature.com/articles/s42256-022-00580-7.

[253] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, Mariama Drame, Quentin Lhoest, and Alexander M. Rush. Huggingface's transformers: State-of-the-art natural language processing, 2020.

[254] Kevin Yang, Kyle Swanson, Wengong Jin, Connor Coley, Philipp Eiden, Hua Gao, Angel Guzman-Perez, Timothy Hopper, Brian Kelley, Miriam Mathea, Andrew Palmer, Volker Settels, Tommi Jaakkola, Klavs Jensen, and Regina Barzilay. Analyzing Learned Molecular Representations for Property Prediction. *Journal of Chemical Information and Modeling*, 59(8):3370–3388, August 2019. ISSN 1549-960X. doi:10.1021/acs.jcim.9b00237.

[255] Philippe Schwaller, Teodoro Laino, Théophile Gaudin, Peter Bolgar, Christopher A Hunter, Costas Bekas, and Alpha A Lee. Molecular Transformer: A Model for Uncertainty-Calibrated Chemical Reaction Prediction. *ACS Central Science*, 5(9):1572–1583, September 2019. ISSN 2374-7943, 2374-7951. doi:10.1021/acscentsci.9b00576. URL https://doi.org/10.1021/acscentsci.9b00576.

[256] Jesse Vig. Bertviz: A tool for visualizing multihead self-attention in the bert model. In *ICLR workshop: Debugging machine learning models*, volume 3, 2019.

[257] Michael A. Skinnider. Invalid SMILES are beneficial rather than detrimental to chemical language models. *Nature Machine Intelligence*, 6(4):437–448, April 2024. ISSN 2522-5839. doi:10.1038/s42256-024-00821-x. URL `https://www.nature.com/articles/s42256-024-00821-x`. Publisher: Nature Publishing Group.

[258] Florence H Vermeire and William H Green. Transfer learning for solvation free energies: From quantum chemistry to experiments. *Chemical Engineering Journal*, 418:129307, 2021.

[259] David L Mobley and J Peter Guthrie. FreeSolv: a database of experimental and calculated hydration free energies, with input files. *J. Comput. Aided Mol. Des.*, 28(7):711–720, July 2014. ISSN 0920-654X,1573-4951. doi:10.1007/s10822-014-9747-x. URL `https://link.springer.com/article/10.1007/s10822-014-9747-x`.

[260] Aleksandr V Marenich, Casey P Kelly, Jason D Thompson, Gregory D Hawkins, Candee C Chambers, David J Giesen, Paul Winget, Christopher J Cramer, and Donald G Truhlar. Minnesota solvation database (mnsol) version 2012. 2020.

[261] Edouard Moine, Romain Privat, Baptiste Sirjean, and Jean-Noël Jaubert. Estimation of solvation quantities from experimental thermodynamic data: Development of the comprehensive compsol databank for pure and mixed solutes. *Journal of Physical and Chemical Reference Data*, 46(3), 2017.

[262] Laura M Grubbs, Mariam Saifullah, E Nohelli, Shulin Ye, Sai S Achi, William E Acree Jr, and Michael H Abraham. Mathematical correlations for describing solute transfer into functionalized alkane solvents containing hydroxyl, ether, ester or ketone solvents. *Fluid phase equilibria*, 298(1):48–53, 2010.

[263] Kevin Yang, Kyle Swanson, Wengong Jin, Connor Coley, Hua Gao, Angel Guzman-Perez, Timothy Hopper, Brian P Kelley, Andrew Palmer, Volker Settels, et al. Are learned molecular representations ready for prime time? 2019. doi:10.26434/chemrxiv.7940594.v2.

[264] Yu Rong, Yatao Bian, Tingyang Xu, Wei-Yang Xie, Ying Wei, Wen-Bing Huang, and Junzhou Huang. Self-Supervised Graph Transformer on Large-Scale Molecular Data. *Adv. Neural Inf. Process. Syst.*, 33:12559–12571, June 2020. ISSN 1049-5258. URL `https://proceedings.neurips.cc/paper_files/paper/2020/hash/94aef38441efa3380a3bed3faf1f9d5d-Abstract.html`.

[265] Benedikt Winter, Clemens Winter, Johannes Schilling, and André Bardow. A smile is all you need: predicting limiting activity coefficients from SMILES with natural language processing. *Digital Discovery*, 1(6):859–869, December 2022. ISSN 2635-098X. doi:10.1039/D2DD00058J. URL `https://pubs.rsc.org/en/content/articlelanding/2022/dd/d2dd00058j`. Publisher: RSC.

[266] Jing Jiang, Yachao Li, Ruisheng Zhang, and Yunwu Liu. INTransformer: Data augmentation-based contrastive learning by injecting noise into transformer for molecular property prediction. *Journal of Molecular Graphics and Modelling*, 128:108703, May 2024. ISSN 1093-3263. doi:10.1016/j.jmgm.2024.108703. URL `https://www.sciencedirect.com/science/article/pii/S1093326324000032`.

[267] Shengchao Liu, Weili Nie, Chengpeng Wang, Jiarui Lu, Zhuoran Qiao, Ling Liu, Jian Tang, Chaowei Xiao, and Anima Anandkumar. Multi-modal molecule structure-text model for text-based retrieval and editing. *arXiv [cs.LG]*, December 2022. URL `http://arxiv.org/abs/2212.10789`.

[268] Sunghwan Kim, Jie Chen, Tiejun Cheng, Asta Gindulyte, Jia He, Siqian He, Qingliang Li, Benjamin A Shoemaker, Paul A Thiessen, Bo Yu, Leonid Zaslavsky, Jian Zhang, and Evan E Bolton. PubChem 2023 update. *Nucleic Acids Res.*, 51(D1):D1373–D1380, January 2023. ISSN 0305-1048, 1362-4962. doi:10.1093/nar/gkac956. URL `http://dx.doi.org/10.1093/nar/gkac956`.

[269] Minghao Xu, Xinyu Yuan, Santiago Miret, and Jian Tang. ProtST: Multi-modality learning of protein sequences and biomedical texts. In *International Conference on Machine Learning*, pages 38749–38767. PMLR, July 2023. URL `https://proceedings.mlr.press/v202/xu23t.html`.

[270] Glen M Hocky. Connecting molecular properties with plain language. *Nat. Mach. Intell.*, 6(3):249–250, March 2024. ISSN 2522-5839,2522-5839. doi:10.1038/s42256-024-00812-y. URL `https://www.nature.com/articles/s42256-024-00812-y`.

[271] Juan Manuel Zambrano Chaves, Eric Wang, Tao Tu, Eeshit Dhaval Vaishnav, Byron Lee, S Sara Mahdavi, Christopher Semturs, David Fleet, Vivek Natarajan, and Shekoofeh Azizi. Tx-LLM: A large language model for therapeutics. *arXiv [cs.CL]*, June 2024. URL `http://arxiv.org/abs/2406.06316`.

[272] Elliot Bolton, Abhinav Venigalla, Michihiro Yasunaga, David Hall, Betty Xiong, Tony Lee, Roxana Daneshjou, Jonathan Frankle, Percy Liang, Michael Carbin, and Christopher D Manning. BioMedLM: A 2.7B parameter language model trained on biomedical text. *arXiv [cs.CL]*, March 2024. URL `http://arxiv.org/abs/2403.18421`.

[273] Botao Yu, Frazier N. Baker, Ziqi Chen, Xia Ning, and Huan Sun. LlaSMol: Advancing Large Language Models for Chemistry with a Large-Scale, Comprehensive, High-Quality Instruction Tuning Dataset, April 2024. URL http://arxiv.org/abs/2402.09391. arXiv:2402.09391 [cs].

[274] Yanis Labrak, Adrien Bazoge, Emmanuel Morin, Pierre-Antoine Gourraud, Mickael Rouvier, and Richard Dufour. BioMistral: A collection of open-source pretrained large language models for medical domains. *arXiv [cs.CL]*, February 2024. URL http://arxiv.org/abs/2402.10373.

[275] Sara Pieri, Sahal Shaji Mullappilly, Fahad Shahbaz Khan, Rao Muhammad Anwer, Salman Khan, Timothy Baldwin, and Hisham Cholakkal. BiMediX: Bilingual medical mixture of experts LLM. *arXiv [cs.CL]*, February 2024. URL http://arxiv.org/abs/2402.13253.

[276] Xuyang Zhao, Qibin Zhao, and Toshihisa Tanaka. EpilepsyLLM: Domain-specific large language model fine-tuned with epilepsy medical knowledge. *arXiv [cs.CL]*, January 2024. URL http://arxiv.org/abs/2401.05908.

[277] Zhihong Chen, Maya Varma, Jean-Benoit Delbrouck, Magdalini Paschali, Louis Blankemeier, Dave Van Veen, Jeya Maria Jose Valanarasu, Alaa Youssef, Joseph Paul Cohen, Eduardo Pontes Reis, Emily B Tsai, Andrew Johnston, Cameron Olsen, Tanishq Mathew Abraham, Sergios Gatidis, Akshay S Chaudhari, and Curtis Langlotz. CheXagent: Towards a foundation model for chest X-ray interpretation. *arXiv [cs.CV]*, January 2024. URL http://arxiv.org/abs/2401.12208.

[278] Gregory W. Kyro, Anton Morgunov, Rafael I. Brent, and Victor S. Batista. ChemSpaceAL: An Efficient Active Learning Methodology Applied to Protein-Specific Molecular Generation. *Journal of Chemical Information and Modeling*, 64(3):653–665, February 2024. ISSN 1549-9596. doi:10.1021/acs.jcim.3c01456. URL https://doi.org/10.1021/acs.jcim.3c01456. Publisher: American Chemical Society.

[279] Yizhen Luo, Jiahuan Zhang, Siqi Fan, Kai Yang, Yushuai Wu, Mu Qiao, and Zaiqing Nie. BioMedGPT: Open multimodal generative pre-trained transformer for BioMedicine. *arXiv [cs.CE]*, August 2023. URL https://github.com/PharMolix/OpenBioMed.

[280] Tong Xie, Yuwei Wan, Wei Huang, Zhenyu Yin, Yixuan Liu, Shaozhou Wang, Qingyuan Linghu, Chunyu Kit, Clara Grazian, Wenjie Zhang, Imran Razzak, and Bram Hoex. DARWIN series: Domain specific large language models for natural science. *arXiv [cs.CL]*, August 2023. URL http://arxiv.org/abs/2308.13565.

[281] Chaoyi Wu, Weixiong Lin, Xiaoman Zhang, Ya Zhang, Yanfeng Wang, and Weidi Xie. PMC-LLaMA: Towards building open-source language models for medicine. *arXiv [cs.CL]*, April 2023. URL http://arxiv.org/abs/2304.14454.

[282] Renqian Luo, Liai Sun, Yingce Xia, Tao Qin, Sheng Zhang, Hoifung Poon, and Tie-Yan Liu. BioGPT: generative pre-trained transformer for biomedical text generation and mining. *Brief. Bioinform.*, 23(6), November 2022. ISSN 1467-5463, 1477-4054. doi:10.1093/bib/bbac409. URL http://dx.doi.org/10.1093/bib/bbac409.

[283] Nathan C. Frey, Ryan Soklaski, Simon Axelrod, Siddharth Samsi, Rafael Gómez-Bombarelli, Connor W. Coley, and Vijay Gadepally. Neural scaling of deep chemical models. *Nature Machine Intelligence*, 5(11):1297–1305, November 2023. ISSN 2522-5839. doi:10.1038/s42256-023-00740-3. URL https://www.nature.com/articles/s42256-023-00740-3. Publisher: Nature Publishing Group.

[284] Viraj Bagal, Rishal Aggarwal, P K Vinod, and U Deva Priyakumar. MolGPT: Molecular generation using a Transformer-Decoder model. *J. Chem. Inf. Model.*, 62(9):2064–2076, May 2022. ISSN 1549-9596, 1549-960X. doi:10.1021/acs.jcim.1c00600. URL http://dx.doi.org/10.1021/acs.jcim.1c00600.

[285] Sanjar Adilov. Generative Pre-Training from molecules. *ChemRxiv*, September 2021. doi:10.26434/chemrxiv-2021-5fwjd. URL https://chemrxiv.org/engage/chemrxiv/article-details/6142f60742198e8c31782e9e.

[286] Jesse W.-H. Li and John C. Vederas. Drug Discovery and Natural Products: End of an Era or an Endless Frontier? *Science*, 325(5937):161–165, July 2009. doi:10.1126/science.1168243. URL https://www.science.org/doi/abs/10.1126/science.1168243. Publisher: American Association for the Advancement of Science.

[287] David J. Newman and Gordon M. Cragg. Natural Products As Sources of New Drugs over the 30 Years from 1981 to 2010. *Journal of Natural Products*, 75(3):311–335, March 2012. ISSN 0163-3864. doi:10.1021/np200906s. URL https://doi.org/10.1021/np200906s. Publisher: American Chemical Society.

[288] Maya A. Farha and Eric D. Brown. Strategies for target identification of antimicrobial natural products. *Natural Product Reports*, 33(5):668–680, 2016. doi:10.1039/C5NP00127G. URL https://pubs.rsc.org/en/content/articlelanding/2016/np/c5np00127g. Publisher: Royal Society of Chemistry.

[289] AkshatKumar Nigam, Robert Pollice, Mario Krenn, Gabriel dos Passos Gomes, and Alán Aspuru-Guzik. Beyond generative models: superfast traversal, optimization, novelty, exploration and discovery (STONED) algorithm for molecules using SELFIES. *Chemical Science*, 12(20):7079–7090, 2021. doi:10.1039/D1SC00231G. URL https://pubs.rsc.org/en/content/articlelanding/2021/sc/d1sc00231g. Publisher: Royal Society of Chemistry.

[290] Heta A Gandhi and Andrew D White. Explaining molecular properties with natural language. *ChemRxiv*, October 2022. doi:10.26434/chemrxiv-2022-v5p6m-v3. URL https://chemrxiv.org/engage/chemrxiv/article-details/633731d1f764e6e535093041.

[291] Yuanqi Du, Arian R Jamasb, Jeff Guo, Tianfan Fu, Charles Harris, Yingheng Wang, Chenru Duan, Pietro Liò, Philippe Schwaller, and Tom L Blundell. Machine learning-aided generative molecular design. *Nat. Mach. Intell.*, pages 1–16, June 2024. ISSN 2522-5839,2522-5839. doi:10.1038/s42256-024-00843-5. URL https://www.nature.com/articles/s42256-024-00843-5.

[292] Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for nlp. In *International conference on machine learning*, pages 2790–2799. PMLR, 2019.

[293] Addis S. Fuhr and Bobby G. Sumpter. Deep generative models for materials discovery and machine learning-accelerated innovation. *FRONTIERS IN MATERIALS*, 9, MAR 22 2022. ISSN 2296-8016. doi:10.3389/fmats.2022.865270.

[294] Ri Han, Hongryul Yoon, Gahee Kim, Hyundo Lee, and Yoonji Lee. Revolutionizing medicinal chemistry: The application of artificial intelligence (ai) in early drug discovery. *PHARMACEUTICALS*, 16(9), SEP 2023. doi:10.3390/ph16091259.

[295] Nikoletta-Maria Koutroumpa, Konstantinos D. Papavasileiou, Anastasios G. Papadiamantis, Georgia Melagraki, and Antreas Afantitis. A systematic review of deep learning methodologies used in the drug discovery process with emphasis on in vivo validation. *INTERNATIONAL JOURNAL OF MOLECULAR SCIENCES*, 24(7), APR 2023. ISSN 1661-6596. doi:10.3390/ijms24076573.

[296] Douglas B. Kell, Soumitra Samanta, and Neil Swainston. Deep learning and generative methods in cheminformatics and chemical biology: navigating small molecule space intelligently. *BIOCHEMICAL JOURNAL*, 477 (23):4559–4580, DEC 2020. ISSN 0264-6021. doi:10.1042/BCJ20200781.

[297] Camille Bilodeau, Wengong Jin, Tommi Jaakkola, Regina Barzilay, and Klavs F. Jensen. Generative models for molecular discovery: Recent advances and challenges. *WILEY INTERDISCIPLINARY REVIEWS-COMPUTATIONAL MOLECULAR SCIENCE*, 12(5), SEP 2022. ISSN 1759-0876. doi:10.1002/wcms.1608.

[298] Amit Gangwal, Azim Ansari, Iqrar Ahmad, Abul Kalam Azad, Vinoth Kumarasamy, Vetriselvan Subramaniyan, and Ling Shing Wong. Generative artificial intelligence in drug discovery: basic framework, recent advances, challenges, and opportunities. *FRONTIERS IN PHARMACOLOGY*, 15, FEB 7 2024. doi:10.3389/fphar.2024.1331062.

[299] Martin Vogt. Using deep neural networks to explore chemical space. *EXPERT OPINION ON DRUG DISCOVERY*, 17(3):297–304, MAR 4 2022. ISSN 1746-0441. doi:10.1080/17460441.2022.2019704.

[300] Sekhar Talluri, Mohammad Amjad Kamal, and Rama Rao Malla. Novel computational methods for cancer drug design. *CURRENT MEDICINAL CHEMISTRY*, 31(5):554–572, 2024. ISSN 0929-8673. doi:10.2174/0929867330666230403100008.

[301] Daniil Polykovskiy, Alexander Zhebrak, Benjamin Sanchez-Lengeling, Sergey Golovanov, Oktai Tatanov, Stanislav Belyaev, Rauf Kurbanov, Aleksey Artamonov, Vladimir Aladinskiy, Mark Veselov, Artur Kadurin, Simon Johansson, Hongming Chen, Sergey Nikolenko, Alán Aspuru-Guzik, and Alex Zhavoronkov. Molecular sets (MOSES): A benchmarking platform for molecular generation models. *Front. Pharmacol.*, 11:565644, December 2020. ISSN 1663-9812. doi:10.3389/fphar.2020.565644. URL https://www.frontiersin.org/journals/pharmacology/articles/10.3389/fphar.2020.565644/full.

[302] Nathan Brown, Marco Fiscato, Marwin H S Segler, and Alain C Vaucher. GuacaMol: Benchmarking models for de novo molecular design. *J. Chem. Inf. Model.*, 59(3):1096–1108, March 2019. ISSN 1549-9596, 1549-960X. doi:10.1021/acs.jcim.8b00839. URL http://dx.doi.org/10.1021/acs.jcim.8b00839.

[303] Kristina Preuer, Philipp Renz, Thomas Unterthiner, Sepp Hochreiter, and Gunter Klambauer. Fréchet chemnet distance: a metric for generative models for molecules in drug discovery. *Journal of chemical information and modeling*, 58(9):1736–1741, 2018.

[304] Suhail Haroon, Hafsath C.a., and Jereesh A.s. Generative Pre-trained Transformer (GPT) based model with relative attention for de novo drug design. *Computational Biology and Chemistry*, 106:107911, October 2023.

ISSN 1476-9271. doi:10.1016/j.compbiolchem.2023.107911. URL https://www.sciencedirect.com/science/article/pii/S1476927123001020.

[305] Jianmin Wang, Jiashun Mao, Meng Wang, Xiangyang Le, and Yunyun Wang. Explore drug-like space with deep generative models. *Methods*, 210:52–59, February 2023. ISSN 1046-2023. doi:10.1016/j.ymeth.2023.01.004. URL https://www.sciencedirect.com/science/article/pii/S1046202323000129.

[306] Jiashun Mao, Jianmin Wang, Amir Zeb, Kwang-Hwi Cho, Haiyan Jin, Jongwan Kim, Onju Lee, Yunyun Wang, and Kyoung Tai No. Transformer-Based Molecular Generative Model for Antiviral Drug Design. *Journal of Chemical Information and Modeling*, 64(7):2733–2745, April 2024. ISSN 1549-9596. doi:10.1021/acs.jcim.3c00536. URL https://doi.org/10.1021/acs.jcim.3c00536. Publisher: American Chemical Society.

[307] Wenyi Zhang, Kaiyue Zhang, and Jing Huang. A Simple Way to Incorporate Target Structural Information in Molecular Generative Models. *Journal of Chemical Information and Modeling*, 63(12):3719–3730, June 2023. ISSN 1549-9596. doi:10.1021/acs.jcim.3c00293. URL https://doi.org/10.1021/acs.jcim.3c00293. Publisher: American Chemical Society.

[308] Xun Wang, Changnan Gao, Peifu Han, Xue Li, Wenqi Chen, Alfonso Rodríguez Patón, Shuang Wang, and Pan Zheng. PETrans: De Novo Drug Design with Protein-Specific Encoding Based on Transfer Learning. *International Journal of Molecular Sciences*, 24(2):1146, January 2023. ISSN 1422-0067. doi:10.3390/ijms24021146. URL https://www.mdpi.com/1422-0067/24/2/1146. Number: 2 Publisher: Multidisciplinary Digital Publishing Institute.

[309] Yasuhiro Yoshikai, Tadahaya Mizuno, Shumpei Nemoto, and Hiroyuki Kusuhara. A novel molecule generative model of VAE combined with Transformer for unseen structure generation, April 2024. URL http://arxiv.org/abs/2402.11950. arXiv:2402.11950 [physics, q-bio].

[310] Xiaoying Yan, Chi Gu, Yuehua Feng, and Jiaxin Han. Predicting Drug-drug Interaction with Graph Mutual Interaction Attention Mechanism. *Methods*, 223:16–25, March 2024. ISSN 1046-2023. doi:10.1016/j.ymeth.2024.01.009. URL https://www.sciencedirect.com/science/article/pii/S1046202324000276.

[311] Tao Shen, Jiale Guo, Zunsheng Han, Gao Zhang, Qingxin Liu, Xinxin Si, Dongmei Wang, Song Wu, and Jie Xia. AutoMolDesigner for Antibiotic Discovery: An AI-Based Open-Source Software for Automated Design of Small-Molecule Antibiotics. *Journal of Chemical Information and Modeling*, 64(3):575–583, February 2024. ISSN 1549-9596. doi:10.1021/acs.jcim.3c01562. URL https://doi.org/10.1021/acs.jcim.3c01562. Publisher: American Chemical Society.

[312] G Richard Bickerton, Gaia V Paolini, Jérémy Besnard, Sorel Muresan, and Andrew L Hopkins. Quantifying the chemical beauty of drugs. *Nature chemistry*, 4(2):90–98, 2012.

[313] Govindan Subramanian, Bharath Ramsundar, Vijay Pande, and Rajiah Aldrin Denny. Computational modeling of β-secretase 1 (bace-1) inhibitors using ligand based approaches. *Journal of chemical information and modeling*, 56(10):1936–1949, 2016.

[314] Guozhang Xu, Marta C Abad, Peter J Connolly, Michael P Neeper, Geoffrey T Struble, Barry A Springer, Stuart L Emanuel, Niranjan Pandey, Robert H Gruninger, Mary Adams, et al. 4-amino-6-arylamino-pyrimidine-5-carbaldehyde hydrazones as potent erbb-2/egfr dual kinase inhibitors. *Bioorganic & medicinal chemistry letters*, 18(16):4615–4619, 2008.

[315] Leiye Yu, Licong He, Bing Gan, Rujuan Ti, Qingjie Xiao, Xin Yang, Hongli Hu, Lizhe Zhu, Sheng Wang, and Ruobing Ren. Structural insights into sphingosine-1-phosphate receptor activation. *Proceedings of the National Academy of Sciences*, 119(16):e2117716119, 2022.

[316] Peiyu Xu, Sijie Huang, Huibing Zhang, Chunyou Mao, X Edward Zhou, Xi Cheng, Icaro A Simon, Dan-Dan Shen, Hsin-Yung Yen, Carol V Robinson, et al. Structural insights into the lipid and ligand regulation of serotonin receptors. *Nature*, 592(7854):469–473, 2021.

[317] Jiangming Sun, Nina Jeliazkova, Vladimir Chupakhin, Jose-Felipe Golib-Dzib, Ola Engkvist, Lars Carlsson, Jörg Wegner, Hugo Ceulemans, Ivan Georgiev, Vedrin Jeliazkov, et al. Excape-db: an integrated large scale dataset facilitating big data analysis in chemogenomics. *Journal of cheminformatics*, 9:1–9, 2017.

[318] Zeyu Han, Chao Gao, Jinyang Liu, Jeff Zhang, and Sai Qian Zhang. Parameter-efficient fine-tuning for large models: A comprehensive survey. *arXiv [cs.LG]*, March 2024. URL http://arxiv.org/abs/2403.14608.

[319] Ning Ding, Yujia Qin, Guang Yang, Fuchao Wei, Zonghan Yang, Yusheng Su, Shengding Hu, Yulin Chen, Chi-Min Chan, Weize Chen, Jing Yi, Weilin Zhao, Xiaozhi Wang, Zhiyuan Liu, Hai-Tao Zheng, Jianfei Chen, Yang Liu, Jie Tang, Juanzi Li, and Maosong Sun. Parameter-efficient fine-tuning of large-scale pre-trained language

models. *Nat. Mach. Intell.*, 5(3):220–235, March 2023. ISSN 2522-5839,2522-5839. doi:10.1038/s42256-023-00626-4. URL https://www.nature.com/articles/s42256-023-00626-4.

[320] Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. *arXiv [cs.CL]*, June 2021. URL http://arxiv.org/abs/2106.09685.

[321] Abimael Guzman-Pando, Graciela Ramirez-Alonso, Carlos Arzate-Quintana, and Javier Camarillo-Cisneros. Deep learning algorithms applied to computational chemistry. *Molecular Diversity*, December 2023. ISSN 1573-501X. doi:10.1007/s11030-023-10771-y. URL https://doi.org/10.1007/s11030-023-10771-y.

[322] Jenna C. Fromer and Connor W. Coley. Computer-aided multi-objective optimization in small molecule discovery. *Patterns*, 4(2):100678, February 2023. ISSN 2666-3899. doi:10.1016/j.patter.2023.100678. URL https://www.sciencedirect.com/science/article/pii/S2666389923000016.

[323] Martin Vogt. Exploring chemical space — Generative models and their evaluation. *Artificial Intelligence in the Life Sciences*, 3:100064, December 2023. ISSN 2667-3185. doi:10.1016/j.ailsci.2023.100064. URL https://www.sciencedirect.com/science/article/pii/S2667318523000089.

[324] Manan Goel, Rishal Aggarwal, Bhuvanesh Sridharan, Pradeep Kumar Pal, and U. Deva Priyakumar. Efficient and enhanced sampling of drug-like chemical space for virtual screening and molecular design using modern machine learning methods. *WIREs Computational Molecular Science*, 13(2):e1637, 2023. ISSN 1759-0884. doi:10.1002/wcms.1637. URL https://onlinelibrary.wiley.com/doi/abs/10.1002/wcms.1637. _eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/wcms.1637.

[325] Shion Honda, Shoi Shi, and Hiroki R Ueda. SMILES transformer: Pre-trained molecular fingerprint for low data drug discovery. *arXiv [cs.LG]*, November 2019. URL http://arxiv.org/abs/1911.04738.

[326] Tatsuya Sagawa and Ryosuke Kojima. ReactionT5: a large-scale pre-trained model towards application of limited reaction data. *arXiv [physics.chem-ph]*, November 2023. URL http://arxiv.org/abs/2311.06708.

[327] Yin Fang, Ningyu Zhang, Zhuo Chen, Lingbing Guo, Xiaohui Fan, and Huajun Chen. Domain-agnostic molecular generation with chemical feedback. *arXiv [cs.LG]*, January 2023. URL http://arxiv.org/abs/2301.11259.

[328] Dimitrios Christofidellis, Giorgio Giannone, Jannis Born, Ole Winther, Teodoro Laino, and Matteo Manica. Unifying molecular and textual representations via multi-task language modelling. *arXiv [cs.LG]*, January 2023. URL http://arxiv.org/abs/2301.12586.

[329] Alain C Vaucher, Federico Zipoli, Joppe Geluykens, Vishnu H Nair, Philippe Schwaller, and Teodoro Laino. Automated extraction of chemical synthesis actions from experimental procedures. *Nat. Commun.*, 11(1): 3601, July 2020. ISSN 2041-1723. doi:10.1038/s41467-020-17266-6. URL http://dx.doi.org/10.1038/s41467-020-17266-6.

[330] Carl Edwards, Tuan Lai, Kevin Ros, Garrett Honke, Kyunghyun Cho, and Heng Ji. Translation between molecules and natural language. *arXiv [cs.CL]*, April 2022. URL http://arxiv.org/abs/2204.11817.

[331] Carl Edwards, ChengXiang Zhai, and Heng Ji. Text2Mol: Cross-modal molecule retrieval with natural language queries. In Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih, editors, *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 595–607, Online and Punta Cana, Dominican Republic, November 2021. Association for Computational Linguistics. doi:10.18653/v1/2021.emnlp-main.47. URL https://aclanthology.org/2021.emnlp-main.47.

[332] Ryan A. Shenvi. Natural product synthesis in the 21st century: Beyond the mountain top. *ACS Central Science*, 10(3):519–528, 2024. doi:10.1021/acscentsci.3c01518. URL https://doi.org/10.1021/acscentsci.3c01518.

[333] Qianxiang Ai, Fanwang Meng, Jiale Shi, Brenden Pelkie, and Connor W. Coley. Extracting Structured Data from Organic Synthesis Procedures Using a Fine-Tuned Large Language Model, April 2024. URL https://chemrxiv.org/engage/chemrxiv/article-details/661064e921291e5d1d2bc860.

[334] E. J. Corey. Robert Robinson Lecture. Retrosynthetic thinking—essentials and examples. *Chemical Society Reviews*, 17(0):111–133, January 1988. ISSN 1460-4744. doi:10.1039/CS9881700111. URL https://pubs.rsc.org/en/content/articlelanding/1988/cs/cs9881700111. Publisher: The Royal Society of Chemistry.

[335] Jennie B. Nerenberg, Deborah T. Hung, Patricia K. Somers, and Stuart L. Schreiber. Total synthesis of the immunosuppressive agent (-)-discodermolide. *Journal of the American Chemical Society*, 115(26): 12621–12622, 1993. doi:10.1021/ja00079a066. URL https://doi.org/10.1021/ja00079a066. _eprint: https://doi.org/10.1021/ja00079a066.

[336] Juno Nam and Jurae Kim. Linking the Neural Machine Translation and the Prediction of Organic Chemistry Reactions, December 2016. URL `http://arxiv.org/abs/1612.09529`. arXiv:1612.09529 [cs].

[337] Bowen Liu, Bharath Ramsundar, Prasad Kawthekar, Jade Shi, Joseph Gomes, Quang Luu Nguyen, Stephen Ho, Jack Sloane, Paul Wender, and Vijay Pande. Retrosynthetic Reaction Prediction Using Neural Sequence-to-Sequence Models. *ACS Central Science*, 3(10):1103–1113, October 2017. ISSN 2374-7943. doi:10.1021/acscentsci.7b00303. URL `https://doi.org/10.1021/acscentsci.7b00303`. Publisher: American Chemical Society.

[338] Nadine Schneider, Nikolaus Stiefl, and Gregory A. Landrum. What's what: The (nearly) definitive guide to reaction role assignment. *Journal of Chemical Information and Modeling*, 56(12):2336–2346, 2016. doi:10.1021/acs.jcim.6b00564. URL `https://doi.org/10.1021/acs.jcim.6b00564`. PMID: 28024398.

[339] Matthew Gunther10 August 2016. Software could revolutionise chemistry. URL `https://www.chemistryworld.com/news/software-could-revolutionise-chemistry/1017236.article`.

[340] Tomasz Klucznik, Barbara Mikulak-Klucznik, Michael P. McCormack, Heather Lima, Sara Szymkuć, Manishabrata Bhowmick, Karol Molga, Yubai Zhou, Lindsey Rickershauser, Ewa P. Gajewska, Alexei Toutchkine, Piotr Dittwald, Michał P. Startek, Gregory J. Kirkovits, Rafał Roszak, Ariel Adamski, Bianka Sieredzińska, Milan Mrksich, Sarah L. J. Trice, and Bartosz A. Grzybowski. Efficient Syntheses of Diverse, Medicinally Relevant Targets Planned by Computer and Executed in the Laboratory. *Chem*, 4(3):522–532, March 2018. ISSN 2451-9294, 2451-9308. doi:10.1016/j.chempr.2018.02.002. URL `https://www.cell.com/chem/abstract/S2451-9294(18)30063-9`. Publisher: Elsevier.

[341] Marwin H S Segler and Mark P Waller. Neural-symbolic machine learning for retrosynthesis and reaction prediction. *Chemistry*, 23(25):5966–5971, May 2017. ISSN 0947-6539, 1521-3765. doi:10.1002/chem.201605499. URL `https://chemistry-europe.onlinelibrary.wiley.com/doi/10.1002/chem.201605499`.

[342] Vignesh Ram Somnath, Charlotte Bunne, Connor W. Coley, Andreas Krause, and Regina Barzilay. Learning Graph Models for Retrosynthesis Prediction, June 2021. URL `http://arxiv.org/abs/2006.07038`. arXiv:2006.07038 [cs, stat].

[343] Geemi P Wellawatte and Philippe Schwaller. Extracting human interpretable structure-property relationships in chemistry using XAI and large language models. *arXiv [physics.chem-ph]*, November 2023. URL `http://arxiv.org/abs/2311.04047`.

[344] Andrea Cadeddu, Elizabeth K. Wylie, Janusz Jurczak, Matthew Wampler-Doty, and Bartosz A. Grzybowski. Organic Chemistry as a Language and the Implications of Chemical Linguistics for Structural and Retrosynthetic Analyses. *Angewandte Chemie International Edition*, 53(31):8108–8112, 2014. ISSN 1521-3773. doi:10.1002/anie.201403708. URL `https://onlinelibrary.wiley.com/doi/abs/10.1002/anie.201403708`. _eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/anie.201403708.

[345] Philippe Schwaller, Théophile Gaudin, Dávid Lányi, Costas Bekas, and Teodoro Laino. "Found in Translation": predicting outcomes of complex organic chemistry reactions using neural sequence-to-sequence models. *Chemical Science*, 9(28):6091–6098, July 2018. ISSN 2041-6539. doi:10.1039/C8SC02339E. URL `https://pubs.rsc.org/en/content/articlelanding/2018/sc/c8sc02339e`. Publisher: The Royal Society of Chemistry.

[346] Wengong Jin, Connor W. Coley, Regina Barzilay, and Tommi Jaakkola. Predicting Organic Reaction Outcomes with Weisfeiler-Lehman Network, December 2017. URL `http://arxiv.org/abs/1709.04555`. arXiv:1709.04555 [cs, stat].

[347] John Bradshaw, Matt J. Kusner, Brooks Paige, Marwin H. S. Segler, and José Miguel Hernández-Lobato. A Generative Model For Electron Paths, March 2019. URL `http://arxiv.org/abs/1805.10970`. arXiv:1805.10970 [physics, stat].

[348] Philippe Schwaller, Riccardo Petraglia, Valerio Zullo, Vishnu H. Nair, Rico Andreas Haeuselmann, Riccardo Pisoni, Costas Bekas, Anna Iuliano, and Teodoro Laino. Predicting retrosynthetic pathways using transformer-based models and a hyper-graph exploration strategy. *Chemical Science*, 11(12):3316–3325, March 2020. ISSN 2041-6539. doi:10.1039/C9SC05704H. URL `https://pubs.rsc.org/en/content/articlelanding/2020/sc/c9sc05704h`. Publisher: The Royal Society of Chemistry.

[349] Shuan Chen and Yousung Jung. Deep Retrosynthetic Reaction Prediction using Local Reactivity and Global Attention. *JACS Au*, 1(10):1612–1620, August 2021. ISSN 2691-3704. doi:10.1021/jacsau.1c00246. URL `https://doi.org/10.1021/jacsau.1c00246`. Publisher: American Chemical Society.

[350] Annie M. Westerlund, Siva Manohar Koki, Supriya Kancharla, Alessandro Tibo, Lakshidaa Saigiridharan, Mikhail Kabeshov, Rocio Mercado, and Samuel Genheden. Do Chemformers Dream of Organic

Matter? Evaluating a Transformer Model for Multistep Retrosynthesis. *JOURNAL OF CHEMICAL INFORMATION AND MODELING*, 64(8):3021–3033, April 2024. ISSN 1549-9596, 1549-960X. doi:10.1021/acs.jcim.3c01685. URL `https://www.webofscience.com/api/gateway?GWVersion=2&SrcAuth=DynamicDOIArticle&SrcApp=UA&KeyAID=10.1021%2Facs.jcim.3c01685&DestApp=DOI&SrcAppSID=USW2EC0C6Dsgh6m89jbQCLUMIbTAg&SrcJTitle=JOURNAL+OF+CHEMICAL+INFORMATION+AND+MODELING&DestDOIRegistrantName=American+Chemical+Society`. Num Pages: 13 Place: Washington Publisher: Amer Chemical Soc Web of Science ID: WOS:001201284600001.

[351] Shuangjia Zheng, Jiahua Rao, Zhongyue Zhang, Jun Xu, and Yuedong Yang. Predicting Retrosynthetic Reactions Using Self-Corrected Transformer Neural Networks. *Journal of Chemical Information and Modeling*, 60(1): 47–55, January 2020. ISSN 1549-9596. doi:10.1021/acs.jcim.9b00949. URL `https://doi.org/10.1021/acs.jcim.9b00949`. Publisher: American Chemical Society.

[352] Daniel Mark Lowe. Extraction of chemical structures and reactions from the literature. October 2012. URL `http://www.dspace.cam.ac.uk/handle/1810/244727`.

[353] Junren Li, Lei Fang, and Jian-Guang Lou. Retro-BLEU: quantifying chemical plausibility of retrosynthesis routes through reaction template sequence analysis. *Digital Discovery*, 3(3):482–490, 2024. doi:10.1039/D3DD00219E. URL `https://pubs.rsc.org/en/content/articlelanding/2024/dd/d3dd00219e`. Publisher: Royal Society of Chemistry.

[354] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. BLEU: a method for automatic evaluation of machine translation. In Pierre Isabelle, Eugene Charniak, and Dekang Lin, editors, *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics - ACL '02*, pages 311–318, Philadelphia, Pennsylvania, USA, July 2002. Association for Computational Linguistics. doi:10.3115/1073083.1073135. URL `http://portal.acm.org/citation.cfm?doid=1073083.1073135`.

[355] Chin-Yew Lin. ROUGE: A Package for Automatic Evaluation of Summaries. In *Text Summarization Branches Out*, pages 74–81, Barcelona, Spain, July 2004. Association for Computational Linguistics. URL `https://aclanthology.org/W04-1013`.

[356] Laurianne David, Amol Thakkar, Rocío Mercado, and Ola Engkvist. Molecular representations in AI-driven drug discovery: a review and practical guide. *J. Cheminform.*, 12(1):56, September 2020. ISSN 1758-2946. doi:10.1186/s13321-020-00460-5. URL `https://doi.org/10.1186/s13321-020-00460-5`.

[357] Gus L W Hart, Tim Mueller, Cormac Toher, and Stefano Curtarolo. Machine learning for alloys. *Nat. Rev. Mater.*, 6(8):730–755, July 2021. ISSN 2058-8437,2058-8437. doi:10.1038/s41578-021-00340-w. URL `https://www.nature.com/articles/s41578-021-00340-w`.

[358] Oscar Eraso, Daniela Bolaños, Nikolas Echeverri, Carolina Orozco Donneys, Tayebeh Ameri, and Jose Dario Perea. A present scenario of the computational approaches for ternary organic solar cells. *J. Renew. Sustain. Energy*, 15(6):062702, November 2023. ISSN 1941-7012. doi:10.1063/5.0172426. URL `https://pubs.aip.org/aip/jrse/article-pdf/doi/10.1063/5.0172426/18235425/062702_1_5.0172426.pdf`.

[359] Yu-Chen Lo, Stefano E Rensi, Wen Torng, and Russ B Altman. Machine learning in chemoinformatics and drug discovery. *Drug Discov. Today*, 23(8):1538–1546, August 2018. ISSN 1359-6446,1878-5832. doi:10.1016/j.drudis.2018.05.010. URL `http://dx.doi.org/10.1016/j.drudis.2018.05.010`.

[360] Daniel Flam-Shepherd and Alán Aspuru-Guzik. Language models can generate molecules, materials, and protein binding sites directly in three dimensions as XYZ, CIF, and PDB files. *arXiv [cs.LG]*, May 2023. URL `http://arxiv.org/abs/2305.05708`.

[361] Kohulan Rajan, Achim Zielesny, and Christoph Steinbeck. DECIMER 1.0: deep learning for chemical image recognition using transformers. *J. Cheminform.*, 13(1):61, August 2021. ISSN 1758-2946,1758-2946. doi:10.1186/s13321-021-00538-8. URL `http://dx.doi.org/10.1186/s13321-021-00538-8`.

[362] J. T. Carstensen and Faraneh Attarchi. Decomposition of Aspirin in the Solid State in the Presence of Limited Amounts of Moisture III: Effect of Temperature and a Possible Mechanism. *Journal of Pharmaceutical Sciences*, 77(4):318–321, April 1988. ISSN 0022-3549. doi:10.1002/jps.2600770407. URL `https://www.sciencedirect.com/science/article/pii/S0022354915476658`.

[363] Philipp Seidl, Andreu Vall, Sepp Hochreiter, and Guenter Klambauer. Enhancing Activity Prediction Models in Drug Discovery with the Ability to Understand Human Language. *Arxiv*, 2023.

[364] Hanwen Xu, Addie Woicik, Hoifung Poon, Russ B. Altman, and Sheng Wang. Multilingual translation for zero-shot biomedical classification using BioTranslator. *NATURE COMMUNICATIONS*, 14(1), February 2023. ISSN 2041-1723. doi:10.1038/s41467-023-36476-2.

[365] Shengchao Liu, Jiongxiao Wang, Yijin Yang, Chengpeng Wang, Ling Liu, Hongyu Guo, and Chaowei Xiao. ChatGPT-powered Conversational Drug Editing Using Retrieval and Domain Feedback. *Arxiv*, 2023.

[366] Zequn Liu, Wei Zhang, Yingce Xia, Lijun Wu, Shufang Xie, Tao Qin, Ming Zhang, and Tie-Yan Liu. MolXPT: Wrapping Molecules with Text for Generative Pre-training. *Arxiv*, 2023.

[367] Pengfei Liu, Yiming Ren, Jun Tao, and Zhixiang Ren. Git-mol: A multi-modal large language model for molecular science with graph, image, and text. *Computers in Biology and Medicine*, 171:108073, March 2024. ISSN 0010-4825. doi:10.1016/j.compbiomed.2024.108073. URL http://dx.doi.org/10.1016/j.compbiomed.2024.108073.

[368] Junfeng Fang, Shuai Zhang, Chang Wu, Zhiyuan Liu, Sihang Li, Kun Wang, Wenjie Du, and Xiang Wang. MolTC: Towards Molecular Relational Modeling In Language Models. *Arxiv*, 2024.

[369] Haohui Zhang, Juntong Wu, Shichao Liu, and Shen Han. A pre-trained multi-representation fusion network for molecular property prediction. *INFORMATION FUSION*, 103, March 2024. ISSN 1566-2535. doi:10.1016/j.inffus.2023.102092.

[370] Huaisheng Zhu, Teng Xiao, and Vasant G Honavar. 3M-Diffusion: Latent Multi-Modal Diffusion for Text-Guided Generation of Molecular Graphs. *Arxiv*, 2024.

[371] Changnan Gao, Wenjie Bao, Shuang Wang, Jianyang Zheng, Lulu Wang, Yongqi Ren, Linfang Jiao, Jianmin Wang, and Xun Wang. DockingGA: enhancing targeted molecule generation using transformer neural network and genetic algorithm with docking simulation. *BRIEFINGS IN FUNCTIONAL GENOMICS*, April 2024. ISSN 2041-2649. doi:10.1093/bfgp/elae011.

[372] Peng Zhou, Jianmin Wang, Chunyan Li, Zixu Wang, Yiping Liu, Siqi Sun, Jianxin Lin, Longyue Wang, and Xiangxiang Zeng. Instruction Multi-Constraint Molecular Generation Using a Teacher-Student Large Language Model. *Arxiv*, 2024.

[373] Haisong Gong, Qiang Liu, Shu Wu, and Liang Wang. Text-Guided Molecule Generation with Diffusion Language Model. *Arxiv*, 2024.

[374] Eduardo Soares, Emilio Vital Brazil, Karen Fiorela Aquino Gutierrez, Renato Cerqueira, Dan Sanders, Kristin Schmidt, and Dmitry Zubarev. Beyond Chemical Language: A Multimodal Approach to Enhance Molecular Property Prediction, June 2023. URL http://arxiv.org/abs/2306.14919. arXiv:2306.14919 [physics, q-bio].

[375] Michael Riedl, Sayak Mukherjee, and Mitch Gauthier. Descriptor-free deep learning QSAR model for the fraction unbound in human plasma. *Mol. Pharm.*, 20(10):4984–4993, October 2023. ISSN 1543-8384,1543-8392. doi:10.1021/acs.molpharmaceut.3c00129. URL http://dx.doi.org/10.1021/acs.molpharmaceut.3c00129.

[376] MIMIC-III documentation. https://mimic.mit.edu/docs/iii/, 2021. URL https://mimic.mit.edu/docs/iii/. Accessed: 2024-3-25.

[377] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. RoBERTa: A robustly optimized BERT pretraining approach. *arXiv [cs.CL]*, July 2019. URL http://arxiv.org/abs/1907.11692.

[378] Hannah Smith, Zeyu Zhang, John Culnan, and Peter Jansen. ScienceExamCER: A high-density fine-grained science-domain corpus for common entity recognition. *arXiv [cs.CL]*, November 2019. URL http://arxiv.org/abs/1911.10436.

[379] Matthew C Swain and Jacqueline M Cole. ChemDataExtractor: A toolkit for automated extraction of chemical information from the scientific literature. *J. Chem. Inf. Model.*, 56(10):1894–1904, October 2016. ISSN 1549-9596, 1549-960X. doi:10.1021/acs.jcim.6b00207. URL http://dx.doi.org/10.1021/acs.jcim.6b00207.

[380] Santiago Ibanez. chemie-turk: Mechanical turk on your own machine for chemical literature annotation. URL https://github.com/asibanez/chemie-turk.

[381] Wei Zhang, Qinggong Wang, Xiangtai Kong, Jiacheng Xiong, Shengkun Ni, Duanhua Cao, Buying Niu, Mingan Chen, Runze Zhang, Yitian Wang, Lehan Zhang, Xutong Li, Zhaoping Xiong, Qian Shi, Ziming Huang, Zunyun Fu, and Mingyue Zheng. Fine-tuning large language models for chemical text mining. *ChemRxiv*, February 2024. doi:10.26434/chemrxiv-2023-k7ct5-v2. URL https://chemrxiv.org/engage/chemrxiv/article-details/65baa07b9138d2316124f224.

[382] Kexin Chen, Hanqun Cao, Junyou Li, Yuyang Du, Menghao Guo, Xin Zeng, Lanqing Li, Jiezhong Qiu, Pheng Ann Heng, and Guangyong Chen. An autonomous large language model agent for chemical literature data mining. *arXiv [cs.IR]*, February 2024. URL http://arxiv.org/abs/2402.12993.

[383] Xin Wang, Liangliang Huang, Shuozhi Xu, and Kun Lu. How does a generative large language model perform on domain-specific information extraction? – a comparison between GPT-4 and a rule-based method on band gap extraction. *J. Chem. Inf. Model.*, October 2024. ISSN 1549-9596,1549-960X. doi:10.1021/acs.jcim.4c00882. URL `https://pubs.acs.org/doi/full/10.1021/acs.jcim.4c00882`.

[384] Mara Schilling-Wilhelmi, Martiño Ríos-García, Sherjeel Shabih, María Victoria Gil, Santiago Miret, Christoph T Koch, José A Márquez, and Kevin Maik Jablonka. From text to insight: Large language models for materials science data extraction. *arXiv [cond-mat.mtrl-sci]*, July 2024. URL `http://arxiv.org/abs/2407.16867`.

[385] Pranav Shetty, Arunkumar Chitteth Rajan, Chris Kuenneth, Sonakshi Gupta, Lakshmi Prerana Panchumarti, Lauren Holm, Chao Zhang, and Rampi Ramprasad. A general-purpose material property data extraction pipeline from large polymer corpora using natural language processing. *NPJ Comput Mater*, 9(1):52, April 2023. ISSN 2057-3960. doi:10.1038/s41524-023-01003-w. URL `http://dx.doi.org/10.1038/s41524-023-01003-w`.

[386] Mohammad Shoeybi, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. Megatron-LM: Training multi-billion parameter language models using model parallelism. *arXiv [cs.CL]*, September 2019. URL `http://arxiv.org/abs/1909.08053`.

[387] Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. SQuAD: 100,000+ questions for machine comprehension of text. *arXiv [cs.CL]*, June 2016. URL `http://arxiv.org/abs/1606.05250`.

[388] Yifan Peng, Qingyu Chen, and Zhiyong Lu. An empirical study of multi-task learning on BERT for biomedical text mining. *arXiv [cs.CL]*, May 2020. URL `http://arxiv.org/abs/2005.02799`.

[389] Johannes Welbl, Nelson F Liu, and Matt Gardner. Crowdsourcing multiple choice science questions. *arXiv [cs.HC]*, July 2017. URL `http://arxiv.org/abs/1707.06209`.

[390] Yu Song, Santiago Miret, Huan Zhang, and Bang Liu. HoneyBee: Progressive instruction finetuning of large language models for materials science. *arXiv [cs.CL]*, October 2023. URL `http://arxiv.org/abs/2310.08511`.

[391] Huan Zhang, Yu Song, Ziyu Hou, Santiago Miret, and Bang Liu. HoneyComb: A flexible LLM-based agent system for materials science. *arXiv [cs.CL]*, August 2024. URL `http://arxiv.org/abs/2409.00135`.

[392] Sören Auer, Dante A C Barone, Cassiano Bartz, Eduardo G Cortes, Mohamad Yaser Jaradeh, Oliver Karras, Manolis Koubarakis, Dmitry Mouromtsev, Dmitrii Pliukhin, Daniil Radyush, Ivan Shilin, Markus Stocker, and Eleni Tsalapati. The SciQA scientific question answering benchmark for scholarly knowledge. *Sci. Rep.*, 13(1):7240, May 2023. ISSN 2045-2322,2045-2322. doi:10.1038/s41598-023-33607-z. URL `https://www.nature.com/articles/s41598-023-33607-z`.

[393] Di Jin, Eileen Pan, Nassim Oufattole, Wei-Hung Weng, Hanyi Fang, and Peter Szolovits. What disease does this patient have? a large-scale open domain question answering dataset from medical exams. *Preprints*, May 2021. URL `http://arxiv.org/abs/2009.13081`.

[394] Ankit Pal, Logesh Kumar Umapathi, and Malaikannan Sankarasubbu. MedMCQA : A large-scale multi-subject multi-choice dataset for medical domain question answering. *arXiv [cs.CL]*, March 2022. URL `http://arxiv.org/abs/2203.14371`.

[395] Cayque Monteiro Castro Nascimento and André Silva Pimentel. Do Large Language Models Understand Chemistry? A Conversation with ChatGPT. *Journal of Chemical Information and Modeling*, 63(6):1649–1655, March 2023. ISSN 1549-9596. doi:10.1021/acs.jcim.3c00285. URL `https://doi.org/10.1021/acs.jcim.3c00285`. Publisher: American Chemical Society.

[396] Tim Humphry and Amy L. Fuller. Potential ChatGPT Use in Undergraduate Chemistry Laboratories. *JOURNAL OF CHEMICAL EDUCATION*, 100(4):1434–1436, April 2023. ISSN 0021-9584. doi:10.1021/acs.jchemed.3c00006.

[397] Mary E. Emenike and Bright U. Emenike. Was This Title Generated by ChatGPT? Considerations for Artificial Intelligence Text-Generation Software Programs for Chemists and Chemistry Educators. *JOURNAL OF CHEMICAL EDUCATION*, 100(4):1413–1418, April 2023. ISSN 0021-9584. doi:10.1021/acs.jchemed.3c00063.

[398] Suzanne Fergus, Michelle Botha, and Mehrnoosh Ostovar. Evaluating Academic Answers Generated Using ChatGPT. *JOURNAL OF CHEMICAL EDUCATION*, 100(4):1672–1675, April 2023. ISSN 0021-9584. doi:10.1021/acs.jchemed.3c00087.

[399] Zhiling Zheng, Ali H. Alawadhi, Saumil Chheda, S. Ephraim Neumann, Nakul Rampal, Shengchao Liu, Ha L. Nguyen, Yen-hsu Lin, Zichao Rong, J. Ilja Siepmann, Laura Gagliardi, Anima Anandkumar, Christian Borgs, Jennifer T. Chayes, and Omar M. Yaghi. Shaping the Water-Harvesting Behavior of Metal-Organic Frameworks Aided by Fine-Tuned GPT Models. *JOURNAL OF THE AMERICAN CHEMICAL SOCIETY*, 145(51):28284–28295, December 2023. ISSN 0002-7863. doi:10.1021/jacs.3c12086.

[400] Zhiling Zheng, Oufan Zhang, Christian Borgs, Jennifer T. Chayes, and Omar M. Yaghi. ChatGPT Chemistry Assistant for Text Mining and the Prediction of MOF Synthesis. *Journal of the American Chemical Society*, 145 (Copyright © 2024 American Chemical Society (ACS). All Rights Reserved.; Copyright © 2024 U.S. National Library of Medicine.):18048–18062, 2023. ISSN 1520-5126. doi:10.1021/jacs.3c05819. Publisher: American Chemical Society.

[401] Zikai Xie, Xenophon Evangelopoulos, Ömer H. Omar, Alessandro Troisi, Andrew I. Cooper, and Linjiang Chen. Fine-tuning GPT-3 for machine learning electronic and functional properties of organic molecules. *Chemical Science*, 15(2):500–510, January 2024. ISSN 2041-6539. doi:10.1039/D3SC04610A. URL `https://pubs.rsc.org/en/content/articlelanding/2024/sc/d3sc04610a`. Publisher: The Royal Society of Chemistry.

[402] Zhiling Zheng, Zichao Rong, Nakul Rampal, Christian Borgs, Jennifer T. Chayes, and Omar M. Yaghi. A GPT-4 Reticular Chemist for Guiding MOF Discovery. *ANGEWANDTE CHEMIE-INTERNATIONAL EDITION*, 62 (46), November 2023. ISSN 1433-7851. doi:10.1002/anie.202311983.

[403] Jyotirmoy Deb, Lakshi Saikia, Kripa Dristi Dihingia, and G. Narahari Sastry. Chatgpt in the material design: Selected case studies to assess the potential of chatgpt. *Journal of Chemical Information and Modeling*, 64 (3):799–811, 2024. doi:10.1021/acs.jcim.3c01702. URL `https://doi.org/10.1021/acs.jcim.3c01702`. PMID: 38237025.

[404] Benjamin S Bloom. *Taxonomy of Educational Objectives: The Classification of Educational Goals; Handbook. Cognitive Domain*. McKay, 1968.

[405] Benjamin Samuel Bloom. *A taxonomy for learning, teaching, and assessing: A revision of Bloom's taxonomy of educational objectives*. Longman, 2010.

[406] Connor W. Coley, Dale A. Thomas, Justin A. M. Lummiss, Jonathan N. Jaworski, Christopher P. Breen, Victor Schultz, Travis Hart, Joshua S. Fishman, Luke Rogers, Hanyu Gao, Robert W. Hicklin, Pieter P. Plehiers, Joshua Byington, John S. Piotti, William H. Green, A. John Hart, Timothy F. Jamison, and Klavs F. Jensen. A robotic platform for flow synthesis of organic compounds informed by AI planning. *Science*, 365(6453): eaax1566, August 2019. doi:10.1126/science.aax1566. URL `https://www.science.org/doi/full/10.1126/science.aax1566`. Publisher: American Association for the Advancement of Science.

[407] Piotr S. Gromski, Jarosław M. Granda, and Leroy Cronin. Universal Chemical Synthesis and Discovery with 'The Chemputer'. *Trends in Chemistry*, 2(1):4–12, January 2020. ISSN 25895974. doi:10.1016/j.trechm.2019.07.004. URL `https://linkinghub.elsevier.com/retrieve/pii/S2589597419301868`.

[408] Francesca Grisoni, Berend J. H. Huisman, Alexander L. Button, Michael Moret, Kenneth Atz, Daniel Merk, and Gisbert Schneider. Combining generative artificial intelligence and on-chip synthesis for de novo drug design. *Science Advances*, 7(24):eabg3338, June 2021. doi:10.1126/sciadv.abg3338. URL `https://www.science.org/doi/full/10.1126/sciadv.abg3338`. Publisher: American Association for the Advancement of Science.

[409] Brian Goldman, Steven Kearnes, Trevor Kramer, Patrick Riley, and W. Patrick Walters. Defining Levels of Automated Chemical Design. *Journal of Medicinal Chemistry*, 65(10):7073–7087, May 2022. ISSN 0022-2623, 1520-4804. doi:10.1021/acs.jmedchem.2c00334. URL `https://pubs.acs.org/doi/10.1021/acs.jmedchem.2c00334`.

[410] Gisbert Schneider. Automating drug discovery. *Nature Reviews Drug Discovery*, 17(2):97–113, February 2018. ISSN 1474-1784. doi:10.1038/nrd.2017.232. URL `https://doi.org/10.1038/nrd.2017.232`.

[411] Jon Paul Janet, Lewis Mervin, and Ola Engkvist. Artificial intelligence in molecular *de novo* design: Integration with experiment. *Current Opinion in Structural Biology*, 80:102575, June 2023. ISSN 0959-440X. doi:10.1016/j.sbi.2023.102575. URL `https://www.sciencedirect.com/science/article/pii/S0959440X23000490`.

[412] Connor W. Coley, Natalie S. Eyke, and Klavs F. Jensen. Autonomous Discovery in the Chemical Sciences Part I: Progress. *Angewandte Chemie International Edition*, 59(51):22858–22893, 2020. ISSN 1521-3773. doi:10.1002/anie.201909987. URL `https://onlinelibrary.wiley.com/doi/abs/10.1002/anie.201909987`. _eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/anie.201909987.

[413] Connor W. Coley, Natalie S. Eyke, and Klavs F. Jensen. Autonomous Discovery in the Chemical Sciences Part II: Outlook. *Angewandte Chemie International Edition*, 59(52):23414–23436, 2020. ISSN 1521-3773. doi:10.1002/anie.201909989. URL `https://onlinelibrary.wiley.com/doi/abs/10.1002/anie.201909989`. _eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/anie.201909989.

[414] Amol Thakkar, Simon Johansson, Kjell Jorner, David Buttar, Jean-Louis Reymond, and Ola Engkvist. Artificial intelligence and automation in computer aided synthesis planning. *Reaction Chemistry & Engineering*, 6(1): 27–51, January 2021. ISSN 2058-9883. doi:10.1039/D0RE00340A. URL https://pubs.rsc.org/en/content/articlelanding/2021/re/d0re00340a. Publisher: The Royal Society of Chemistry.

[415] Yuning Shen, Julia E. Borowski, Melissa A. Hardy, Richmond Sarpong, Abigail G. Doyle, and Tim Cernak. Automation and computer-assisted planning for chemical synthesis. *Nature Reviews Methods Primers*, 1(1): 23, March 2021. ISSN 2662-8449. doi:10.1038/s43586-021-00022-5. URL https://www.nature.com/articles/s43586-021-00022-5.

[416] Yuxiao Liu, Lingyu Sun, Hui Zhang, Luoran Shang, and Yuanjin Zhao. Microfluidics for Drug Development: From Synthesis to Evaluation. *Chemical Reviews*, 121(13):7468–7529, July 2021. ISSN 0009-2665. doi:10.1021/acs.chemrev.0c01289. URL https://doi.org/10.1021/acs.chemrev.0c01289. Publisher: American Chemical Society.

[417] Kourosh Darvish, Marta Skreta, Yuchi Zhao, Naruki Yoshikawa, Sagnik Som, Miroslav Bogdanovic, Yang Cao, Han Hao, Haoping Xu, Alán Aspuru-Guzik, Animesh Garg, and Florian Shkurti. ORGANA: A robotic assistant for automated chemistry experimentation and characterization. *arXiv [cs.RO]*, January 2024. URL http://arxiv.org/abs/2401.06949.

[418] Tomáš Šalamon. *Design of Agent-based Models: Developing Computer Simulations for a Better Understanding of Social Processes*. Tomáš Bruckner, Repin, CZE, 2011. ISBN 9788090466111. URL https://play.google.com/store/books/details?id=2rCdKnltaH8C.

[419] Zhiheng Xi, Wenxiang Chen, Xin Guo, Wei He, Yiwen Ding, Boyang Hong, Ming Zhang, Junzhe Wang, Senjie Jin, Enyu Zhou, Rui Zheng, Xiaoran Fan, Xiao Wang, Limao Xiong, Yuhao Zhou, Weiran Wang, Changhao Jiang, Yicheng Zou, Xiangyang Liu, Zhangyue Yin, Shihan Dou, Rongxiang Weng, Wensen Cheng, Qi Zhang, Wenjuan Qin, Yongyan Zheng, Xipeng Qiu, Xuanjing Huang, and Tao Gui. The rise and potential of large language model based agents: A survey. *arXiv [cs.AI]*, September 2023. URL http://arxiv.org/abs/2309.07864.

[420] Shanghua Gao, Ada Fang, Yepeng Huang, Valentina Giunchiglia, Ayush Noori, Jonathan Richard Schwarz, Yasha Ektefaie, Jovana Kondic, and Marinka Zitnik. Empowering biomedical discovery with AI agents. *arXiv [cs.AI]*, April 2024. URL http://arxiv.org/abs/2404.02831.

[421] Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Jirong Wen. A survey on large language model based autonomous agents. *Front. Comput. Sci.*, 18(6), December 2024. ISSN 2095-2228,2095-2236. doi:10.1007/s11704-024-40231-1. URL http://dx.doi.org/10.1007/s11704-024-40231-1.

[422] Lilian Weng. LLM powered autonomous agents. https://lilianweng.github.io/posts/2023-06-23-agent/, June 2023. URL https://lilianweng.github.io/posts/2023-06-23-agent/. Accessed: 2024-1-22.

[423] Theodore R Sumers, Shunyu Yao, Karthik Narasimhan, and Thomas L Griffiths. Cognitive architectures for language agents. *arXiv [cs.AI]*, September 2023. URL http://arxiv.org/abs/2309.02427.

[424] Bing Wang, Xinnian Liang, Jian Yang, Hui Huang, Shuangzhi Wu, Peihao Wu, Lu Lu, Zejun Ma, and Zhoujun Li. Enhancing large language model with self-controlled memory framework. *arXiv [cs.CL]*, April 2023. URL http://arxiv.org/abs/2304.13343.

[425] Yi Zhang, Zhongyang Yu, Wanqi Jiang, Yufeng Shen, and Jin Li. Long-term memory for large language models through topic-based vector database. In *2023 International Conference on Asian Language Processing (IALP)*. IEEE, November 2023. doi:10.1109/ialp61005.2023.10337079. URL https://ieeexplore.ieee.org/document/10337079/.

[426] Xizhou Zhu, Yuntao Chen, Hao Tian, Chenxin Tao, Weijie Su, Chenyu Yang, Gao Huang, Bin Li, Lewei Lu, Xiaogang Wang, Yu Qiao, Zhaoxiang Zhang, and Jifeng Dai. Ghost in the minecraft: Generally capable agents for open-world environments via large language models with text-based knowledge and memory. *arXiv [cs.AI]*, May 2023. URL http://arxiv.org/abs/2305.17144.

[427] Wanjun Zhong, Lianghong Guo, Qiqi Gao, He Ye, and Yanlin Wang. MemoryBank: Enhancing large language models with long-term memory. *arXiv [cs.CL]*, May 2023. URL http://arxiv.org/abs/2305.10250.

[428] Yikun Han, Chunjiang Liu, and Pengfei Wang. A comprehensive survey on vector database: Storage and retrieval technique, challenge. *arXiv [cs.DB]*, October 2023. URL http://arxiv.org/abs/2310.11703.

[429] Andrew Zhao, Daniel Huang, Quentin Xu, Matthieu Lin, Yong-Jin Liu, and Gao Huang. ExpeL: LLM agents are experiential learners. *arXiv [cs.LG]*, August 2023. URL http://arxiv.org/abs/2308.10144.

[430] ANN-benchmarks. https://ann-benchmarks.com/. URL https://ann-benchmarks.com/. Accessed: 2024-2-1.

[431] Kostas Hatalis, Despina Christou, Joshua Myers, Steven Jones, Keith Lambert, Adam Amos-Binks, Zohreh Dannenhauer, and Dustin Dannenhauer. Memory matters: The need to improve Long-Term memory in LLM-Agents. *AAAI-SS*, 2(1):277–280, 2023. ISSN 2994-4317, 2994-4317. doi:10.1609/aaaiss.v2i1.27688. URL https://ojs.aaai.org/index.php/AAAI-SS/article/view/27688.

[432] Joon Sung Park, Joseph C O'Brien, Carrie J Cai, Meredith Ringel Morris, Percy Liang, and Michael S Bernstein. Generative agents: Interactive simulacra of human behavior. *arXiv [cs.HC]*, April 2023. URL http://arxiv.org/abs/2304.03442.

[433] S S Raman, V Cohen, E Rosen, and I Idrees. Planning with large language models via corrective re-prompting. *Foundation Models*, November 2022. URL https://openreview.net/pdf?id=cMDMRBe1TKs.

[434] Shehzaad Dhuliawala, Mojtaba Komeili, Jing Xu, Roberta Raileanu, Xian Li, Asli Celikyilmaz, and Jason Weston. Chain-of-verification reduces hallucination in large language models. *arXiv [cs.CL]*, September 2023. URL http://arxiv.org/abs/2309.11495.

[435] Wenlong Huang, Fei Xia, Ted Xiao, Harris Chan, Jacky Liang, Pete Florence, Andy Zeng, Jonathan Tompson, Igor Mordatch, Yevgen Chebotar, Pierre Sermanet, Noah Brown, Tomas Jackson, Linda Luu, Sergey Levine, Karol Hausman, and Brian Ichter. Inner monologue: Embodied reasoning through planning with language models. *arXiv [cs.RO]*, July 2022. URL http://arxiv.org/abs/2207.05608.

[436] Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large language models are zero-shot reasoners. *arXiv [cs.CL]*, pages 22199–22213, May 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/file/8bb0d291acd4acf06ef112099c16f326-Paper-Conference.pdf.

[437] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, E Chi, F Xia, Quoc Le, and Denny Zhou. Chain of thought prompting elicits reasoning in large language models. *Neural Inf Process Syst*, abs/2201.11903, January 2022. URL http://arxiv.org/abs/2201.11903.

[438] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. *arXiv [cs.CL]*, March 2022. URL http://arxiv.org/abs/2203.11171.

[439] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. *arXiv [cs.CL]*, October 2022. URL http://arxiv.org/abs/2210.03629.

[440] Shibo Hao, Yi Gu, Haodi Ma, Joshua Jiahua Hong, Zhen Wang, Daisy Zhe Wang, and Zhiting Hu. Reasoning with language model is planning with world model. *arXiv [cs.CL]*, May 2023. URL http://arxiv.org/abs/2305.14992.

[441] Hao Liu, Carmelo Sferrazza, and Pieter Abbeel. Chain of hindsight aligns language models with feedback. *arXiv [cs.LG]*, February 2023. URL http://arxiv.org/abs/2302.02676.

[442] Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. *arXiv [cs.CL]*, May 2023. URL http://arxiv.org/abs/2305.10601.

[443] Jikun Kang, Romain Laroche, Xingdi Yuan, Adam Trischler, Xue Liu, and Jie Fu. Think before you act: Decision transformers with working memory. *arXiv [cs.LG]*, May 2023. URL http://arxiv.org/abs/2305.16338.

[444] Cheng Qian, Shihao Liang, Yujia Qin, Yining Ye, Xin Cong, Yankai Lin, Yesai Wu, Zhiyuan Liu, and Maosong Sun. Investigate-consolidate-exploit: A general strategy for inter-task agent self-evolution. *arXiv [cs.CL]*, January 2024. URL http://arxiv.org/abs/2401.13996.

[445] Rohan Anil, Andrew M Dai, Orhan Firat, Melvin Johnson, Dmitry Lepikhin, Alexandre Passos, Siamak Shakeri, Emanuel Taropa, Paige Bailey, Zhifeng Chen, Eric Chu, Jonathan H Clark, Laurent El Shafey, Yanping Huang, Kathy Meier-Hellstern, Gaurav Mishra, Erica Moreira, Mark Omernick, Kevin Robinson, Sebastian Ruder, Yi Tay, Kefan Xiao, Yuanzhong Xu, Yujing Zhang, Gustavo Hernandez Abrego, Junwhan Ahn, Jacob Austin, Paul Barham, Jan Botha, James Bradbury, Siddhartha Brahma, Kevin Brooks, Michele Catasta, Yong Cheng, Colin Cherry, Christopher A Choquette-Choo, Aakanksha Chowdhery, Clément Crepy, Shachi Dave, Mostafa Dehghani, Sunipa Dev, Jacob Devlin, Mark Díaz, Nan Du, Ethan Dyer, Vlad Feinberg, Fangxiaoyu Feng, Vlad Fienber, Markus Freitag, Xavier Garcia, Sebastian Gehrmann, Lucas Gonzalez, Guy Gur-Ari, Steven Hand, Hadi Hashemi, Le Hou, Joshua Howland, Andrea Hu, Jeffrey Hui, Jeremy Hurwitz, Michael Isard, Abe Ittycheriah, Matthew Jagielski, Wenhao Jia, Kathleen Kenealy, Maxim Krikun, Sneha Kudugunta, Chang Lan, Katherine

Lee, Benjamin Lee, Eric Li, Music Li, Wei Li, Yaguang Li, Jian Li, Hyeontaek Lim, Hanzhao Lin, Zhongtao Liu, Frederick Liu, Marcello Maggioni, Aroma Mahendru, Joshua Maynez, Vedant Misra, Maysam Moussalem, Zachary Nado, John Nham, Eric Ni, Andrew Nystrom, Alicia Parrish, Marie Pellat, Martin Polacek, Alex Polozov, Reiner Pope, Siyuan Qiao, Emily Reif, Bryan Richter, Parker Riley, Alex Castro Ros, Aurko Roy, Brennan Saeta, Rajkumar Samuel, Renee Shelby, Ambrose Slone, Daniel Smilkov, David R So, Daniel Sohn, Simon Tokumine, Dasha Valter, Vijay Vasudevan, Kiran Vodrahalli, Xuezhi Wang, Pidong Wang, Zirui Wang, Tao Wang, John Wieting, Yuhuai Wu, Kelvin Xu, Yunhan Xu, Linting Xue, Pengcheng Yin, Jiahui Yu, Qiao Zhang, Steven Zheng, Ce Zheng, Weikang Zhou, Denny Zhou, Slav Petrov, and Yonghui Wu. PaLM 2 technical report. *arXiv [cs.CL]*, May 2023. URL `http://arxiv.org/abs/2305.10403`.

[446] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv [cs.LG]*, October 2021. URL `http://arxiv.org/abs/2110.14168`.

[447] Zhanming Jie, Jierui Li, and Wei Lu. Learning to reason deductively: Math word problem solving as complex relation extraction. *arXiv [cs.CL]*, March 2022. URL `http://arxiv.org/abs/2203.10316`.

[448] Yihuai Lan, Lei Wang, Qiyuan Zhang, Yunshi Lan, Bing Tian Dai, Yan Wang, Dongxiang Zhang, and Ee-Peng Lim. MWPToolkit: An open-source framework for deep learning-based math word problem solvers. *arXiv [cs.CL]*, September 2021. URL `http://arxiv.org/abs/2109.00799`.

[449] Maciej Besta, Nils Blach, Ales Kubicek, Robert Gerstenberger, Michal Podstawski, Lukas Gianinazzi, Joanna Gajda, Tomasz Lehmann, Hubert Niewiadomski, Piotr Nyczyk, and Torsten Hoefler. Graph of thoughts: Solving elaborate problems with large language models. *arXiv [cs.CL]*, August 2023. URL `http://arxiv.org/abs/2308.09687`.

[450] Bilgehan Sel, Ahmad Al-Tawaha, Vanshaj Khattar, Ruoxi Jia, and Ming Jin. Algorithm of thoughts: Enhancing exploration of ideas in large language models. *arXiv [cs.CL]*, August 2023. URL `http://arxiv.org/abs/2308.10379`.

[451] Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang, Yan Wang, Rui Wang, Yujiu Yang, Zhaopeng Tu, and Shuming Shi. Encouraging divergent thinking in large language models through multi-agent debate. *arXiv [cs.CL]*, May 2023. URL `http://arxiv.org/abs/2305.19118`.

[452] Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate. *arXiv [cs.CL]*, May 2023. URL `http://arxiv.org/abs/2305.14325`.

[453] Chi-Min Chan, Weize Chen, Yusheng Su, Jianxuan Yu, Wei Xue, Shan Zhang, Jie Fu, and Zhiyuan Liu. ChatEval: Towards better LLM-based evaluators through multi-agent debate. *ArXiv*, abs/2308.07201, August 2023. ISSN 2331-8422. doi:10.48550/arXiv.2308.07201. URL `http://dx.doi.org/10.48550/arXiv.2308.07201`.

[454] Chan Hee Song, Jiaman Wu, Clayton Washington, Brian M Sadler, Wei-Lun Chao, and Yu Su. LLM-planner: Few-shot grounded planning for embodied agents with large language models. *arXiv [cs.AI]*, December 2022. URL `http://arxiv.org/abs/2212.04088`.

[455] Bo Liu, Yuqian Jiang, Xiaohan Zhang, Qiang Liu, Shiqi Zhang, Joydeep Biswas, and Peter Stone. LLM+P: Empowering large language models with optimal planning proficiency. *arXiv [cs.AI]*, April 2023. URL `http://arxiv.org/abs/2304.11477`.

[456] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Shashank Gupta, Bodhisattwa Prasad Majumder, Katherine Hermann, Sean Welleck, Amir Yazdanbakhsh, and Peter Clark. Self-refine: Iterative refinement with self-feedback. *arXiv [cs.CL]*, March 2023. URL `http://arxiv.org/abs/2303.17651`.

[457] Zhiheng Xi, Senjie Jin, Yuhao Zhou, Rui Zheng, Songyang Gao, Tao Gui, Qi Zhang, and Xuanjing Huang. Self-polish: Enhance reasoning in large language models via problem refinement. *arXiv [cs.CL]*, May 2023. URL `http://arxiv.org/abs/2305.14497`.

[458] Zihao Wang, Shaofei Cai, Guanzhou Chen, Anji Liu, Xiaojian Ma, and Yitao Liang. Describe, explain, plan and select: Interactive planning with large language models enables open-world multi-task agents. *arXiv [cs.AI]*, February 2023. URL `http://arxiv.org/abs/2302.01560`.

[459] Ameet Deshpande, Vishvak Murahari, Tanmay Rajpurohit, Ashwin Kalyan, and Karthik Narasimhan. Toxicity in ChatGPT: Analyzing persona-assigned language models. *arXiv [cs.CL]*, April 2023. URL `http://arxiv.org/abs/2304.05335`.

[460] Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu Zheng, Yuheng Cheng, Ceyao Zhang, Jinlin Wang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, Liyang Zhou, Chenyu Ran, Lingfeng Xiao, Chenglin Wu, and Jurgen

Schmidhuber. MetaGPT: Meta programming for a multi-agent collaborative framework. *arXiv [cs.AI]*, August 2023. URL `http://arxiv.org/abs/2308.00352`.

[461] Guohao Li, Hasan Abed Al Kader Hammoud, Hani Itani, Dmitrii Khizbullin, and Bernard Ghanem. CAMEL: Communicative agents for "mind" exploration of large language model society. *arXiv [cs.AI]*, March 2023. URL `http://arxiv.org/abs/2303.17760`.

[462] Shi Jinxin, Zhao Jiabao, Wang Yilei, Wu Xingjiao, Li Jiawen, and He Liang. CGMI: Configurable general multi-agent interaction framework. *arXiv [cs.AI]*, August 2023. URL `http://arxiv.org/abs/2308.12503`.

[463] Chen Qian, Xin Cong, Wei Liu, Cheng Yang, Weize Chen, Yusheng Su, Yufan Dang, Jiahao Li, Juyuan Xu, Dahai Li, Zhiyuan Liu, and Maosong Sun. Communicative agents for software development. *arXiv [cs.SE]*, July 2023. URL `http://arxiv.org/abs/2307.07924`.

[464] Yunfan Shao, Linyang Li, Junqi Dai, and Xipeng Qiu. Character-LLM: A trainable agent for role-playing. *arXiv [cs.CL]*, October 2023. URL `http://arxiv.org/abs/2310.10158`.

[465] Pei Ke, Bosi Wen, Zhuoer Feng, Xiao Liu, Xuanyu Lei, Jiale Cheng, Shengyuan Wang, Aohan Zeng, Yuxiao Dong, Hongning Wang, Jie Tang, and Minlie Huang. CritiqueLLM: Scaling LLM-as-critic for effective and explainable evaluation of large language model generation. *arXiv [cs.CL]*, November 2023. URL `http://arxiv.org/abs/2311.18702`.

[466] Lei Wang, Jingsen Zhang, Hao Yang, Zhiyuan Chen, Jiakai Tang, Zeyu Zhang, Xu Chen, Yankai Lin, Ruihua Song, Wayne Xin Zhao, Jun Xu, Zhicheng Dou, Jun Wang, and Ji-Rong Wen. User behavior simulation with large language model based agents. *arXiv [cs.IR]*, June 2023. URL `http://arxiv.org/abs/2306.02552`.

[467] Lisa P Argyle, Ethan C Busby, Nancy Fulda, Joshua R Gubler, Christopher Rytting, and David Wingate. Out of one, many: Using language models to simulate human samples. *Polit. Anal.*, 31(3):337–351, July 2023. ISSN 1047-1987, 1476-4989. doi:10.1017/pan.2023.2. URL `https://www.cambridge.org/core/journals/political-analysis/article/out-of-one-many-using-language-models-to-simulate-human-samples/035D7C8A55B237942FB6DBAD7CAA4E49`.

[468] Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C Berg, Wan-Yen Lo, Piotr Dollár, and Ross Girshick. Segment anything. *arXiv [cs.CV]*, April 2023. URL `http://arxiv.org/abs/2304.02643`.

[469] AI Open. GPT-4V(ision) system card, 2023. URL `https://cdn.openai.com/papers/GPTV_System_Card.pdf`.

[470] Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. *arXiv [cs.CV]*, April 2023. URL `http://arxiv.org/abs/2304.08485`.

[471] Yang Zhao, Zhijie Lin, Daquan Zhou, Zilong Huang, Jiashi Feng, and Bingyi Kang. BuboGPT: Enabling visual grounding in multi-modal LLMs. *arXiv [cs.CV]*, July 2023. URL `http://arxiv.org/abs/2307.08581`.

[472] Chenyang Lyu, Minghao Wu, Longyue Wang, Xinting Huang, Bingshuai Liu, Zefeng Du, Shuming Shi, and Zhaopeng Tu. Macaw-LLM: Multi-modal language modeling with image, audio, video, and text integration. *arXiv [cs.CL]*, June 2023. URL `http://arxiv.org/abs/2306.09093`.

[473] Chenyu Wang, Weixin Luo, Qianyu Chen, Haonan Mai, Jindi Guo, Sixun Dong, Xiaohua, Xuan, Zhengxin Li, Lin Ma, and Shenghua Gao. Tool-LMM: A large multi-modal model for tool agent learning. *arXiv [cs.CV]*, January 2024. URL `http://arxiv.org/abs/2401.10727`.

[474] Difei Gao, Lei Ji, Luowei Zhou, Kevin Qinghong Lin, Joya Chen, Zihan Fan, and Mike Zheng Shou. AssistGPT: A general multi-modal assistant that can plan, execute, inspect, and learn. *arXiv [cs.CV]*, June 2023. URL `http://arxiv.org/abs/2306.08640`.

[475] Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv [cs.AI]*, May 2023. URL `http://arxiv.org/abs/2305.16291`.

[476] Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Chuyuan Fu, Keerthana Gopalakrishnan, Karol Hausman, Alex Herzog, Daniel Ho, Jasmine Hsu, Julian Ibarz, Brian Ichter, Alex Irpan, Eric Jang, Rosario Jauregui Ruano, Kyle Jeffrey, Sally Jesmonth, Nikhil J Joshi, Ryan Julian, Dmitry Kalashnikov, Yuheng Kuang, Kuang-Huei Lee, Sergey Levine, Yao Lu, Linda Luu, Carolina Parada, Peter Pastor, Jornell Quiambao, Kanishka Rao, Jarek Rettinghouse, Diego Reyes, Pierre Sermanet, Nicolas Sievers, Clayton Tan, Alexander Toshev, Vincent Vanhoucke, Fei Xia, Ted Xiao, Peng Xu, Sichun Xu, Mengyuan Yan, and Andy Zeng. Do as I can, not as I say: Grounding language in robotic affordances. *arXiv [cs.RO]*, April 2022. URL `http://arxiv.org/abs/2204.01691`.

[477] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating large language models trained on code. *arXiv [cs.LG]*, July 2021. URL `http://arxiv.org/abs/2107.03374`.

[478] Yujia Qin, Shengding Hu, Yankai Lin, Weize Chen, Ning Ding, Ganqu Cui, Zheni Zeng, Yufei Huang, Chaojun Xiao, Chi Han, Yi Ren Fung, Yusheng Su, Huadong Wang, Cheng Qian, Runchu Tian, Kunlun Zhu, Shihao Liang, Xingyu Shen, Bokai Xu, Zhen Zhang, Yining Ye, Bowen Li, Ziwei Tang, Jing Yi, Yuzhang Zhu, Zhenning Dai, Lan Yan, Xin Cong, Yaxi Lu, Weilin Zhao, Yuxiang Huang, Junxi Yan, Xu Han, Xian Sun, Dahai Li, Jason Phang, Cheng Yang, Tongshuang Wu, Heng Ji, Zhiyuan Liu, and Maosong Sun. Tool learning with foundation models. *arXiv [cs.CL]*, April 2023. URL `https://github.com/OpenBMB/BMTools`.

[479] Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, Xu Jiang, Karl Cobbe, Tyna Eloundou, Gretchen Krueger, Kevin Button, Matthew Knight, Benjamin Chess, and John Schulman. WebGPT: Browser-assisted question-answering with human feedback. *arXiv [cs.CL]*, December 2021. URL `http://arxiv.org/abs/2112.09332`.

[480] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. *arXiv [cs.CL]*, February 2023. URL `http://arxiv.org/abs/2302.04761`.

[481] Aaron Parisi, Yao Zhao, and Noah Fiedel. TALM: Tool augmented language models. *arXiv [cs.CL]*, May 2022. URL `http://arxiv.org/abs/2205.12255`.

[482] Cheng Qian, Chenyan Xiong, Zhenghao Liu, and Zhiyuan Liu. Toolink: Linking toolkit creation and using through chain-of-solving on open-source model. *arXiv [cs.CL]*, October 2023. URL `http://arxiv.org/abs/2310.05155`.

[483] Ehud Karpas, Omri Abend, Yonatan Belinkov, Barak Lenz, Opher Lieber, Nir Ratner, Yoav Shoham, Hofit Bata, Yoav Levine, Kevin Leyton-Brown, Dor Muhlgay, Noam Rozen, Erez Schwartz, Gal Shachaf, Shai Shalev-Shwartz, Amnon Shashua, and Moshe Tenenholtz. MRKL systems: A modular, neuro-symbolic architecture that combines large language models, external knowledge sources and discrete reasoning. *arXiv [cs.CL]*, May 2022. URL `http://arxiv.org/abs/2205.00445`.

[484] Zhipeng Chen, Kun Zhou, Beichen Zhang, Zheng Gong, Wayne Xin Zhao, and Ji-Rong Wen. ChatCoT: Tool-augmented chain-of-thought reasoning on chat-based large language models. *arXiv [cs.CL]*, May 2023. URL `http://arxiv.org/abs/2305.14323`.

[485] Tianle Cai, Xuezhi Wang, Tengyu Ma, Xinyun Chen, and Denny Zhou. Large language models as tool makers. *arXiv [cs.LG]*, May 2023. URL `http://arxiv.org/abs/2305.17126`.

[486] Cheng Qian, Chi Han, Yi R Fung, Yujia Qin, Zhiyuan Liu, and Heng Ji. CREATOR: Disentangling abstract and concrete reasonings of large language models through tool creation. *arXiv [cs.CL]*, May 2023. URL `http://arxiv.org/abs/2305.14318`.

[487] Lifan Yuan, Yangyi Chen, Xingyao Wang, Yi R Fung, Hao Peng, and Heng Ji. CRAFT: Customizing LLMs by creating and retrieving from specialized toolsets. *arXiv [cs.CL]*, September 2023. URL `http://arxiv.org/abs/2309.17428`.

[488] Hilmy Abiyyu A. Flaticon. `https://www.flaticon.com/authors/hilmy-abiyyu-a`. Accessed: 2024-5-1.

[489] Laisa Islam Ani. Flaticon. `https://www.flaticon.com/authors/laisa-islam-ani`. Accessed: 2024-5-1.

[490] Freepik. Flaticon. `https://www.flaticon.com/authors/freepik`. Accessed: 2024-5-1.

[491] Kiranshastry. Flaticon. `https://www.flaticon.com/authors/kiranshastry`. Accessed: 2024-5-1.

[492] Michael D Skarlinski, Sam Cox, Jon M Laurent, James D Braza, Michaela Hinks, Michael J Hammerling, Manvitha Ponnapati, Samuel G Rodriques, and Andrew D White. Language agents achieve superhuman synthesis of scientific knowledge. *arXiv [cs.CL]*, September 2024. URL `http://arxiv.org/abs/2409.13740`.

[493] Yuan Chiang, Elvis Hsieh, Chia-Hong Chou, and Janosh Riebesell. LLaMP: Large language model made powerful for high-fidelity materials knowledge retrieval and distillation. *arXiv [cs.CL]*, January 2024. URL `http://arxiv.org/abs/2401.17244`.

[494] Pingchuan Ma, Tsun-Hsuan Wang, Minghao Guo, Zhiqing Sun, Joshua B Tenenbaum, Daniela Rus, Chuang Gan, and Wojciech Matusik. LLM and simulation as bilevel optimizers: A new paradigm to advance physical scientific discovery. *arXiv [cs.LG]*, May 2024. URL http://arxiv.org/abs/2405.09783.

[495] Yuanhao Qu, Kaixuan Huang, Henry Cousins, William A Johnson, Di Yin, Mihir Shah, Denny Zhou, Russ Altman, Mengdi Wang, and Le Cong. CRISPR-GPT: An LLM agent for automated design of gene-editing experiments. *bioRxiv*, April 2024. doi:10.1101/2024.04.25.591003. URL http://dx.doi.org/10.48550/arXiv.2404.18021.

[496] Haoyang Liu, Yijiang Li, Jinglin Jian, Yuxuan Cheng, Jianrong Lu, Shuyi Guo, Jinglei Zhu, Mianchen Zhang, Miantong Zhang, and Haohan Wang. Toward a team of AI-made scientists for scientific discovery from gene expression data. *arXiv [q-bio.GN]*, February 2024. URL http://arxiv.org/abs/2402.12391.

[497] Henry W Sprueill, Carl Edwards, Khushbu Agarwal, Mariefel V Olarte, Udishnu Sanyal, Conrad Johnston, Hongbin Liu, Heng Ji, and Sutanay Choudhury. ChemReasoner: Heuristic search over a large language model's knowledge space using quantum-chemical feedback. *arXiv [physics.chem-ph]*, February 2024. URL http://arxiv.org/abs/2402.10980.

[498] Yubo Ma, Zhibin Gou, Junheng Hao, Ruochen Xu, Shuohang Wang, Liangming Pan, Yujiu Yang, Yixin Cao, and Aixin Sun. SciAgent: Tool-augmented language models for scientific reasoning. *arXiv [cs.CL]*, February 2024. URL http://arxiv.org/abs/2402.11451.

[499] Yijia Shao, Yucheng Jiang, Theodore A Kanell, Peter Xu, Omar Khattab, and Monica S Lam. Assisting in writing wikipedia-like articles from scratch with large language models. *arXiv [cs.CL]*, February 2024. URL http://arxiv.org/abs/2402.14207.

[500] Christoph Völker, Tehseen Rug, Kevin Maik Jablonka, and Sabine Kruschwitz. LLMs can design sustainable concrete – a systematic benchmark. February 2024. URL https://www.researchsquare.com/article/rs-3913272/v1.

[501] A Ghafarollahi and M J Buehler. ProtAgents: Protein discovery via large language model multi-agent collaborations combining physics and machine learning. *arXiv [cond-mat.soft]*, January 2024. URL http://arxiv.org/abs/2402.04268.

[502] Jakub Lála, Odhran O'Donoghue, Aleksandar Shtedritski, Sam Cox, Samuel G Rodriques, and Andrew D White. PaperQA: Retrieval-augmented generative agent for scientific research. *arXiv [cs.CL]*, December 2023. URL http://arxiv.org/abs/2312.07559.

[503] Sam Cox, Michael Hammerling, Jakub Lála, Jon Laurent, Sam Rodriques, Matt Rubashkin, Andrew White. WikiCrow: Automating synthesis of human scientific knowledge. https://www.futurehouse.org/wikicrow, 2023. URL https://www.futurehouse.org/wikicrow. Accessed: 2024-2-15.

[504] Mehrad Ansari and Seyed Mohamad Moosavi. Agent-based learning of materials datasets from scientific literature. *arXiv [cs.AI]*, December 2023. URL http://arxiv.org/abs/2312.11690.

[505] Michael H Prince, Henry Chan, Aikaterini Vriza, Tao Zhou, Varuni K Sastry, Matthew T Dearing, Ross J Harder, Rama K Vasudevan, and Mathew J Cherukara. Opportunities for retrieval and tool augmented large language models in scientific facilities. *arXiv [cs.CE]*, December 2023. URL http://arxiv.org/abs/2312.01291.

[506] Yiren Liu, Si Chen, Haocong Cheng, Mengxia Yu, Xiao Ran, Andrew Mo, Yiliu Tang, and Yun Huang. CoQuest: Exploring research question co-creation with an LLM-based agent. *arXiv [cs.HC]*, October 2023. URL http://arxiv.org/abs/2310.06155.

[507] Odhran O'Donoghue, Aleksandar Shtedritski, John Ginger, Ralph Abboud, Ali Essa Ghareeb, Justin Booth, and Samuel G Rodriques. BioPlanner: Automatic evaluation of LLMs on protocol planning in biology. *arXiv [cs.CL]*, October 2023. URL http://arxiv.org/abs/2310.10632.

[508] Nikita Janakarajan, Tim Erdmann, Sarath Swaminathan, Teodoro Laino, and Jannis Born. Language models in molecular discovery. *arXiv [physics.chem-ph]*, September 2023. URL http://arxiv.org/abs/2309.16235.

[509] Yeonghun Kang and Jihan Kim. ChatMOF: An autonomous AI system for predicting and generating metal-organic frameworks. *arXiv [cs.CL]*, July 2023. URL http://arxiv.org/abs/2308.01423.

[510] Beatriz Mouriño, Elias Moubarak, Joren Van Herck, Sauradeep Majumdar, and Xiaoqi Zhang. i-digest: v1.0, 2023. URL https://zenodo.org/record/8080962.

[511] Bojana Rankovic, Andres M Bran, and Philippe Schwaller. BOLLaMa: BOLLaMA interface working with CHAOS, 2023. URL https://zenodo.org/record/8096827.

[512] Sabine Kruschwitz, Christoph Völker, and Ghezal Ahmad Zia. Text2Concrete, 2023. URL https://zenodo.org/record/8091195.

[513] Mayk Caldas Ramos, Sam Cox, and Andrew White. MAPI_LLM: MAPI_LLM first release, 2023. URL https://zenodo.org/record/8097336.

[514] Mayk Caldas Ramos, Shane S Michtavy, Marc D Porosoff, and Andrew D White. Bayesian optimization of catalysts with in-context learning. *arXiv [physics.chem-ph]*, April 2023. URL http://arxiv.org/abs/2304.05341.

[515] Glen M Hocky and Andrew D White. Natural language processing models that automate programming will transform chemistry research and teaching. *Digit. Discov.*, 1(2):79–83, April 2022. ISSN 2635-098X. doi:10.1039/d1dd00009h. URL https://pubs.rsc.org/en/content/articlelanding/2022/dd/d1dd00009h.

[516] Andrew D White, Glen M Hocky, Heta A Gandhi, Mehrad Ansari, Sam Cox, Geemi P Wellawatte, Subarna Sasmal, Ziyue Yang, Kangxin Liu, Yuvraj Singh, and Willmor J Peña Ccoa. Assessment of chemistry knowledge in large language models that generate code. *Digital Discovery*, 2(2):368–376, April 2023. ISSN 2635-098X. doi:10.1039/D2DD00087C. URL https://pubs.rsc.org/en/content/articlelanding/2023/dd/d2dd00087c.

[517] Agustinus Kristiadi, Felix Strieth-Kalthoff, Marta Skreta, Pascal Poupart, Alán Aspuru-Guzik, and Geoff Pleiss. A sober look at LLMs for material discovery: Are they actually good for bayesian optimization over molecules? *arXiv [cs.LG]*, February 2024. URL http://arxiv.org/abs/2402.05015.

[518] Bojana Ranković and Philippe Schwaller. BoChemian: Large language model embeddings for bayesian optimization of chemical reactions. December 2023. URL https://openreview.net/pdf?id=A1RVn1m3J3.

[519] K Jablonka, Qianxiang Ai, Alexander H Al-Feghali, S Badhwar, Joshua D Bocarsly Andres Bran, S Bringuier, L Brinson, K Choudhary, Defne Circi, Sam Cox, W D Jong, Matthew L Evans, Nicolas Gastellu, Jerome Genzling, M Gil, Ankur Gupta, Zhi Hong, A Imran, S Kruschwitz, A Labarre, Jakub L'ala, Tao Liu, Steven Ma, Sauradeep Majumdar, G Merz, N Moitessier, E Moubarak, B Mouriño, Brenden G Pelkie, M Pieler, M C Ramos, Bojana Rankovi'c, Samuel G Rodriques, J N Sanders, P Schwaller, Marcus Schwarting, Jia-Xin Shi, B Smit, Benn Smith, J V Heck, C Volker, Logan T Ward, S Warren, B Weiser, Sylvester Zhang, Xiaoqi Zhang, Ghezal Ahmad Jan Zia, A Scourtas, K Schmidt, Ian T Foster, Andrew D White, and B Blaiszik. 14 examples of how LLMs can transform materials science and chemistry: a reflection on a large language model hackathon. *Digit. Discov.*, 2(5):1233–1250, June 2023. ISSN 2635-098X. doi:10.1039/D3DD00113J. URL https://pubs.rsc.org/en/content/articlelanding/2023/dd/d3dd00113j.

[520] Bojana Ranković, Ryan-Rhys Griffiths, Henry B Moss, and Philippe Schwaller. Bayesian optimisation for additive screening and yield improvements in chemical reactions – beyond one-hot encoding. *ChemRxiv*, June 2023. doi:10.26434/chemrxiv-2022-nll2j-v3. URL https://chemrxiv.org/engage/chemrxiv/article-details/6489f95c4f8b1884b74b69c8.

[521] Benjamin Weiser, Jerome Genzling, Nicolas Gastellu, Sylvester Zhang, Tao Liu, Alexander Al-Feghali, Nicolas Moitessier, Anne Labarre, and Steven Ma. LLM-Guided-GA: LLM-Guided-GA digital discovery release, 2023. URL https://zenodo.org/record/8125541.

[522] Defne Circi and Shruti Badhwar. InsightGraph: InsightGraph, 2023. URL https://zenodo.org/record/8092575.

[523] Marwa Zaabi, Walid Hariri, and Nadia Smaoui. A review study of ChatGPT applications in education. In *2023 International Conference on Innovations in Intelligent Systems and Applications (INISTA)*, pages 1–5. IEEE, September 2023. doi:10.1109/inista59065.2023.10310439. URL http://dx.doi.org/10.1109/INISTA59065.2023.10310439.

[524] Enkelejda Kasneci, Kathrin Sessler, Stefan Küchemann, Maria Bannert, Daryna Dementieva, Frank Fischer, Urs Gasser, Georg Groh, Stephan Günnemann, Eyke Hüllermeier, Stephan Krusche, Gitta Kutyniok, Tilman Michaeli, Claudia Nerdel, Jürgen Pfeffer, Oleksandra Poquet, Michael Sailer, Albrecht Schmidt, Tina Seidel, Matthias Stadler, Jochen Weller, Jochen Kuhn, and Gjergji Kasneci. ChatGPT for good? on opportunities and challenges of large language models for education. *Learn. Individ. Differ.*, 103(102274):102274, April 2023. ISSN 1041-6080, 1873-3425. doi:10.1016/j.lindif.2023.102274. URL https://www.sciencedirect.com/science/article/pii/S1041608023000195.

[525] Arto Hellas, Juho Leinonen, Sami Sarsa, Charles Koutcheme, Lilja Kujanpaa, and Juha Sorva. Exploring the responses of large language models to beginner programmers' help requests. *arXiv [cs.CY]*, June 2023. URL http://arxiv.org/abs/2306.05715.

[526] Yuhao Dan, Zhikai Lei, Yiyang Gu, Yong Li, Jianghao Yin, Jiaju Lin, Linhao Ye, Zhiyan Tie, Yougen Zhou, Yilei Wang, Aimin Zhou, Ze Zhou, Qin Chen, Jie Zhou, Liang He, and Xipeng Qiu. EduChat: A large-

scale language model-based chatbot system for intelligent education. *arXiv [cs.CL]*, August 2023. URL http://arxiv.org/abs/2308.02773.

[527] Jie Tian, Jixin Hou, Zihao Wu, Peng Shu, Zhengliang Liu, Yujie Xiang, Beikang Gu, Nicholas Filla, Yiwei Li, Ning Liu, Xianyan Chen, Keke Tang, Tianming Liu, and Xianqiao Wang. Assessing large language models in mechanical engineering education: A study on mechanics-focused conceptual understanding. *arXiv [cs.CL]*, January 2024. URL http://arxiv.org/abs/2401.12983.

[528] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Kuttler, Mike Lewis, Wen-Tau Yih, Tim Rocktaschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. *arXiv [cs.CL]*, May 2020. URL http://arxiv.org/abs/2005.11401.

[529] Kexin Chen, Junyou Li, Kunyi Wang, Yuyang Du, Jiahui Yu, Jiamin Lu, Lanqing Li, Jiezhong Qiu, Jianzhang Pan, Yi Huang, Qun Fang, Pheng Ann Heng, and Guangyong Chen. Chemist-X: Large language model-empowered agent for reaction condition recommendation in chemical synthesis. *arXiv [cs.IR]*, November 2023. URL http://arxiv.org/abs/2311.10776.

[530] John B Ingraham, Max Baranov, Zak Costello, Karl W Barber, Wujie Wang, Ahmed Ismail, Vincent Frappier, Dana M Lord, Christopher Ng-Thow-Hing, Erik R Van Vlack, Shan Tie, Vincent Xue, Sarah C Cowles, Alan Leung, João V Rodrigues, Claudio L Morales-Perez, Alex M Ayoub, Robin Green, Katherine Puentes, Frank Oplinger, Nishant V Panwar, Fritz Obermeyer, Adam R Root, Andrew L Beam, Frank J Poelwijk, and Gevorg Grigoryan. Illuminating protein space with a programmable generative model. *Nature*, 623 (7989):1070–1078, November 2023. ISSN 0028-0836,1476-4687. doi:10.1038/s41586-023-06728-8. URL http://dx.doi.org/10.1038/s41586-023-06728-8.

[531] Kevin E Wu, Kevin K Yang, Rianne van den Berg, Sarah Alamdari, James Y Zou, Alex X Lu, and Ava P Amini. Protein structure generation via folding diffusion. *Nat. Commun.*, 15(1):1059, February 2024. ISSN 2041-1723. doi:10.1038/s41467-024-45051-2. URL http://dx.doi.org/10.1038/s41467-024-45051-2.

[532] Xiangru Tang, Qiao Jin, Kunlun Zhu, Tongxin Yuan, Yichi Zhang, Wangchunshu Zhou, Meng Qu, Yilun Zhao, Jian Tang, Zhuosheng Zhang, Arman Cohan, Zhiyong Lu, and Mark Gerstein. Prioritizing safeguarding over autonomy: Risks of LLM agents for science. *arXiv [cs.CY]*, February 2024. URL http://arxiv.org/abs/2402.04247.

[533] Yixiang Ruan, Chenyin Lu, Ning Xu, Jian Zhang, Jun Xuan, Jianzhang Pan, Qun Fang, Hanyu Gao, Xiaodong Shen, Ning Ye, Qiang Zhang, and Yiming Mo. Accelerated end-to-end chemical synthesis development with large language models. *ChemRxiv*, May 2024. doi:10.26434/chemrxiv-2024-6wmg4. URL https://chemrxiv.org/engage/chemrxiv/article-details/6634f02021291e5d1d58702c.

[534] Zhiling Zheng, Oufan Zhang, Christian Borgs, Jennifer T Chayes, and Omar M Yaghi. ChatGPT Chemistry Assistant for Text Mining and the Prediction of MOF Synthesis. *J. Am. Chem. Soc.*, 145(Copyright © 2024 American Chemical Society (ACS). All Rights Reserved.; Copyright © 2024 U.S. National Library of Medicine.): 18048–18062, 2023. ISSN 0002-7863,1520-5126. doi:10.1021/jacs.3c05819. URL http://dx.doi.org/10.1021/jacs.3c05819.

[535] Zhiling Zheng, Zhiguo He, Omar Khattab, Nakul Rampal, Matei A Zaharia, Christian Borgs, Jennifer T Chayes, and Omar M Yaghi. Image and data mining in reticular chemistry powered by GPT-4V. *Digit. Discov.*, 3(3): 491–501, 2024. ISSN 2635-098X. doi:10.1039/d3dd00239j. URL https://pubs.rsc.org/en/content/articlelanding/2024/dd/d3dd00239j.

[536] Zhiling Zheng, Oufan Zhang, Ha L Nguyen, Nakul Rampal, Ali H Alawadhi, Zichao Rong, Teresa Head-Gordon, Christian Borgs, Jennifer T Chayes, and Omar M Yaghi. ChatGPT research group for optimizing the crystallinity of MOFs and COFs. *ACS Cent. Sci.*, 9(11):2161–2170, November 2023. ISSN 2374-7943,2374-7951. doi:10.1021/acscentsci.3c01087. URL https://pubs.acs.org/doi/full/10.1021/acscentsci.3c01087.

[537] Nathalia Nascimento, Paulo Alencar, and Donald Cowan. Self-adaptive large language model (LLM)-based multiagent systems. In *2023 IEEE International Conference on Autonomic Computing and Self-Organizing Systems Companion (ACSOS-C)*, pages 104–109. IEEE, September 2023. ISBN 9798350337464. doi:10.1109/ACSOS-C58168.2023.00048. URL https://ieeexplore.ieee.org/document/10336211.

[538] Sarfaraz K Niazi and Zamara Mariam. Recent advances in Machine-Learning-Based chemoinformatics: A comprehensive review. *Int. J. Mol. Sci.*, 24(14), July 2023. ISSN 1422-0067. doi:10.3390/ijms241411488. URL http://dx.doi.org/10.3390/ijms241411488.

[539] Andrew D McNaughton, Gautham Ramalaxmi, Agustin Kruel, Carter R Knutson, Rohith A Varikoti, and Neeraj Kumar. CACTUS: Chemistry agent connecting tool-usage to science. *arXiv [cs.CL]*, May 2024. URL http://arxiv.org/abs/2405.00972.

[540] Yeonghun Kang, Hyunsoo Park, Berend Smit, and Jihan Kim. A multi-modal pre-training transformer for universal transfer learning in metal–organic frameworks. *Nature Machine Intelligence*, 5(3):309–318, March 2023. ISSN 2522-5839, 2522-5839. doi:10.1038/s42256-023-00628-2. URL https://www.nature.com/articles/s42256-023-00628-2.

[541] Matteo Manica, Jannis Born, Joris Cadow, Dimitrios Christofidellis, Ashish Dave, Dean Clarke, Yves Gaetan Nana Teukam, Giorgio Giannone, Samuel C Hoffman, Matthew Buchan, Vijil Chenthamarakshan, Timothy Donovan, Hsiang Han Hsu, Federico Zipoli, Oliver Schilter, Akihiro Kishimoto, Lisa Hamada, Inkit Padhi, Karl Wehden, Lauren McHugh, Alexy Khrabrov, Payel Das, Seiji Takeda, and John R Smith. Accelerating material design with the generative toolkit for scientific discovery. *npj Computational Materials*, 9(1):1–6, May 2023. ISSN 2057-3960, 2057-3960. doi:10.1038/s41524-023-01028-1. URL https://www.nature.com/articles/s41524-023-01028-1.

[542] rxn4chemistry: Python wrapper for the IBM RXN for chemistry API. URL https://github.com/rxn4chemistry/rxn4chemistry.

[543] Piotr Gaiński, Lukasz Maziarka, Tomasz Danel, and Stanislaw Jastrzebski. HuggingMolecules: An Open-Source library for Transformer-Based molecular property prediction (student abstract). *AAAI*, 36(11):12949–12950, June 2022. ISSN 2374-3468, 2374-3468. doi:10.1609/aaai.v36i11.21611. URL https://ojs.aaai.org/index.php/AAAI/article/view/21611.

[544] Greg Landrum. Rdkit documentation. *Release*, 1(1-79):4, 2013. ISSN 1047-935X, 1533-3752. URL https://media.readthedocs.org/pdf/rdkit/latest/rdkit.pdf.

[545] Yizhen Zheng, Huan Yee Koh, Jiaxin Ju, Anh T N Nguyen, Lauren T May, Geoffrey I Webb, and Shirui Pan. Large language models for scientific synthesis, inference and explanation. *arXiv [cs.AI]*, October 2023. URL http://arxiv.org/abs/2310.07984.

[546] Qingyun Wang, Doug Downey, Heng Ji, and Tom Hope. SciMON: Scientific inspiration machines optimized for novelty. *arXiv [cs.CL]*, May 2023. URL http://arxiv.org/abs/2305.14259.

[547] Xuemei Gu and Mario Krenn. Generation and human-expert evaluation of interesting research ideas using knowledge graphs and large language models. *arXiv [cs.AI]*, May 2024. URL http://arxiv.org/abs/2405.17044.

[548] Henry W Sprueill, Carl Edwards, Mariefel V Olarte, Udishnu Sanyal, Heng Ji, and Sutanay Choudhury. Monte carlo thought search: Large language model querying for complex scientific reasoning in catalyst design. *arXiv [cs.AI]*, October 2023. URL http://arxiv.org/abs/2310.14420.

[549] Gengmo Zhou, Zhifeng Gao, Qiankun Ding, Hang Zheng, Hongteng Xu, Zhewei Wei, Linfeng Zhang, and Guolin Ke. Uni-Mol: A universal 3D molecular representation learning framework. *ChemRxiv*, March 2023. doi:10.26434/chemrxiv-2022-jjm0j-v4. URL https://chemrxiv.org/engage/chemrxiv/article-details/6402990d37e01856dc1d1581.

[550] Santiago Miret and N M Anoop Krishnan. Are LLMs ready for real-world materials discovery? *arXiv [cond-mat.mtrl-sci]*, February 2024. URL http://arxiv.org/abs/2402.05200.

[551] Yuanqi Du, Chenru Duan, Andres Bran, Anna Sotnikova, Yi Qu, Heather Kulik, Antoine Bosselut, Jinjia Xu, and Philippe Schwaller. Large language models are catalyzing chemistry education. *ChemRxiv*, June 2024. doi:10.26434/chemrxiv-2024-h722v. URL https://chemrxiv.org/engage/api-gateway/chemrxiv/assets/orp/resource/item/66772be25101a2ffa8412ee0/original/large-language-models-are-catalyzing-chemistry-education.pdf.

[552] Jinlu Zhang, Yin Fang, Xin Shao, Huajun Chen, Ningyu Zhang, and Xiaohui Fan. The future of molecular studies through the lens of large language models. *J. Chem. Inf. Model.*, 64(3):563–566, February 2024. ISSN 1549-9596, 1549-960X. doi:10.1021/acs.jcim.3c01977. URL http://dx.doi.org/10.1021/acs.jcim.3c01977.

[553] OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, Red Avila, Igor Babuschkin, Suchir Balaji, Valerie Balcom, Paul Baltescu, Haiming Bao, Mohammad Bavarian, Jeff Belgum, Irwan Bello, Jake Berdine, Gabriel Bernadett-Shapiro, Christopher Berner, Lenny Bogdonoff, Oleg Boiko, Madelaine Boyd, Anna-Luisa Brakman, Greg Brockman, Tim Brooks, Miles Brundage, Kevin Button, Trevor Cai, Rosie Campbell, Andrew Cann, Brittany Carey, Chelsea Carlson, Rory Carmichael, Brooke Chan, Che Chang, Fotis Chantzis, Derek Chen, Sully Chen, Ruby Chen, Jason Chen, Mark Chen, Ben Chess, Chester Cho, Casey Chu, Hyung Won Chung,

Dave Cummings, Jeremiah Currier, Yunxing Dai, Cory Decareaux, Thomas Degry, Noah Deutsch, Damien Deville, Arka Dhar, David Dohan, Steve Dowling, Sheila Dunning, Adrien Ecoffet, Atty Eleti, Tyna Eloundou, David Farhi, Liam Fedus, Niko Felix, Simón Posada Fishman, Juston Forte, Isabella Fulford, Leo Gao, Elie Georges, Christian Gibson, Vik Goel, Tarun Gogineni, Gabriel Goh, Rapha Gontijo-Lopes, Jonathan Gordon, Morgan Grafstein, Scott Gray, Ryan Greene, Joshua Gross, Shixiang Shane Gu, Yufei Guo, Chris Hallacy, Jesse Han, Jeff Harris, Yuchen He, Mike Heaton, Johannes Heidecke, Chris Hesse, Alan Hickey, Wade Hickey, Peter Hoeschele, Brandon Houghton, Kenny Hsu, Shengli Hu, Xin Hu, Joost Huizinga, Shantanu Jain, Shawn Jain, Joanne Jang, Angela Jiang, Roger Jiang, Haozhun Jin, Denny Jin, Shino Jomoto, Billie Jonn, Heewoo Jun, Tomer Kaftan, Lukasz Kaiser, Ali Kamali, Ingmar Kanitscheider, Nitish Shirish Keskar, Tabarak Khan, Logan Kilpatrick, Jong Wook Kim, Christina Kim, Yongjik Kim, Jan Hendrik Kirchner, Jamie Kiros, Matt Knight, Daniel Kokotajlo, Lukasz Kondraciuk, Andrew Kondrich, Aris Konstantinidis, Kyle Kosic, Gretchen Krueger, Vishal Kuo, Michael Lampe, Ikai Lan, Teddy Lee, Jan Leike, Jade Leung, Daniel Levy, Chak Ming Li, Rachel Lim, Molly Lin, Stephanie Lin, Mateusz Litwin, Theresa Lopez, Ryan Lowe, Patricia Lue, Anna Makanju, Kim Malfacini, Sam Manning, Todor Markov, Yaniv Markovski, Bianca Martin, Katie Mayer, Andrew Mayne, Bob McGrew, Scott Mayer McKinney, Christine McLeavey, Paul McMillan, Jake McNeil, David Medina, Aalok Mehta, Jacob Menick, Luke Metz, Andrey Mishchenko, Pamela Mishkin, Vinnie Monaco, Evan Morikawa, Daniel Mossing, Tong Mu, Mira Murati, Oleg Murk, David Mély, Ashvin Nair, Reiichiro Nakano, Rajeev Nayak, Arvind Neelakantan, Richard Ngo, Hyeonwoo Noh, Long Ouyang, Cullen O'Keefe, Jakub Pachocki, Alex Paino, Joe Palermo, Ashley Pantuliano, Giambattista Parascandolo, Joel Parish, Emy Parparita, Alex Passos, Mikhail Pavlov, Andrew Peng, Adam Perelman, Filipe de Avila Belbute Peres, Michael Petrov, Henrique Ponde de Oliveira Pinto, Michael, Pokorny, Michelle Pokrass, Vitchyr H Pong, Tolly Powell, Alethea Power, Boris Power, Elizabeth Proehl, Raul Puri, Alec Radford, Jack Rae, Aditya Ramesh, Cameron Raymond, Francis Real, Kendra Rimbach, Carl Ross, Bob Rotsted, Henri Roussez, Nick Ryder, Mario Saltarelli, Ted Sanders, Shibani Santurkar, Girish Sastry, Heather Schmidt, David Schnurr, John Schulman, Daniel Selsam, Kyla Sheppard, Toki Sherbakov, Jessica Shieh, Sarah Shoker, Pranav Shyam, Szymon Sidor, Eric Sigler, Maddie Simens, Jordan Sitkin, Katarina Slama, Ian Sohl, Benjamin Sokolowsky, Yang Song, Natalie Staudacher, Felipe Petroski Such, Natalie Summers, Ilya Sutskever, Jie Tang, Nikolas Tezak, Madeleine B Thompson, Phil Tillet, Amin Tootoonchian, Elizabeth Tseng, Preston Tuggle, Nick Turley, Jerry Tworek, Juan Felipe Cerón Uribe, Andrea Vallone, Arun Vijayvergiya, Chelsea Voss, Carroll Wainwright, Justin Jay Wang, Alvin Wang, Ben Wang, Jonathan Ward, Jason Wei, C J Weinmann, Akila Welihinda, Peter Welinder, Jiayi Weng, Lilian Weng, Matt Wiethoff, Dave Willner, Clemens Winter, Samuel Wolrich, Hannah Wong, Lauren Workman, Sherwin Wu, Jeff Wu, Michael Wu, Kai Xiao, Tao Xu, Sarah Yoo, Kevin Yu, Qiming Yuan, Wojciech Zaremba, Rowan Zellers, Chong Zhang, Marvin Zhang, Shengjia Zhao, Tianhao Zheng, Juntang Zhuang, William Zhuk, and Barret Zoph. GPT-4 technical report. March 2023. URL http://arxiv.org/abs/2303.08774.

[554] Wenhao Gao and Connor W Coley. The synthesizability of molecules proposed by generative models. *J. Chem. Inf. Model.*, (12):5714–5723, February 2020. ISSN 1549-9596,1549-960X. doi:10.1021/acs.jcim.0c00174. URL https://pubs.acs.org/doi/10.1021/acs.jcim.0c00174.

[555] Pengfei Liu, Jun Tao, and Zhixiang Ren. Scientific language modeling: A quantitative review of large language models in molecular science. *arXiv [cs.LG]*, February 2024. URL http://arxiv.org/abs/2402.04119.

[556] Philippe Schwaller, Benjamin Hoover, Jean-Louis Reymond, Hendrik Strobelt, and Teodoro Laino. Extraction of organic chemistry grammar from unsupervised learning of chemical reactions. *Sci. Adv.*, 7(15):eabe4166, April 2021. ISSN 2375-2548. doi:10.1126/sciadv.abe4166. URL http://dx.doi.org/10.1126/sciadv.abe4166.

[557] Oliver Schilter, Marvin Alberts, Federico Zipoli, Alain C Vaucher, Philippe Schwaller, and Teodoro Laino. Unveiling the secrets of ^{1}H-NMR spectroscopy: A novel approach utilizing attention mechanisms. November 2023. URL https://openreview.net/pdf?id=4ilKwquW51.

[558] Shiva Aryal, Tuyen Do, Bisesh Heyojoo, Sandeep Chataut, Bichar Dip Shrestha Gurung, Venkataramana Gadhamshetty, and Etienne Gnimpieba. Leveraging multi-AI agents for cross-domain knowledge discovery. *arXiv [cs.AI]*, April 2024. URL http://arxiv.org/abs/2404.08511.

[559] Florian Bohm, Yang Gao, Christian M Meyer, Ori Shapira, Ido Dagan, and Iryna Gurevych. Better rewards yield better summaries: Learning to summarise without references. *arXiv [cs.CL]*, September 2019. URL http://arxiv.org/abs/1909.01214.

[560] Sarah Pan, Vladislav Lialin, Sherin Muckatira, and Anna Rumshisky. Let's reinforce step by step. *arXiv [cs.CL]*, November 2023. URL http://arxiv.org/abs/2311.05821.

[561] Bin Hu, Chenyang Zhao, Pu Zhang, Zihao Zhou, Yuanhang Yang, Zenglin Xu, and Bin Liu. Enabling intelligent interactions between an agent and an LLM: A reinforcement learning approach. *arXiv [cs.AI]*, June 2023. URL http://arxiv.org/abs/2306.03604.

[562] Zelai Xu, Chao Yu, Fei Fang, Yu Wang, and Yi Wu. Language agents with reinforcement learning for strategic play in the werewolf game. *arXiv [cs.AI]*, October 2023. URL `http://arxiv.org/abs/2310.18940`.

[563] Meredith Ringel Morris, Jascha Sohl-dickstein, Noah Fiedel, Tris Warkentin, Allan Dafoe, Aleksandra Faust, Clement Farabet, and Shane Legg. Levels of AGI for operationalizing progress on the path to AGI. *arXiv [cs.AI]*, November 2023. URL `http://arxiv.org/abs/2311.02462`.

[564] Types of artificial intelligence. `https://www.ibm.com/think/topics/artificial-intelligence-types`, August 2024. Accessed: 2024-10-9.

[565] Pat Langley, John E Laird, and Seth Rogers. Cognitive architectures: Research issues and challenges. *Cogn. Syst. Res.*, 10(2):141–160, June 2009. ISSN 1389-0417,2214-4366. doi:10.1016/j.cogsys.2006.07.004. URL `http://dx.doi.org/10.1016/j.cogsys.2006.07.004`.

[566] Ben Goertzel. Artificial general intelligence: Concept, state of the art, and future prospects. *J. Artif. Gen. Intell.*, 5(1):1–48, December 2014. ISSN 1946-0163. doi:10.2478/jagi-2014-0001. URL `https://sciendo.com/pdf/10.2478/jagi-2014-0001`.

[567] Yangjun Ruan, Honghua Dong, Andrew Wang, Silviu Pitis, Yongchao Zhou, Jimmy Ba, Yann Dubois, Chris J Maddison, and Tatsunori Hashimoto. Identifying the risks of LM agents with an LM-emulated sandbox. *arXiv [cs.AI]*, September 2023. URL `http://arxiv.org/abs/2309.15817`.

[568] Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. Survey of hallucination in natural language generation. *ACM Comput. Surv.*, 55(12):1–38, March 2023. ISSN 0360-0300. doi:10.1145/3571730. URL `https://doi.org/10.1145/3571730`.

[569] Zefan Cai, Baobao Chang, and Wenjuan Han. Human-in-the-loop through chain-of-thought. *arXiv [cs.CL]*, June 2023. URL `http://arxiv.org/abs/2306.07932`.

[570] Hengjia Xiao and Peng Wang. LLM a*: Human in the loop large language models enabled a* search for robotics. *arXiv [cs.RO]*, December 2023. URL `http://arxiv.org/abs/2312.01797`.

[571] Iddo Drori and Dov Te'eni. Human-in-the-Loop AI reviewing: Feasibility, opportunities, and risks. *Journal of the Association for Information Systems*, 25(1):98–109, 2024. ISSN 1536-9323. doi:10.17705/1jais.00867. URL `https://aisel.aisnet.org/jais/vol25/iss1/7/`.

[572] Nathan J. Szymanski, Bernardus Rendy, Yuxing Fei, Rishi E. Kumar, Tanjin He, David Milsted, Matthew J. McDermott, Max Gallant, Ekin Dogus Cubuk, Amil Merchant, Haegyeom Kim, Anubhav Jain, Christopher J. Bartel, Kristin Persson, Yan Zeng, and Gerbrand Ceder. An autonomous laboratory for the accelerated synthesis of novel materials. *Nature*, 624(7990):86–91, December 2023. ISSN 1476-4687. doi:10.1038/s41586-023-06734-w. URL `https://www.nature.com/articles/s41586-023-06734-w`. Publisher: Nature Publishing Group.

[573] Mark Peplow. Robot chemist sparks row with claim it created new materials. *Nature*, December 2023. doi:10.1038/d41586-023-03956-w. URL `https://www.nature.com/articles/d41586-023-03956-w`. Bandiera_abtest: a Cg_type: News Publisher: Nature Publishing Group Subject_term: Scientific community, Materials science, Machine learning, Chemistry.

[574] Sirui Hong, Yizhang Lin, Bang Liu, Bangbang Liu, Binhao Wu, Danyang Li, Jiaqi Chen, Jiayi Zhang, Jinlin Wang, Li Zhang, Lingyao Zhang, Min Yang, Mingchen Zhuge, Taicheng Guo, Tuo Zhou, Wei Tao, Wenyi Wang, Xiangru Tang, Xiangtao Lu, Xiawu Zheng, Xinbing Liang, Yaying Fei, Yuheng Cheng, Zongze Xu, and Chenglin Wu. Data interpreter: An LLM agent for data science. *arXiv [cs.AI]*, February 2024. URL `http://arxiv.org/abs/2402.18679`.

[575] Danrui Qi and Jiannan Wang. CleanAgent: Automating data standardization with LLM-based agents. *arXiv [cs.LG]*, March 2024. URL `http://arxiv.org/abs/2403.08291`.

[576] Xin Yee Tai, Hao Zhang, Zhiqiang Niu, Steven D R Christie, and Jin Xuan. The future of sustainable chemistry and process: Convergence of artificial intelligence, data and hardware. *Energy and AI*, 2(100036):100036, November 2020. ISSN 2666-5468. doi:10.1016/j.egyai.2020.100036. URL `http://dx.doi.org/10.1016/j.egyai.2020.100036`.

[577] Chenyu Zheng, Ishita Jalan, Patrick Cost, Kyle Oliver, Anju Gupta, Scott Misture, Jeremy A. Cody, and Christopher J. Collison. Impact of Alkyl Chain Length on Small Molecule Crystallization and Nanomorphology in Squaraine-Based Solution Processed Solar Cells. *The Journal of Physical Chemistry C*, 121(14):7750–7760, April 2017. ISSN 1932-7447. doi:10.1021/acs.jpcc.7b01339. URL `http://dx.doi.org/10.1021/acs.jpcc.7b01339`.